

Outline

1. Nonlinear classifiers
2. Kernel trick and kernel SVM
3. **Ensemble Methods - Boosting, Random Forests**
4. Classification Summary

Ensemble Classifiers


- *Why trust only one expert?*
 - In real life, we may consult several experts, or go with the "wisdom of the crowd"
 - In machine learning, *why trust only one classifier?*
- Ensemble methods aim to combine multiple classifiers together to form a better classifier.
- Examples:
 - **boosting** - training multiple classifiers, each focusing on errors made by previous classifiers.
 - **bagging** - training multiple classifiers from random selection of training data

AdaBoost - Adaptive Boosting

- Base classifier is a "weak learner"
 - A simple classifier that can be slightly better than random chance (>50%)
 - Example: *decision stump classifier*
 - check if feature value is above (or below) a threshold.
 - $y = h(x) = \begin{cases} +1, & x_j \geq T \\ -1, & x_j < T \end{cases}$

In [4]: wlf1g

Out [4]:



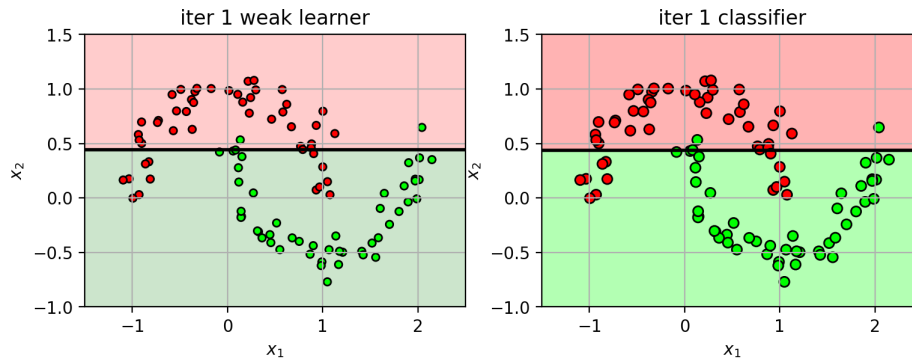
- **Idea:** train weak classifiers sequentially
- In each iteration,
 - Pick a weak learner $h_t(\mathbf{x})$ that best carves out the input space.
 - The weak learner should focus on data that is misclassified.
 - Apply weights to each sample in the training data.
 - Higher weights give more priority to difficult samples.
 - Combine all the weak learners into a strong classifier: $f_t(\mathbf{x}) = f_{t-1}(\mathbf{x}) + \alpha_t h_t(\mathbf{x})$
 - α_t is a weight for each weak learner.

Iteration 1

- Initially, weights for all training samples are equal: $w_i = 1/N$

In [7]: `plts[1]`

Out [7]:




Iteration 2 (part 1)

- points are re-weighted based on the previous errors:
 - increase weights for misclassified samples: $w_i = w_i e^{\alpha}$
 - decrease weights for correctly classified samples: $w_i = w_i e^{-\alpha}$
 - $\alpha = 0.5 \log \frac{1-err}{err}$ is based on the weighted error of the previous weak learner.
 - (larger circles indicates higher weight)

In [8]: `plts[2]`

Out [8]:

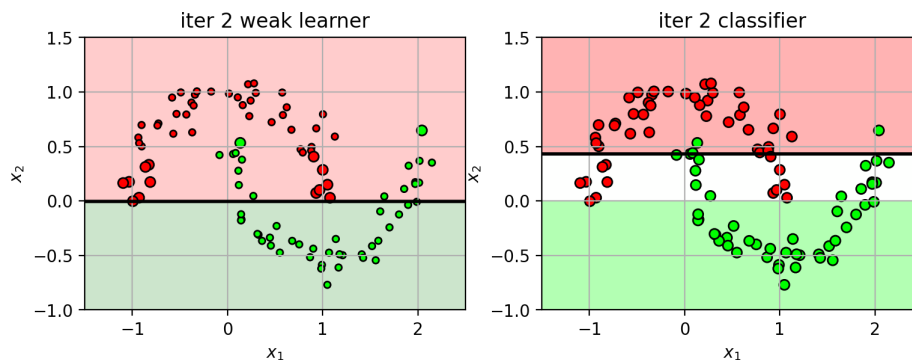


Iteration 2 (part 2)

- using the weighted data, train another weak learner $h_2(\mathbf{x})$.
- the classifier function is the weighted sum of weak learners
 - $f_2(\mathbf{x}) = f_1(\mathbf{x}) + \alpha_2 h_2(\mathbf{x})$

In [9]: `plts[2]`


Out [9]:



Keep iterating...


In [10]: `plts[3]`

Out[10]:




In [11]: `plts[4]`

Out[11]:



In [12]: `plts[5]`

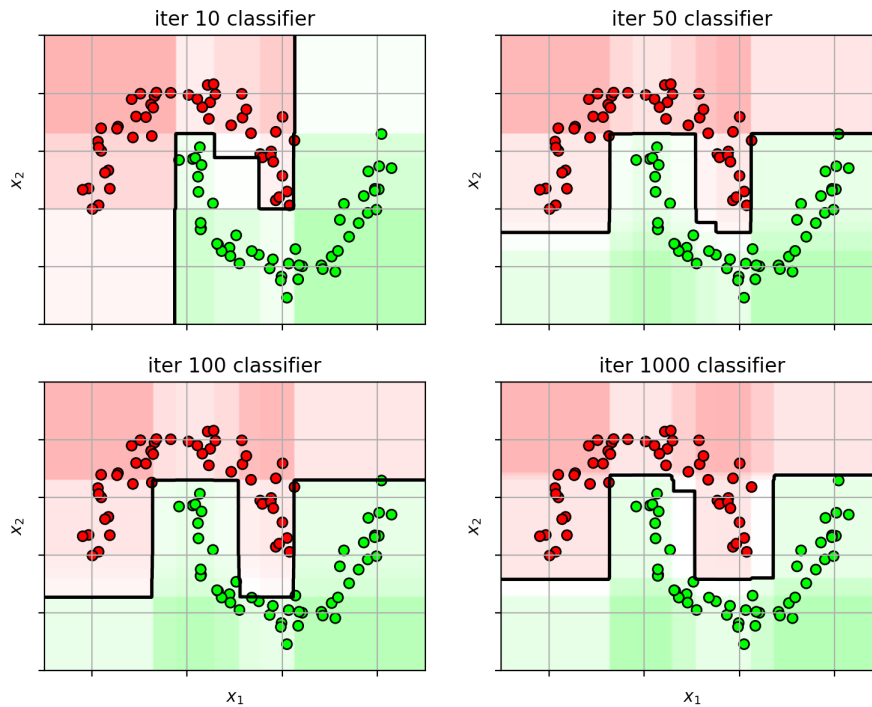
Out[12]:



- After many iterations...

In [14]: `adafig`

Out [14]:



Adaboost Algorithm

- Given data $\{(\mathbf{x}_i, y_i)\}$.
- Initialize data weights, $w_i = 1/N, \forall i$.
- For $t = 1$ to T ,
 - choose weak learner $h_t(\mathbf{x})$
 - minimize the weighted classification error: $\epsilon_t = \sum_{i=1}^N w_i 1(h_t(\mathbf{x}_i) \neq y_i)$.
 - Set the weak learner weight: $\alpha_t = \frac{1}{2} \log(\frac{1-\epsilon_t}{\epsilon_t})$
 - Add to ensemble: $f_t(\mathbf{x}) = f_{t-1}(\mathbf{x}) + \alpha_t h_t(\mathbf{x})$.
 - Update data weights:
 - for all \mathbf{x}_i misclassified, increase weight: $w_i \leftarrow w_i e^{\alpha_t}$.
 - for all \mathbf{x}_i correctly classified, decrease weight: $w_i \leftarrow w_i e^{-\alpha_t}$.
 - normalize weights, so that $\sum_i w_i = 1$.

Adaboost loss function

- It can be shown that Adaboost is minimizing:

$$\min_f \sum_i e^{-y_i f(\mathbf{x}_i)}$$

- Thus, it is an exponential loss function
 - $L(z_i) = e^{-z_i}$
 - $z_i = y_i f(\mathbf{x}_i)$
 - very sensitive to misclassified outliers.

In [16]: `lossfig`

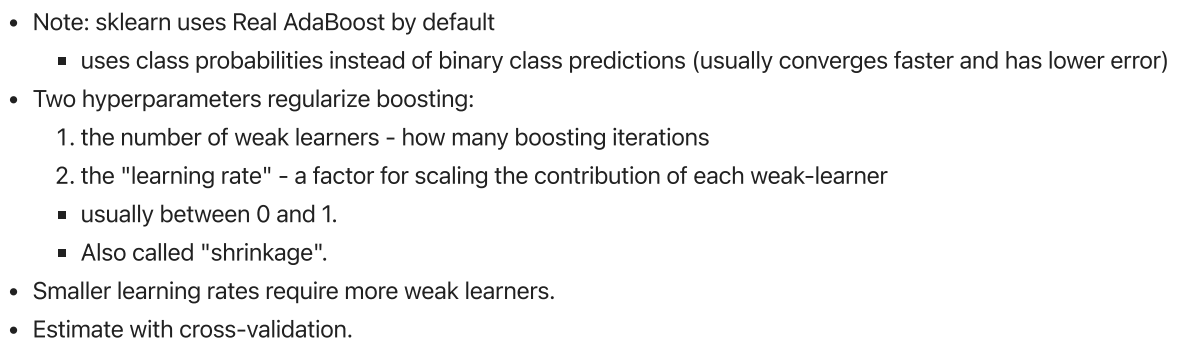
The graph shows the exponential loss function $\text{loss}(z_i)$ plotted against z_i . A vertical dashed line at $z_i = 0$ separates the "incorrectly classified" region (left) from the "correctly classified" region (right). The loss is high for misclassified samples and decreases exponentially for correctly classified samples.

Annotations on the graph include:

- exponential loss** (title)
- incorrectly classified** (left of $z_i = 0$)
- correctly classified** (right of $z_i = 0$)
- loss increases significantly for misclassified samples** (pointing to the steep part of the curve on the left)
- non-zero loss for samples near margin** (pointing to the curve just right of $z_i = 0$)
- loss approaches zero for correctly classified samples** (pointing to the curve far right of $z_i = 0$)

- Too many weak-learners and AdaBoost carves out space for the outliers.

Out[20]:

[illegible]

```
# run cross-validation (train for each split)
adacv.fit(trainX, trainY);

print("best params:", adacv.best_params_)

{'learning_rate': array([1.00000000e-06, 2.06913808e-06, 4.28133240e-06, 8.8586679
0e-06,
      1.83298071e-05, 3.79269019e-05, 7.84759970e-05, 1.62377674e-04,
      3.35981829e-04, 6.95192796e-04, 1.43844989e-03, 2.97635144e-03,
      6.15848211e-03, 1.27427499e-02, 2.63665090e-02, 5.45559478e-02,
      1.12883789e-01, 2.33572147e-01, 4.83293024e-01, 1.00000000e+00]), 'n_estima
tors': array([ 1,  2,  3,  5, 10, 15, 20, 25, 50, 100, 200,
      500, 1000])}
best params: {'learning_rate': 1e-06, 'n_estimators': 1}
```


```
In [24]: # predict from the model
predY = adacv.predict(testX)

# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy =", acc)


test accuracy = 0.82
```

```
In [26]: ifig2
```

```
Out[26]:
```




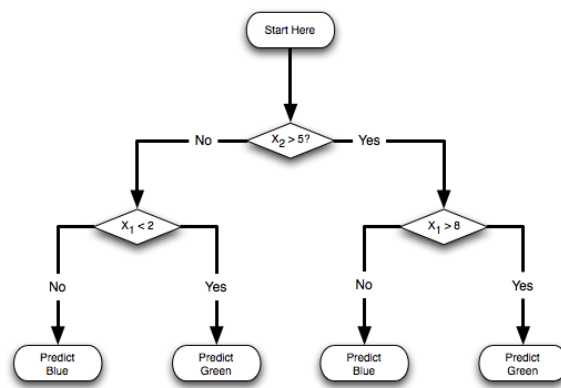
- Boosting can do feature selection
 - each decision stump classifier looks at one feature
- One of the original face detection methods (Viola-Jones) used Boosting.
 - extract a lot of image features from the face
 - during training, Boosting learns which ones are the most useful.



Gradient Boosting

- Variant of boosting
 - each iteration fits the residual between the current predictions and the true labels.

- the residual is computed as the gradient of the loss function.
- It's a gradient descent algorithm
 - in each iteration, the weak learner fits the gradient of the loss
 - $h_t(\mathbf{x}) \approx \frac{dL}{d\mathbf{f}}$
- and adds it to the function:
 - $f_t(\mathbf{x}) = f_{t-1}(\mathbf{x}) - \alpha_t h_t(\mathbf{x}) \approx f_{t-1}(\mathbf{x}) - \alpha_t \frac{dL}{d\mathbf{f}_{t-1}}$
- Generalizes boosting to other loss functions
- Typically uses decision trees for the weak learner:
 - At each node, move down the tree based on that node's criteria.
 - leaf node contains the prediction

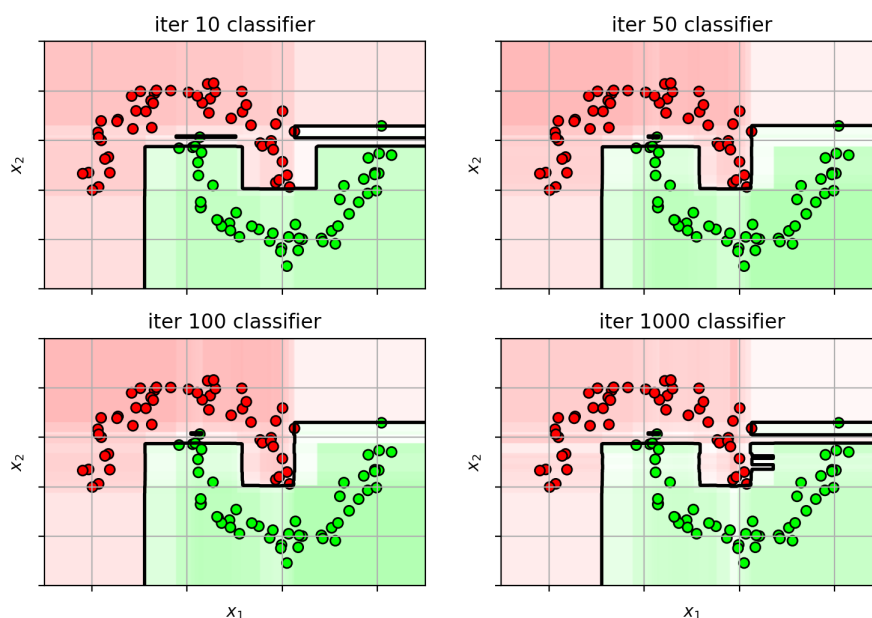


Example

- more iterations tends to overfit severely
 - because the "weak" classifier is actually strong (decision tree).

In [28]: `xgbfig`

Out [28]:



Cross-validation

- select the best hyperparameters
 - number of estimators
 - learning rate (shrinkage term)

```
In [29]: # use the XGBoost package, compatible with sklearn
import xgboost as xgb

# use "multi:softprob" for multi-class classification
xclf = xgb.XGBClassifier(objective="binary:logistic", eval_metric='logloss',
                        random_state=4487, use_label_encoder=False)

# setup the list of parameters to try
paramgrid = {'learning_rate': logspace(-6,0,20),
             'n_estimators': array([1, 2, 3, 5, 10, 15, 20, 25, 50, 100, 200, 500, 1000])
            }
print(paramgrid)

# setup the cross-validation object
xgbcv = model_selection.GridSearchCV(xclf, paramgrid, cv=5, n_jobs=-1)

# run cross-validation (train for each split)
xgbcv.fit(X3, Y3);

print("best params:", xgbcv.best_params_)

{'learning_rate': array([1.00000000e-06, 2.06913808e-06, 4.28133240e-06, 8.8586679
0e-06,
      1.83298071e-05, 3.79269019e-05, 7.84759970e-05, 1.62377674e-04,
      3.35981829e-04, 6.95192796e-04, 1.43844989e-03, 2.97635144e-03,
      6.15848211e-03, 1.27427499e-02, 2.63665090e-02, 5.45559478e-02,
      1.12883789e-01, 2.33572147e-01, 4.83293024e-01, 1.00000000e+00]), 'n_estima
tors': array([ 1,    2,    3,    5,   10,   15,   20,   25,   50,  100,  200,
      500, 1000])}
```

```

/Users/abc/miniforge3/envs/py39np/lib/python3.9/site-packages/xgboost/compat.py:3
6: FutureWarning: pandas.Int64Index is deprecated and will be removed from pandas
in a future version. Use pandas.Index with the appropriate dtype instead.
    from pandas import MultiIndex, Int64Index
/Users/abc/miniforge3/envs/py39np/lib/python3.9/site-packages/xgboost/compat.py:3
6: FutureWarning: pandas.Int64Index is deprecated and will be removed from pandas
in a future version. Use pandas.Index with the appropriate dtype instead.
    from pandas import MultiIndex, Int64Index
/Users/abc/miniforge3/envs/py39np/lib/python3.9/site-packages/xgboost/compat.py:3
6: FutureWarning: pandas.Int64Index is deprecated and will be removed from pandas
in a future version. Use pandas.Index with the appropriate dtype instead.
    from pandas import MultiIndex, Int64Index
/Users/abc/miniforge3/envs/py39np/lib/python3.9/site-packages/xgboost/compat.py:3
6: FutureWarning: pandas.Int64Index is deprecated and will be removed from pandas
in a future version. Use pandas.Index with the appropriate dtype instead.
    from pandas import MultiIndex, Int64Index
/Users/abc/miniforge3/envs/py39np/lib/python3.9/site-packages/xgboost/compat.py:3
6: FutureWarning: pandas.Int64Index is deprecated and will be removed from pandas
in a future version. Use pandas.Index with the appropriate dtype instead.
    from pandas import MultiIndex, Int64Index
/Users/abc/miniforge3/envs/py39np/lib/python3.9/site-packages/xgboost/compat.py:3
6: FutureWarning: pandas.Int64Index is deprecated and will be removed from pandas
in a future version. Use pandas.Index with the appropriate dtype instead.
    from pandas import MultiIndex, Int64Index
/Users/abc/miniforge3/envs/py39np/lib/python3.9/site-packages/xgboost/compat.py:3
6: FutureWarning: pandas.Int64Index is deprecated and will be removed from pandas
in a future version. Use pandas.Index with the appropriate dtype instead.
    from pandas import MultiIndex, Int64Index

```



```

in a future version. Use pandas.Index with the appropriate dtype instead.
from pandas import MultiIndex, Int64Index
/Users/abc/miniforge3/envs/py39np/lib/python3.9/site-packages/xgboost/compat.py:3
6: FutureWarning: pandas.Int64Index is deprecated and will be removed from pandas
in a future version. Use pandas.Index with the appropriate dtype instead.
from pandas import MultiIndex, Int64Index


```

```
best params: {'learning_rate': 0.012742749857031322, 'n_estimators': 1000}
```

```


In [30]: (avgscores, pnames, bestind) = extract_grid_scores(xgbcv, paramgrid)
paramfig = plt.figure()
plt.imshow(avgscores, interpolation='nearest')
plt.plot(bestind[1], bestind[0], '*w', markersize=12)
plt.ylabel(pnames[0] + ' index'); plt.xlabel(pnames[1] + ' index')
plt.grid(True)
plt.title('accuracy for different parameters')
plt.colorbar()
plt.axis('image');

```



```
In [32]: ifig2
```

```
Out[32]:
```



- Since decision trees are used, there are a lot of hyperparameters to tune for the decision tree.
 - `max_depth` : maximum depth of the tree
 - `gamma` : minimum loss reduction in order to split a leaf.
 - `colsample_bytree` : fraction of features to randomly subsample when building a tree.
 - `subsample` : fraction of training data to subsample during each boosting iteration (for each tree).
- **Problem:** Too many parameters to use grid-search!
- **Solution:** use randomized search
 - specify probability distributions for the parameters to try
 - `stats.uniform(a, b)` = uniform distribution between [a, a+b]
 - `stats.randint(a,b)` = random integer between [a, b]

```

In [33]: # setup dictionary of distributions for each parameter
paramsampler = {
    "colsample_bytree": stats.uniform(0.7, 0.3), # default=1

```

```

"gamma": stats.uniform(0, 0.5), # default=0
"max_depth": stats.randint(2, 6), # default=6
"subsample": stats.uniform(0.6, 0.4), # default=1
"learning_rate": stats.uniform(.001, 1), # default=1 (could also use loguniform)
"n_estimators": stats.randint(10, 1000),
}

xclf = xgb.XGBClassifier(objective="binary:logistic", eval_metric='logloss',
                        random_state=4487, use_label_encoder=False)

# cross-validation via random search
# n_iter = number of parameter combinations to try
xgbrcv = model_selection.RandomizedSearchCV(xclf, param_distributions=paramsampler,
                                           random_state=4487, n_iter=200, cv=5,
                                           verbose=1, n_jobs=-1)


xgbrcv.fit(X3, Y3)
print("best params:", xgbrcv.best_params_)

```

Fitting 5 folds for each of 200 candidates, totalling 1000 fits
best params: {'colsample_bytree': 0.9682214619643752, 'gamma': 0.4341101816965796
7, 'learning_rate': 0.014847933781299671, 'max_depth': 4, 'n_estimators': 152, 'su
bsample': 0.6743715045033899}

In [35]: ifig2

Out[35]:



Boosting Summary


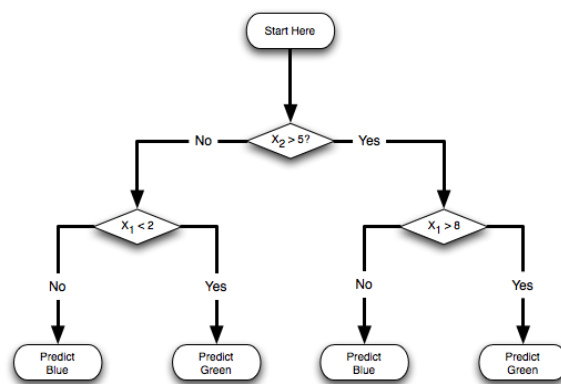
- **Ensemble Classifier:**
 - Combine the outputs of many "weak" classifiers to make a "strong" classifier
- **Training:**
 - In each iteration,
 - training data is re-weighted based on whether it is correctly classified or not.
 - weak classifier focuses on misclassified data from previous iterations.
 - Use cross-validation to pick number of weak learners.
- **Advantages:**
 - Good generalization performance
 - Built-in features selection - decision stump selects one feature at a time.
- **Disadvantages:**
 - Sensitive to outliers.

Outline

1. Nonlinear classifiers
2. Kernel trick and kernel SVM
3. **Ensemble Methods - Boosting, Random Forests**
4. Classification Summary


Decision Tree

- Simple "Rule-based" classifier
 - At each node, move down the tree based on that node's criteria.
 - leaf node contains the prediction
- **Advantage:** can create complex conjunction of rules
- **Disadvantage:** easy to overfit by itself
 - can fix with bagging!



Random Forest Classifier

- Use **bagging** to make an ensemble of Decision Tree Classifiers
 - for each *Decision Tree Classifier*
 - create a new training set by randomly sampling from the training set
 - for each split in a tree, select a random subset of features to use
- for a test sample, the prediction is aggregated over all trees.



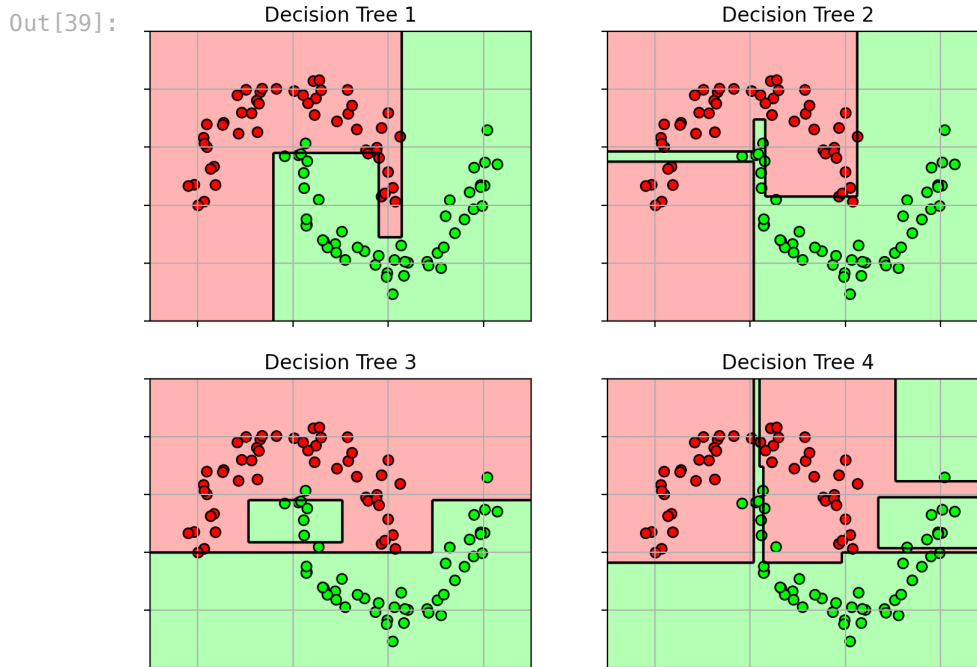
```
In [36]: # learn a RF classifier
         # use 4 trees
```

```
clf = ensemble.RandomForestClassifier(n_estimators=4, random_state=4487, n_jobs=-1)
clf.fit(X3, y3)
```

```
Out[36]: RandomForestClassifier
RandomForestClassifier(n_estimators=4, n_jobs=-1, random_state=4487)
```


- Here are the 4 decision trees
 - each uses a different random sampling of original training set

```
In [39]: dtfig
```



- and the aggregated classifier

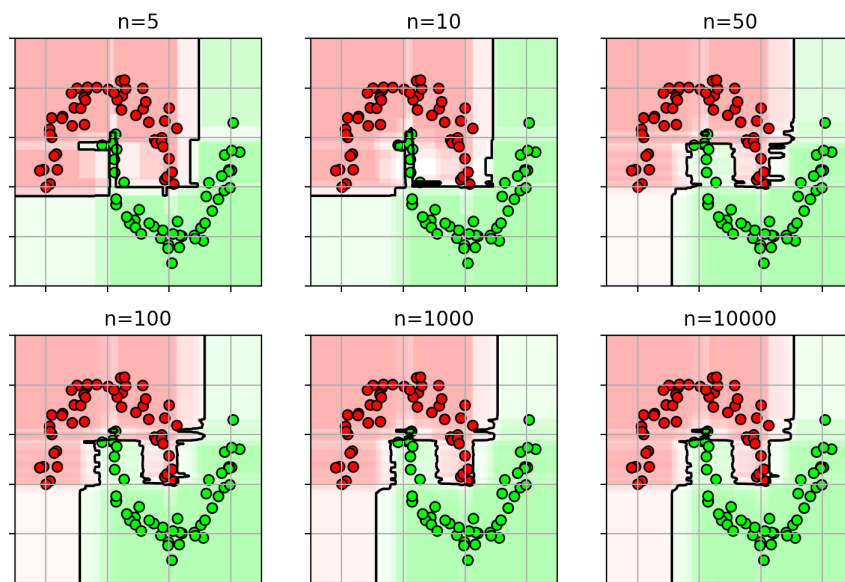
```
In [40]: rffig
```



- Using more trees

```
In [42]: rffig
```

Out [42]:



- Try on the iris data


In [43]:

```
# learn RF classifiers for different n_estimators
clfs = {}
for i,n in enumerate([5, 10, 50, 100, 500, 1000]):
    clfs[n] = ensemble.RandomForestClassifier(n_estimators=n, random_state=4487, n_jobs=-1)
    clfs[n].fit(trainX, trainY)
```

In [45]:

rfnfig

Out [45]:



In [46]:

```
# predict from the model
predY = clfs[1000].predict(testX)


# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy =", acc)
```

test accuracy = 0.78

In [48]:

```
# classifier boundary w/ training and test data
ifig3
```

Out [48]:



- Important parameters for cross-validation
 - `max_features` - maximum number of features used for each split
 - `max_depth` - maximum depth of a decision tree
 - `min_samples_split` - minimum fraction of samples to split a node.
 - `min_samples_leaf` - min fraction of samples in a leaf node.

In [49]:

```
# setup the list of parameters to try
paramsampler = {'max_features': stats.uniform(0,1.0),
                'max_depth': stats.randint(1,5),
                'min_samples_split': stats.uniform(0,0.5),
                'min_samples_leaf': stats.uniform(0,0.5),
                }

# setup the cross-validation object
rfrcv = model_selection.RandomizedSearchCV(
    ensemble.RandomForestClassifier(n_estimators=100, random_state=4),
    param_distributions=paramsampler,
    random_state=4487, n_iter=1000, cv=5,
    verbose=1, n_jobs=-1)

# run cross-validation (train for each split)
rfrcv.fit(trainX, trainY);


print("best params:", rfrcv.best_params_)
```

Fitting 5 folds for each of 1000 candidates, totalling 5000 fits
best params: {'max_depth': 4, 'min_samples_leaf': 0.013009207368046005, 'min_samples_split': 0.033821887640189785}

- Result

In [51]: ifig3

Out [51]:



In [52]:

```
# predict from the model
predY = rfrcv.predict(testX)
```

```
# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy =", acc)

test accuracy = 0.8
```

Random Forest Summary

- **Ensemble Classifier & Training:**
 - aggregate predictions over several decision trees
 - trained using different subsets of data, and different subsets of features.
- **Advantages**
 - non-linear decision boundary.
 - can do feature selection.
 - good generalization.
 - fast.
- **Disadvantages**
 - can be sensitive to outliers
 - based on trees -- cannot well represent "diagonal" decision boundaries.