

# CS5489 - Machine Learning

## Lecture 4c - Classification Summary

Prof. Antoni B. Chan

Dept. of Computer Science, City University of Hong Kong

### Outline

1. Nonlinear classifiers
2. Kernel trick and kernel SVM
3. Ensemble Methods - Boosting, Random Forests
4. **Classification Summary**


### Feature Pre-processing

- Some classifiers, such as SVM and LR, are sensitive to the scale of the feature values.
  - feature dimensions with larger values may dominate the objective function.
- Common practice is to *standardize* or *normalize* each feature dimension before learning the classifier.
  - Two Methods...
- **Method 1:** scale each feature dimension so the mean is 0 and variance is 1.
  - $\tilde{x}_d = \frac{1}{s}(x_d - m)$
  - $s$  is the standard deviation of feature values.
  - $m$  is the mean of the feature values.
- **NOTE:** the parameters for scaling the features should be estimated from the training set!
  - same scaling is applied to the test set.

```
In [5]: # using the iris data
scaler = preprocessing.StandardScaler() # make scaling object
trainXn = scaler.fit_transform(trainX) # use training data to fit scaling parameters
testXn = scaler.transform(testX) # apply scaling to test data
```

```
In [7]: nfig1
```


```
Out [7]:
```



- **Method 2:** scale features to a fixed range, -1 to 1.
  - $\tilde{x}_d = 2 * (x_d - min) / (max - min) - 1$
  - $max$  and  $min$  are the maximum and minimum features values.

```
In [8]: # using the iris data
scaler = preprocessing.MinMaxScaler(feature_range=(-1,1)) # make scaling object
trainXn = scaler.fit_transform(trainX) # use training data to fit scaling parameters
testXn = scaler.transform(testX) # apply scaling to test data
```

```
In [10]: nfig2
```



## Data Representation and Feature Engineering

- How to represent data as a vector of numbers?
  - the encoding of the data into a feature vector should make sense
  - inner-products or distances calculated between feature vectors should be meaningful in terms of the data.
- Categorical variables
  - Example:  $x$  has 3 possible category labels: cat, dog, horse
  - We could encode this as:  $x = 0$ ,  $x = 1$ , and  $x = 2$ .
    - Suppose we have two data points:  $x = \text{cat}$ ,  $x' = \text{horse}$ .
    - What is the meaning of  $x * x' = 2$ ?

## One-hot encoding

- encode a categorical variable as a vector of ones and zeros
  - if there are  $K$  categories, then the vector is  $K$  dimensions.
- Example:
  - $x = \text{cat} \rightarrow x = [1 \ 0 \ 0]$
  - $x = \text{dog} \rightarrow x = [0 \ 1 \ 0]$
  - $x = \text{horse} \rightarrow x = [0 \ 0 \ 1]$

```
In [11]: # one-hot encoding example
X = [['cat'], ['dog'], ['cat'], ['bird'], ['dog']] # each row is a sample
ohe = preprocessing.OneHotEncoder(sparse=False)
ohe.fit(X) # map the categories to one-hot vectors
print(ohe.categories_)
ohe.transform(X) # transform to one-hot-encoding

[array(['bird', 'cat', 'dog'], dtype=object)]
```

```
Out[11]: array([[0., 1., 0.],
                [0., 0., 1.],
                [0., 1., 0.],
                [1., 0., 0.],
                [0., 0., 1.]])
```

## Binning

- encode a real value as a vector of ones and zeros
  - assign each feature value to a bin, and then use one-hot-encoding

```
In [12]: # example
X = [[-3], [0.5], [1.5], [2.5]] # the data
bins = [-2,-1,0,1,2] # define the bin edges

# map from value to bin number
Xbins = digitize(X, bins=bins)

# map from bin number (0..5) to 0-1 vector
ohe = preprocessing.OneHotEncoder(categories=[arange(6)], sparse=False)
ohe.fit(Xbins)
ohe.transform(Xbins)
```

```
Out[12]: array([[1., 0., 0., 0., 0., 0.],
               [0., 0., 0., 1., 0., 0.],
               [0., 0., 0., 0., 1., 0.],
               [0., 0., 0., 0., 0., 1.]])
```

## Data transformations - polynomials

- Represent interactions between features using polynomials
- Example:
  - 2nd-degree polynomial models pair-wise interactions
    - $[x_1, x_2] \rightarrow [x_1^2, x_1x_2, x_2^2]$
  - Combine with other degrees:
    - $[x_1, x_2] \rightarrow [1, x_1, x_2, x_1^2, x_1x_2, x_2^2]$

```
In [13]: X = [[0,1], [1,2], [3,4]]
pf = preprocessing.PolynomialFeatures(degree=2)
pf.fit(X)
pf.transform(X)
```

```
Out[13]: array([[ 1.,  0.,  1.,  0.,  0.,  1.],
               [ 1.,  1.,  2.,  1.,  2.,  4.],
               [ 1.,  3.,  4.,  9., 12., 16.]])
```

## Data transformations - univariate


- Apply a non-linear transformation to the feature
  - e.g.,  $x \rightarrow \log(x)$
  - useful if the dynamic range of  $x$  is very large

## Unbalanced Data

- For some classification tasks that data will be unbalanced
  - many more examples in one class than the other.
- **Example:** detecting credit card fraud
  - credit card fraud is rare
    - 50 examples of fraud, 5000 examples of legitimate transactions.

```
In [15]: udatafig
```


Out [15]:



- Unbalanced data can cause problems when training the classifier
  - classifier will focus more on the class with more points.
  - decision boundary is pushed away from class with more points

In [17]: udatafig1

Out [17]:



- **Solution:** apply weights on the classes during training.
  - weights are inversely proportional to the class size.


```
In [18]: clf = svm.SVC(kernel='linear', C=10, class_weight='balanced')
clf.fit(X, Y)

print("class weights =", clf.class_weight_)

class weights = [0.55 5.5 ]
```

In [20]: udatafig2

Out [20]:




## Classifier Imbalance

- In some tasks, errors on certain classes cannot be tolerated.

- **Example:** detecting spam vs non-spam
  - non-spam should *definitely not* be marked as spam
    - okay to mark some spam as non-spam

In [22]: udatafig3




...

- Class weighting can be used to make the classifier focus on certain classes
  - e.g., weight non-spam class higher than spam class
    - classifier will try to correctly classify all non-spam samples, at the expense of making errors on spam samples.

```
In [23]: # dictionary (key,value) = (class name, class weight)
cw = {0: 0.2,
      1: 5} # class 1 is 25 times more important!

clf = svm.SVC(kernel='linear', C=10, class_weight=cw)
clf.fit(X, Y);
```

In [25]: udatafig4



...

## Classification Summary

- **Classification task**
  - Observation  $\mathbf{x}$ : typically a real vector of feature values,  $\mathbf{x} \in \mathbb{R}^d$ .
  - Class  $y$ : from a set of possible classes, e.g.,  $\mathcal{Y} = \{0, 1\}$
  - **Goal:** given an observation  $\mathbf{x}$ , predict its class  $y$ .

Name	Type	Classes	Decision function	Training	Advantages	Disadvantages
Bayes' classifier	generative	multi-class	non-linear	estimate class-conditional densities $p(\mathbf{x} y)$ by maximizing likelihood of data.	- works well with small amounts of data. - multi-class.	- depends on the data correctly fitting the class-conditional.


					- minimum probability of error if probability models are correct.	
logistic regression	discriminative	binary	linear	maximize likelihood of data in $p(y x)$ .	- well-calibrated probabilities. - efficient to learn.	- linear decision boundary. - sensitive to $C$ parameter.
support vector machine (SVM)	discriminative	binary	linear	maximize the margin (distance between decision surface and closest point).	- works well in high-dimension. - good generalization.	- linear decision boundary. - sensitive to $C$ parameter.
kernel SVM	discriminative	binary	non-linear (kernel function)	maximize the margin.	- non-linear decision boundary. - can be applied to non-vector data using appropriate kernel.	- sensitive to kernel function and hyperparameters. - high memory usage for large datasets
AdaBoost	discriminative	binary	non-linear (ensemble of weak learners)	train successive weak learners to focus on misclassified points.	- non-linear decision boundary. - can do feature selection. - good generalization.	- sensitive to outliers.
XGBoost	discriminative	binary	non-linear (ensemble of decision trees)	train successive learners to focus on gradient of the loss.	- non-linear decision boundary. - good generalization.	- sensitive to outliers.
Random Forest	discriminative	multi-class	non-linear (ensemble of)	aggregate predictions over several decision trees, trained using different	- non-linear decision boundary. - can do feature selection. - good generalization.	- sensitive to outliers.

## Loss functions

- The classifiers differ in their loss functions, which influence how they work.
  - $z_i = y_i f(\mathbf{x}_i)$

In [27]: `lossfig`

Out [27]:

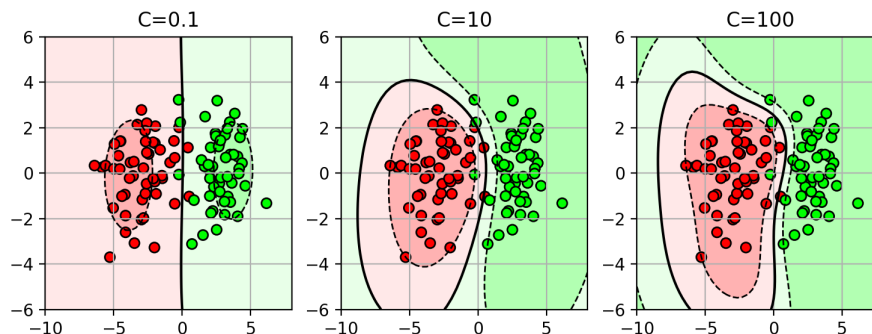


## Regularization and Overfitting

- Some models have terms to prevent overfitting the training data.
  - this can improve *generalization* to new data.
- There is a parameter to control the regularization effect.
  - select this parameter using cross-validation on the training set.

In [29]: `ofig`

Out [29]:



## Structural Risk Minimization

- A general framework for balancing data fit and model complexity.

- Many learning problems can be written as a combination of data-fit and regularization term:

$$f^* = \operatorname{argmin}_f \sum_i L(y_i, f(\mathbf{x}_i)) + \lambda \Omega(f)$$

- assume  $f$  within some class of functions, e.g., linear functions  $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$ .
- $L$  is the loss function, e.g., logistic loss.
- $\Omega$  is the regularization function on  $f$ , e.g.,  $\|\mathbf{w}\|^2$
- $\lambda$  is the tradeoff parameter, e.g.,  $1/C$ .

## Other things

- *Multiclass classification*
  - can use binary classifiers to do multi-class using *1-vs-rest* formulation.
- *Feature normalization*
  - normalize each feature dimension so that some feature dimensions with larger ranges do not dominate the optimization process.
- *Unbalanced data*
  - if more data in one class, then apply weights to each class to balance objectives.
- *Class imbalance*
  - mistakes on some classes are more critical.
  - reweight class to focus classifier on correctly predicting one class at the expense of others.

## Applications

- Web document classification, spam classification
- Face gender recognition, face detection, digit classification

## Features

- Choice of features is important!
  - using uninformative features may confuse the classifier.
  - use domain knowledge to pick the best features to extract from the data.

## Which classifier is best?

- **"No Free Lunch" Theorem** (Wolpert and Macready)

"If an algorithm performs well on a certain class of problems then it necessarily pays for that with degraded performance on the set of all remaining problems."

- In other words, there is no *best* classifier for all tasks. The best classifier depends on the particular problem.