

# CS5489 - Machine Learning

## Lecture 7a - Unsupervised Learning - Clustering

Prof. Antoni B. Chan

Dept. of Computer Science, City University of Hong Kong

### Outline


1. Unsupervised Learning
2. Parametric clustering
  - A. K-means
  - B. Gaussian mixture models (GMMs)
  - C. Bayesian GMMs
3. Non-parametric clustering and Mean-shift
4. Spectral clustering

### Unsupervised Learning

- Unsupervised learning only considers the input data  $\mathbf{x}$ .
  - There are no output values.
- **Goal:** Try to discover inherent properties in the data.
  - Clustering
  - Dimensionality Reduction
  - Manifold Embedding

### Clustering

- Find clusters of similar items in the data.
  - Find a representative item that can represent all items in the cluster.
- **For example:** grouping iris flowers by their measurements.
  - Features are sepal width and petal length.




### Clustering

- Data is set of vectors  $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ 
  - Each data point is a vector  $\mathbf{x} \in \mathbb{R}^d$ .
- **Goal:** group similar data together.
  - groups are also called clusters.
  - each data point is assigned with a cluster index ( $y \in \{1, \dots, K\}$ )
    - $K$  is the number of clusters

In [3]: `clusterfig`

Out [3]:




## K-Means Clustering

- **Idea:**
  - there are  $K$  clusters.
  - each cluster is represented by a *cluster center*.
    - $\mathbf{c}_j \in \mathbb{R}^d, j \in \{1, \dots, K\}$
  - assign each data point to the closest cluster center.
    - assignment variable  $z_i \in \{1, \dots, K\}$  indexes the cluster center of  $\mathbf{x}_i$ .

In [5]: `kmfig`

Out [5]:



## K-means Clustering Objective

- **Objective:** minimize the total sum-squared difference between points and their centers

$$\min_{\mathbf{c}_1, \dots, \mathbf{c}_K, z_1, \dots, z_n} \sum_{i=1}^n \|\mathbf{x}_i - \mathbf{c}_{z_i}\|^2$$

## K-means Clustering Objective

- *How to pick the cluster centers?*

- Assume the assignments  $z_i$  are known.
- Pick the cluster centers that minimize the squared distance to all its cluster members.

$$\min_{\mathbf{c}_1, \dots, \mathbf{c}_K} \sum_{i=1}^n \|\mathbf{x}_i - \mathbf{c}_{z_i}\|^2$$

• Solution:

- if the assignments  $\{z_i\}$  are known...
  - let  $C_j$  be the set of points assigned to cluster  $j$ 
    - $C_j = \{\mathbf{x}_i | z_i = j\}$
  - For each cluster, we have
    - $\mathbf{c}_j = \operatorname{argmin} \sum_{i \in C_j} \|\mathbf{x}_i - \mathbf{c}_j\|^2$
  - Take the derivative and set to 0:

$$\begin{aligned} \frac{d}{d\mathbf{c}_j} \sum_{i \in C_j} \|\mathbf{x}_i - \mathbf{c}_j\|^2 &= \sum_{i \in C_j} 2(\mathbf{x}_i - \mathbf{c}_j)(-1) = 0 \\ \Rightarrow \sum_{i \in C_j} \mathbf{x}_i - |C_j| \mathbf{c}_j &= 0 \end{aligned}$$

- Cluster center is the mean of the points in the cluster
  - $\mathbf{c}_j = \frac{1}{|C_j|} \sum_{\mathbf{x}_i \in C_j} \mathbf{x}_i$

## K-means Clustering Objective

- How to pick the assignments?
  - Assume the clusters  $\{\mathbf{c}_1, \dots, \mathbf{c}_K\}$  are known.
  - Pick the assignments that minimize the squared distance to the clusters.

$$\min_{z_1, \dots, z_n} \sum_{i=1}^n \|\mathbf{x}_i - \mathbf{c}_{z_i}\|^2$$

- For each data point,
  - $z_i = \operatorname{argmin}_{j \in \{1, \dots, K\}} \|\mathbf{x}_i - \mathbf{c}_j\|^2$
  - i.e., assign point  $\mathbf{x}_i$  to its closest cluster.

## Chicken and Egg Problem


- Cluster assignment of each point depends on the cluster centers.
- Location of cluster center depends on which points are assigned to it.
- **Solution:** just iterate between the two steps.
  - Note: in each step we are holding one set of variables fixed, while minimizing over the others.
  - Thus we are always minimizing the original objective!

## K-means Algorithm

- Pick initial cluster centers
- Repeat:
  - 1. calculate assignment  $z_i$  for each point  $\mathbf{x}_i$ : closest cluster center using Euclidean distance.
  - 2. calculate cluster center  $\mathbf{c}_j$  as average of points assigned to cluster  $j$ .
- This procedure will converge eventually.

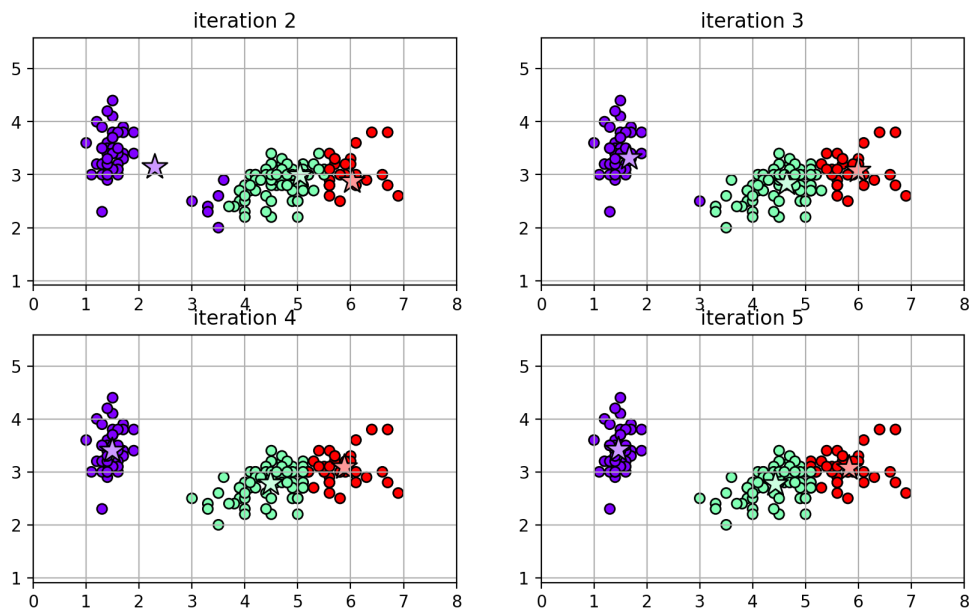
```
In [8]: kmifig1
```

```
Out[8]:
```



```
In [9]: kmifig2
```

```
Out[9]:
```




## Important Note

- The final result depends on the initial cluster centers!
  - Some bad initializations will yield poor clustering results!
  - (Technically, there are multiple local minimums in the objective function)

```
In [11]: kmbadfig
```

```
Out[11]:
```



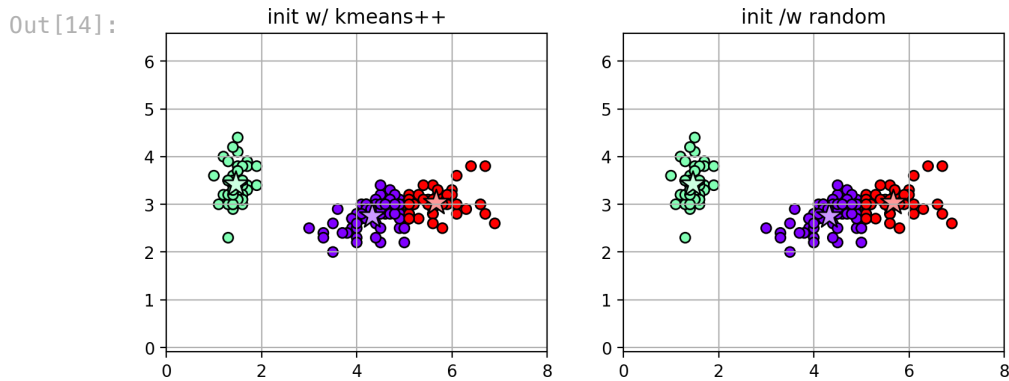
- **Solution:**
  - Try several times using different initializations.
  - Pick the answer with lowest objective score.
- In scikit-learn,
  - 1. use a smart initialization method called "k-means++", which speeds up convergence. (default)
  - 2. use multiple random initializations.

```
In [12]: # K-Means with 3 clusters (w/ Kmeans++ init)
km = cluster.KMeans(n_clusters=3, random_state=4487)
Yp = km.fit_predict(X) # cluster data, and return labels

cc = km.cluster_centers_ # the cluster centers
cl = km.labels_ # labels also stored here

# K-means w/ random initialization (automatically does 10 trials)
kmr = cluster.KMeans(n_clusters=3, random_state=4487, init='random')
Ypr = kmr.fit_predict(X) # cluster data, and return labels
```

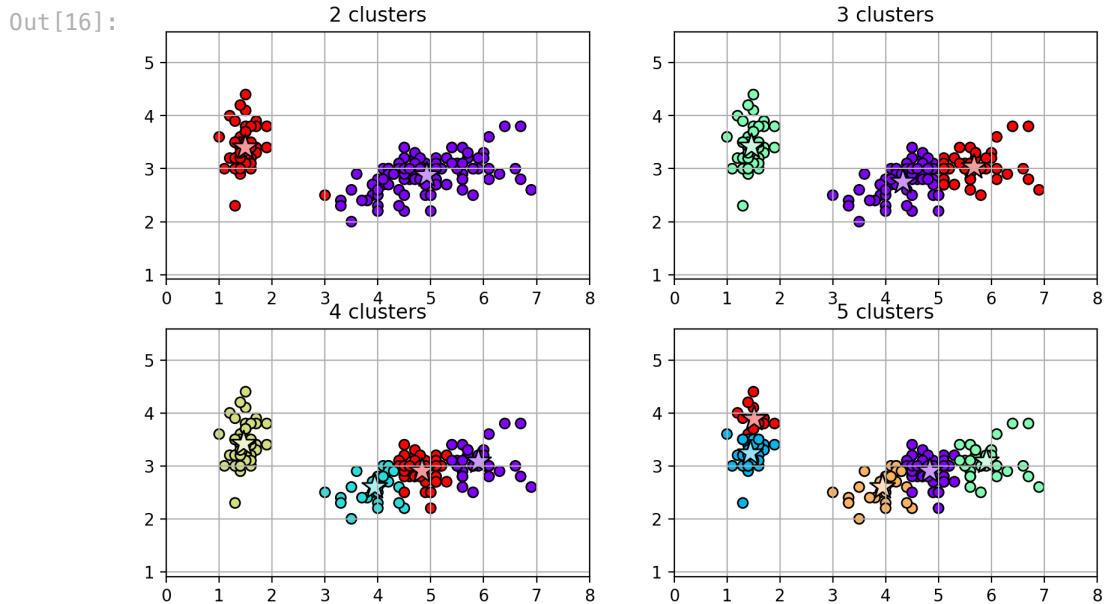
```
In [14]: kmfig
```



## For different K

- We need to choose the appropriate K

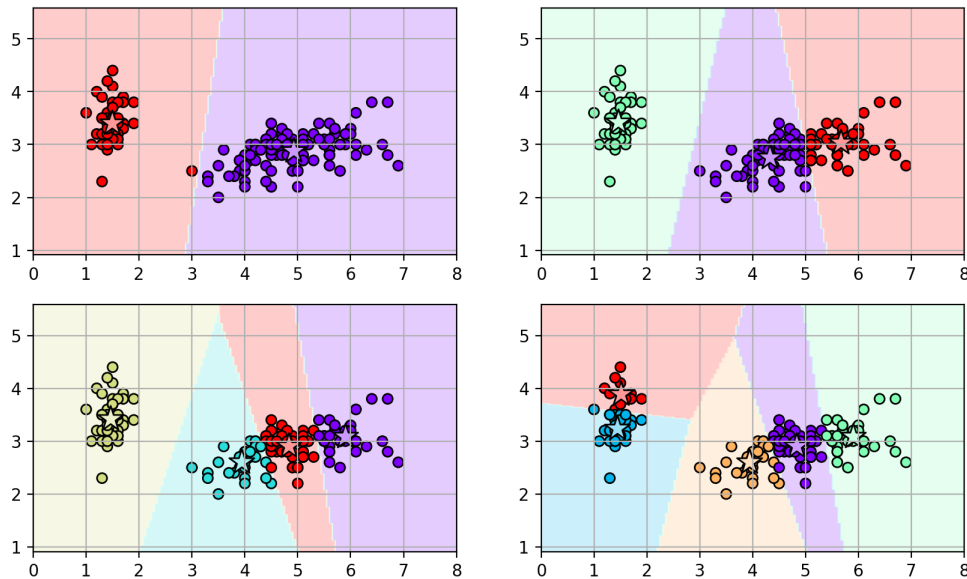
```
In [16]: ksfig
```



- K-means partitions the input space into non-overlapping regions belonging to each cluster
  - assign each location in the space to the closest center.

```
In [18]: krfig
```

Out [18]:



## Bag-of-X Representation

- K-means can partition the input space into regions.
- Create a new quantized representation for a set of samples.
  - 1. learn the partition space of the samples using K-means (i.e., the vocabulary)
  - 2. assign each sample to its closest center (i.e., the word)
  - 3. count the number of assignments for each center and form a histogram (i.e., the bag-of-words)
- Called a "bag-of-X" representation.
  - "X" is whatever modality you are using.

## Example: Bag-of-visual-words

- Procedure:
  - 1. extract small patches from images
  - 2. learn the visual words using K-means
  - 3. assign patches to visual words, and form a histogram for each image.
- Use the bag-of-words model as the new feature vector.

```
In [19]: # load images
oli = datasets.fetch_olivetti_faces(data_home=".")
img = oli.images
print(img.shape)

# extract 8x8 window patches with step size of 4x4
patches = skimage.util.view_as_windows(img, (1,8,8), step=(1,4,4))
print(patches.shape)


# reshape patches into 64-dim vectors
patchesall = patches.reshape((prod(patches.shape[0:3]),8*8))
print(patchesall.shape)

# run k-means (mini-batch version can handle large datasets)
# fit k-means, and predict the cluster index for each sample
K = 100
bows = cluster.MinibatchKMeans(n_clusters=K, random_state=5489, n_init=10)
wordsall = bows.fit_predict(patchesall)

(400, 64, 64)
(400, 15, 15, 1, 8, 8)
(90000, 64)
```

```
In [21]: # get the visual words, and show them
visualwords = bows.cluster_centers_.reshape((K,8,8))
```

```
plt.figure(figsize=(10,4))
plt.imshow(image_montage(visualwords, maxw=20), cmap='gray', interpolation='nearest');
```

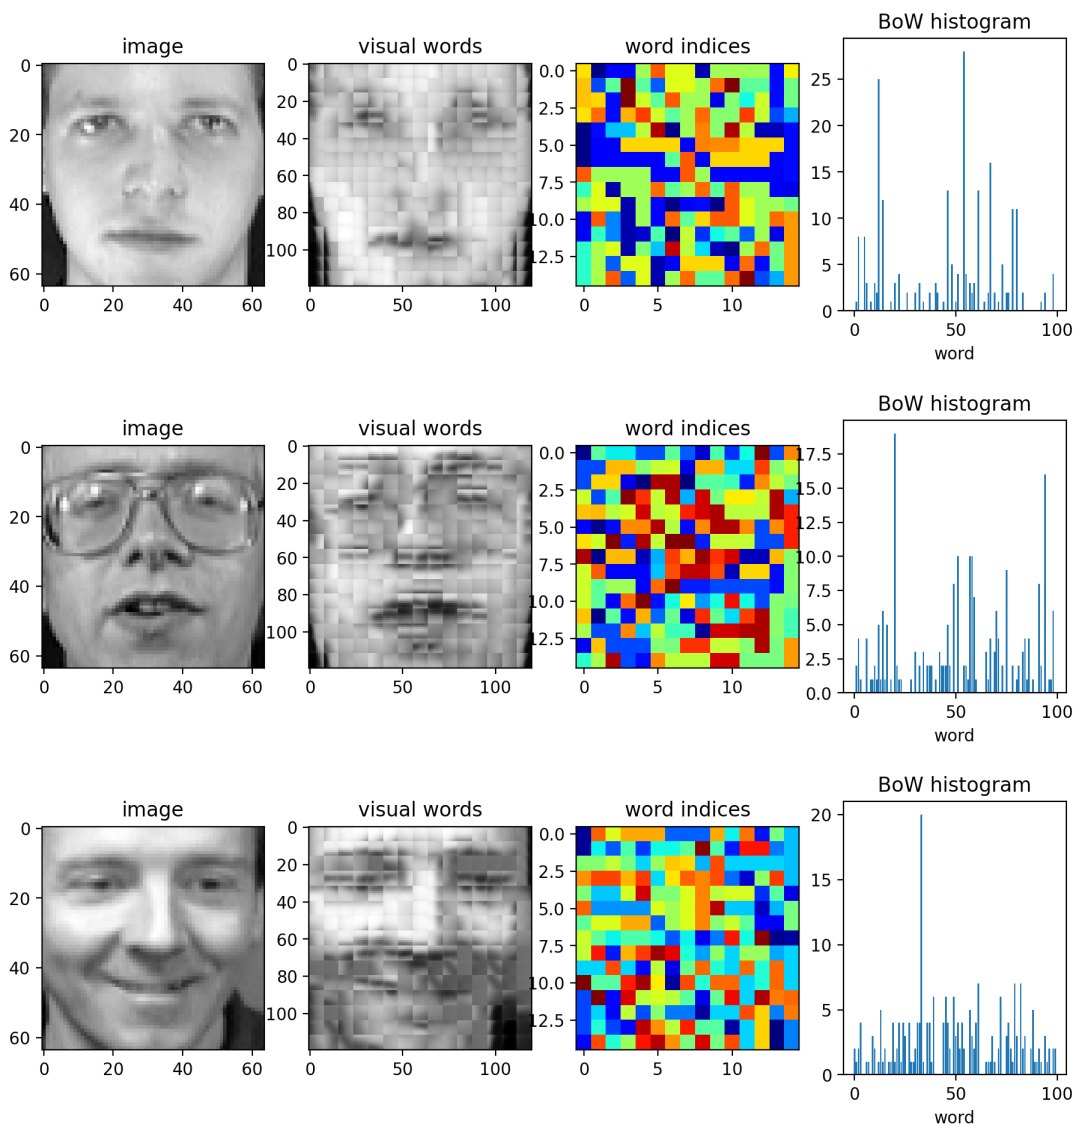


```
In [22]: # reshape the predicted words into small images
print(wordsall.shape)
patches_words = wordsall.reshape(patches.shape[0:3])
print(patches_words.shape)

# build the BoW histogram for each image
bowhist = [bincount(wds.ravel(), minlength=K) for wds in patches_words]

(90000,)
(400, 15, 15)
```

```
In [24]: # view some reconstructions, words, and histograms
for t in [0,10,20]:
    plt.figure(figsize=(11,3))
    show_bow_recon(visualwords, img[t], patches_words[t], bowhist[t])
    plt.show()
```




# Circular clusters

- One problem with K-means is that it assumes that each cluster has a circular shape.
  - based on Euclidean distance to each center
  - Kmeans cannot handle skewed (elliptical) clusters.

In [26]: efig

Out [26]:



## Gaussian mixture model (GMM)

- A multivariate Gaussian can model a cluster with an elliptical shape.
  - the ellipse shape is controlled by the covariance matrix of the Gaussian
  - the location of the cluster is controlled by the mean.
- Gaussian mixture model is a weighted sum of Gaussians

$$p(\mathbf{x}) = \sum_{j=1}^K \pi_j N(\mathbf{x} | \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)$$


- Each Gaussian represents one elliptical cluster
  - $\boldsymbol{\mu}_j$  is the mean of the j-th Gaussian. (the location)
  - $\boldsymbol{\Sigma}_j$  is the covariance matrix of the j-th Gaussian. (the ellipse shape)
  - $\pi_j$  is the prior weight of the j-th Gaussian. (how likely is this cluster)

## 1-D example of GMM

- $p(x) = 0.25 * N(x | -2, 4) + 0.75 * N(x | 3, 1)$
- each Gaussian is a "mountain"

In [28]: eggmm

Out [28]:



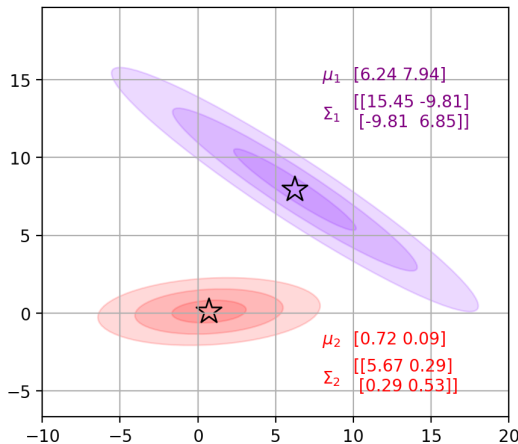
## 2D example of GMM



- Each Gaussian defines a "mountain"
  - contours are ellipses

In [30]: gmm2fig

Out[30]:



## Clustering with GMMs

- Using the data, learn a GMM using maximum likelihood estimation:

$$\max_{\pi, \mu, \Sigma} \sum_{i=1}^N \log \sum_{j=1}^K \pi_j N(\mathbf{x}_i | \mu_j, \Sigma_j)$$

- The learned  $\{\mu_j, \Sigma_j\}$  are the cluster center and ellipse shape.
- The learned  $\pi_j$  is the cluster population (percentage of points).
- This is difficult to optimize because the "sum" is inside the "log".

## Expectation Maximization (EM) Algorithm

- An algorithm for finding the MLE solution when there are hidden (unseen) variables  $z$ .
  - Goal:  $\hat{\theta} = \arg\max_{\theta} p_{\theta}(\mathbf{x})$ 
    - where  $p_{\theta}(\mathbf{x}) = \sum_z p(\mathbf{x}|z)p(z)$
    - $\theta$  are the parameters.
- Solution:
  - iterate between estimating the hidden variables  $z$  and maximizing w.r.t the parameters  $\theta$ .
- For GMMs:
  - $z$  is the assignment of  $\mathbf{x}$  to one of the Gaussian components.
  - prior probability:  $p(z = j) = \pi_j$
  - Gaussian component:  $p(\mathbf{x}|z = j) = \mathcal{N}(\mathbf{x} | \mu_j, \Sigma_j)$
  - likelihood:  $p(\mathbf{x}) = \sum_z p(\mathbf{x}|z)p(z) = \sum_j \pi_j \mathcal{N}(\mathbf{x} | \mu_j, \Sigma_j)$

## EM Algorithm

- 1. Select an initial model  $\hat{\theta}$
- 2. **E-step**: estimate the hidden variables as their expected value.
  - $Q(\theta) = E_{z|\mathbf{x}, \hat{\theta}}[\log p_{\theta}(\mathbf{x}, z)]$
  - the expectation uses current model parameters  $\hat{\theta}$ .
- 3. **M-step**: maximize w.r.t  $\theta$ 
  - $\hat{\theta} = \arg\max_{\theta} Q(\theta)$

- 4. repeat E- and M-steps until convergence.
  - It's guaranteed to converge.

## Joint likelihood for GMMs

- Joint likelihood of  $(\mathbf{x}_i, z_i)$

$$p(\mathbf{x}_i, z_i) = p(\mathbf{x}_i | z_i) p(z_i) = \pi_{z_i} \mathcal{N}(\mathbf{x}_i | \mu_{z_i}, \Sigma_{z_i})$$

- annoying because  $z_i$  is indexing the parameters.

## Indicator Variable Trick

- Define  $z_{ij}$  as the indicator variable that  $z_i = j$ 
  - $z_{ij} = \begin{cases} 1, & z_i = j \\ 0, & \text{otherwise} \end{cases}$
- Joint likelihood of  $(\mathbf{x}_i, z_i)$ 
  - $z_{ij}$  selects the correct Gaussian component.

$$p(\mathbf{x}_i, z_i) = \prod_{j=1}^K (p(\mathbf{x}_i | z_i = j) p(z_i = j))^{z_{ij}} = \prod_{j=1}^K (\pi_j \mathcal{N}(\mathbf{x}_i | \mu_j, \Sigma_j))^{z_{ij}}$$

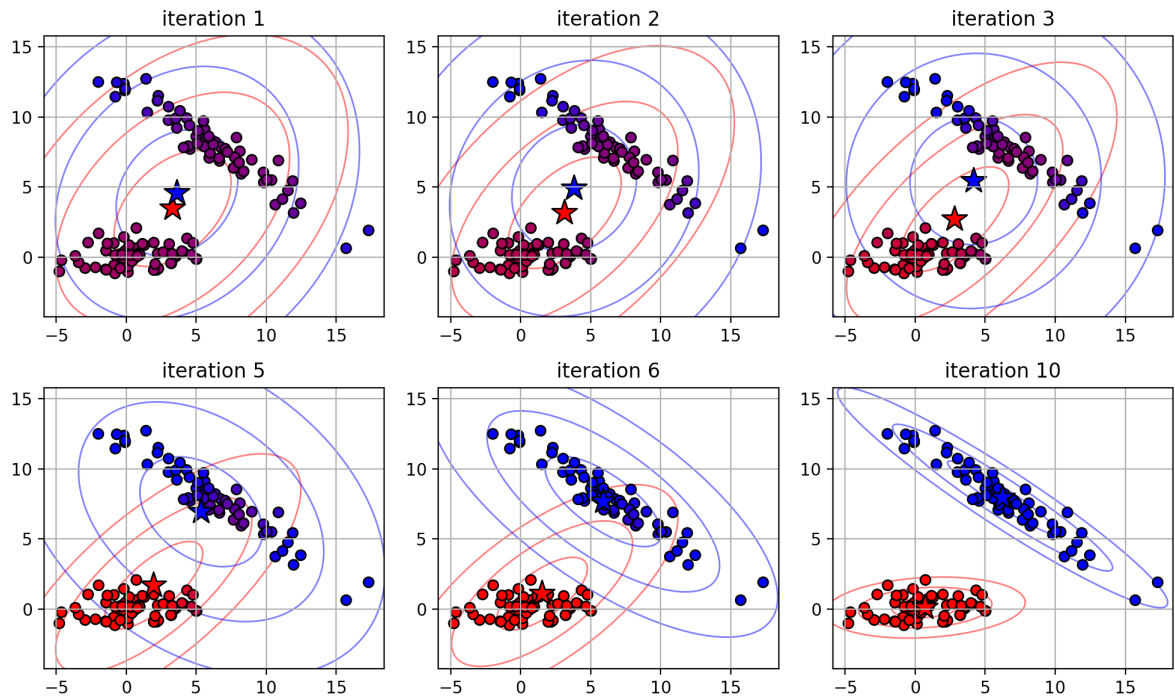
$$\Rightarrow \log p(\mathbf{x}_i, z_i) = \sum_{j=1}^K z_{ij} \log(\pi_j \mathcal{N}(\mathbf{x}_i | \mu_j, \Sigma_j))$$

## EM algorithm for GMMs

- **E-step:** Calculate cluster membership
  - assignment of point  $i$  to cluster  $j$ 
    - $\hat{z}_{ij} = p(z_i = j | \mathbf{x}_i) = \frac{\pi_j \mathcal{N}(\mathbf{x}_i | \mu_j, \Sigma_j)}{\sum_k \pi_k \mathcal{N}(\mathbf{x}_i | \mu_k, \Sigma_k)}$
  - uses "soft" assignment - a data point can have a fractional assignment to different clusters.
- **M-step:** Update each Gaussian cluster (mean, covariance, and weight)
  - uses "soft" weighting
    - "soft" count of points in cluster  $j$ :  $N_j = \sum_{i=1}^N \hat{z}_{ij}$
    - weight:  $\pi_j = N_j / N$
    - mean:  $\mu_j = \frac{1}{N_j} \sum_{i=1}^N \hat{z}_{ij} \mathbf{x}_i$
    - variance:  $\Sigma_j = \frac{1}{N_j} \sum_{i=1}^N \hat{z}_{ij} (\mathbf{x}_i - \mu_j)^2$
- Similar to K-means, except uses "soft" assignments, rather than "hard" assignments.
- EM for GMM Example
  - red points - assigned to cluster 1
  - blue points - assigned to cluster 2
  - purple points - fractionally assigned between clusters 1 and 2

In [32]: efig

Out [32]:



In [33]:

```
# fit a GMM
gmm = mixture.GaussianMixture(n_components=4, random_state=4487, n_init=10)


gmm.fit(X)
Y = gmm.predict(X)

cc = gmm.means_      # the cluster centers
```

In [35]:

```
efig
```

Out [35]:



## Covariance matrix

- The covariance matrix is a  $d \times d$  matrix.

$$\Sigma_j = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

- For high-dimensional data, it can be very large.
  - requires a lot of data to learn effectively.
- Solution:
  - use *diagonal* covariance matrices ( $d$  parameters):

$$\Sigma_j = \text{diag}(\mathbf{a}) = \begin{bmatrix} a_1 & 0 & 0 \\ 0 & a_2 & 0 \\ 0 & 0 & a_3 \end{bmatrix}$$

- Axes of ellipses will be aligned with the axes.
- use *spherical* covariance matrices (1 parameter)

$$\Sigma_j = a\mathbf{I} = \begin{bmatrix} a & 0 & 0 \\ 0 & a & 0 \\ 0 & 0 & a \end{bmatrix}$$

```
In [37]: # full covariance (d*d parameters)
gmfm = mixture.GaussianMixture(n_components=2,
                                covariance_type='full',
                                random_state=4487, n_init=10)

gmfm.fit(X)

# diagonal covariance (d parameters)
gmmd = mixture.GaussianMixture(n_components=2,
                                covariance_type='diag',
                                random_state=4487, n_init=10)

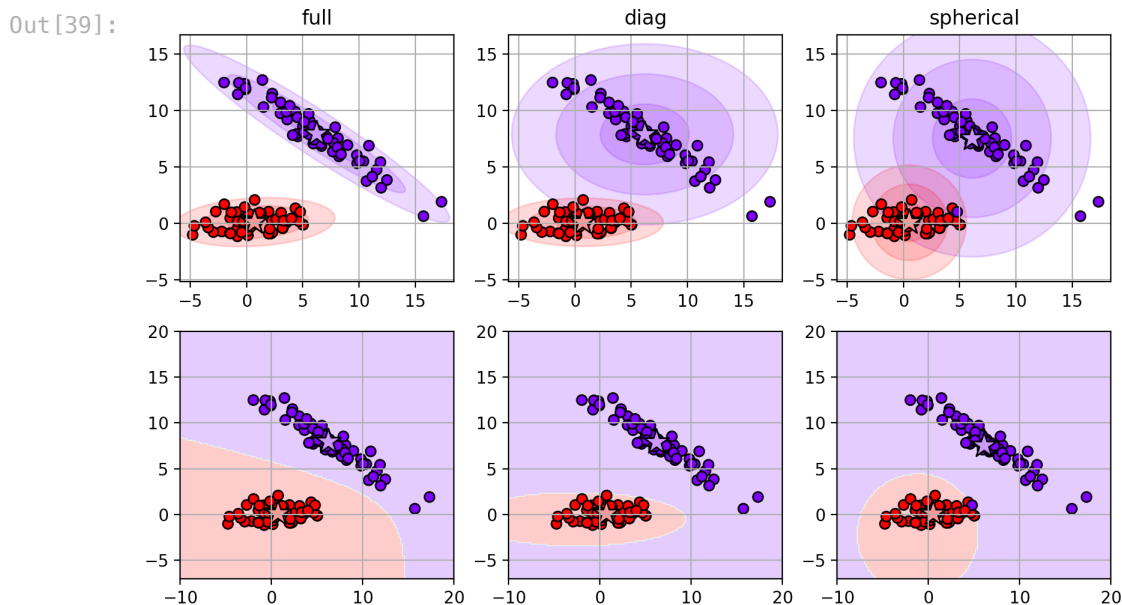
gmmd.fit(X)

# spherical covariance (1 parameter)
gmms = mixture.GaussianMixture(n_components=2,
                                covariance_type='spherical',
                                random_state=4487, n_init=10)

gmms.fit(X)
```

```
Out [37]: GaussianMixture
GaussianMixture(covariance_type='spherical', n_components=2, n_init=10,
                 random_state=4487)
```

```
In [39]: efig
```



## How to select K?

- Clustering results depends on the number of clusters used.
- We don't typically know this information beforehand.

## Bayesian GMM

- A Bayesian version of GMM that automatically selects the value of  $K$ 
  - use a Dirichlet distribution to model the prior probabilities,  $p(\pi)$ .
- *concentration* parameter  $\alpha$  controls the range of  $\pi$  values that are preferred
  - higher values encourage more clusters ( $\pi_j$  are uniform)
  - lower values encourage less clusters (more  $\pi_j = 0$ )
  - Mathematically,  $\alpha$  adds to the number of samples in the cluster,  $\pi_j = \frac{N_j + \alpha}{N + K\alpha}$ 
    - similar to Laplace smoothing.
- (more details in the lecture on graphical models)


```
In [40]: # alpha = concentration parameter
# n_components = the max number of components to consider
dpgmm = mixture.BayesianGaussianMixture(covariance_type='full',
                                         weight_concentration_prior=1,
                                         n_components=5, max_iter=100, random_state=4487)

dpgmm.fit(X)
Y = dpgmm.predict(X)
cl = unique(Y)           # find active clusters
newK = len(cl)           # number of clusters
cc = dpgmm.means_[cl]    # get means
```

- Bayesian GMM automatically selects 2 components from 5
  - for comparison, GMM with 5 clusters looks messy

```
In [42]: dfig
```

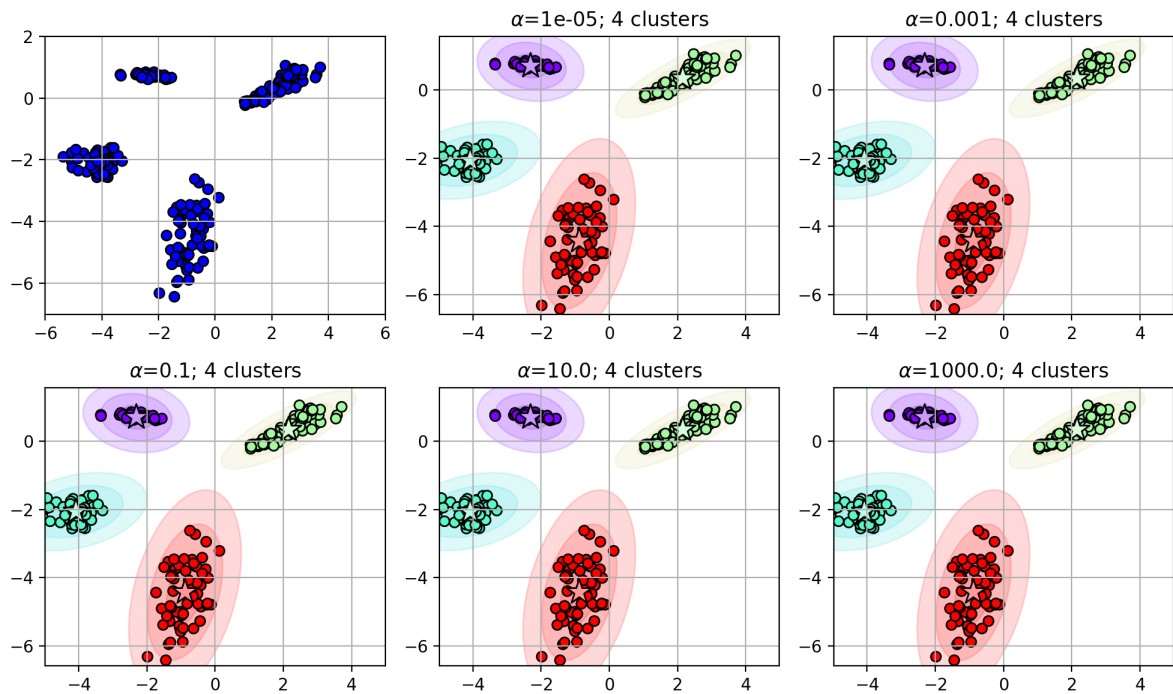
```
Out [42]:
```



- For different concentration parameter  $\alpha$ 
  - larger  $\alpha$  may yield more clusters

```
In [45]: dppfig
```


Out [45]:



- Choice of  $\alpha$  is not that critical
  - same number of clusters for large ranges of  $\alpha$

In [47]: `dpgmmfig`

Out [47]:



In [ ]: