

CS5489 - Machine Learning

Lecture 7b - Unsupervised Learning - Clustering

Prof. Antoni B. Chan

Dept. of Computer Science, City University of Hong Kong

Outline


1. Unsupervised Learning
2. Parametric clustering
 - A. K-means
 - B. Gaussian mixture models (GMMs)
 - C. Bayesian GMMs
3. **Non-parametric clustering and Mean-shift**
4. Spectral clustering

Non-parametric densities

- Suppose we have samples $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$
 - We want to estimate a probability density without assuming a parametric model (e.g., Gaussian)

In [12]: `fig`

Out [12]:




Non-parametric estimation

- Idea: put a small Gaussian at each data point, and sum it up.
 - each point contributes locally to the probability density.
 - $p(\mathbf{x}) = \frac{1}{N} \sum_{i=1}^N \mathcal{N}(\mathbf{x}|\mathbf{x}_i, \sigma^2 \mathbf{I})$
 - σ is the bandwidth of the Gaussian.

In [14]: `fig`

Out [14] :




Kernel density estimator

- This is called a *kernel density estimator*
 - the *kernel* is the small Gaussian.

In [15] :

```
kde = neighbors.KernelDensity(bandwidth=3.0).fit(X)
plot_scores(kde, axbox, 'magma')
plt.scatter(X[:,0], X[:,1], c='b', s=ssize, edgecolors='k');
```



Clustering using KDE

- The modes of the KDE can be considered the cluster centers.
 - A mode is a local maximum of the probability density.
 - The number of clusters is selected automatically according to the data.
- *How to find the cluster centers?*
 - mode: $\mu = \operatorname{argmax} p(\mathbf{x})$
 - select a point, and run gradient ascent on $p(\mathbf{x})$.
 - $\hat{\mu} \leftarrow \hat{\mu} + \eta \frac{d}{d\mathbf{x}} p(\mathbf{x})$
 - using a clever choice of η , we get an algorithm that is guaranteed to converge called the "Mean-shift" algorithm".


Mean-shift algorithm

- **Idea:** iteratively shift towards the largest concentration of points.
 - start from an initial point \mathbf{x} (e.g., one of the data points).
 - repeat until \mathbf{x} is unchanged:
 - 1. find the nearest neighbors to \mathbf{x} within some radius (bandwidth)
 - according to the Gaussian kernels.


- 2. set \mathbf{x} to be the mean of the neighbor points.

In [17]: `efigs[0]`

Out[17]:




iteration 1




In [18]: `efigs[1]`

Out[18]:




iteration 3




In [19]: `efigs[2]`

Out[19]:




iteration 5



In [20]: `efigs[3]`

Out[20]:



Getting the clusters


- Run the mean-shift algorithm for many initial points $\{\mathbf{x}_i\}$.
 - the set of converged points contain the cluster centers.
 - need to remove the duplicate centers.
 - data points that converge to the same center belong to the same cluster.
 - different initializations can run in parallel (`n_jobs`)

```
In [21]: # bin_seeding=True -- coarsely uses data points as initial points
ms = cluster.MeanShift(bandwidth=5, bin_seeding=True, n_jobs=-1)
Y = ms.fit_predict(X)

cc = ms.cluster_centers_ # cluster centers

plot_scores(kde, axbox, 'magma', showcontour=False)
plot_clusters(ms, axbox, X, Y, rbow, rbow2)
```


Out[21]: (2,)



- Cluster partitions
 - assign point based on convergence to same cluster center

```
In [22]: plot_clusters(ms, axbox, X, Y, rbow, rbow2, showregions=True)
```

Out[22]: (2,)




Number of clusters

- Number of clusters is implicitly controlled by the bandwidth (radius of the nearest-neighbors)
 - larger bandwidth creates less clusters
 - focuses on global large groups
 - smaller bandwidth creates more clusters
 - focuses on local groups.

In [24]: `msfig`


Out [24]:



- Cluster partitions: assign points based on convergence to same cluster center.

In [26]: `msfig`

Out [26]:




Non-compact clusters

- K-means, GMM, and Mean-Shift assume that all clusters are compact.
 - i.e., circles or ellipses
- What about clusters of other shapes?
 - e.g., clusters not defined by compact distance to a "center"

In [28]:

`tiefig`


Out [28]:



In [30]:

`circfig`

Out [30]:




Spectral Clustering

- Estimate clusters using pair-wise affinity between points.
- Affinity (similarity) between points
 - kernel function: $k(\mathbf{x}_i, \mathbf{x}_j)$ -- RBF kernel
 - number of nearest neighbors within a radius (bandwidth)

In [32]: `afig`

Out [32]:




Spectral Clustering

- **Idea:** clustering with a graph formulation
 - each data point is a node in a graph
 - edge weight between two nodes is the affinity $k(\mathbf{x}_i, \mathbf{x}_j)$
 - (darker colors indicate stronger weights)

In [34]: `graphfig`

Out [34]:



- **Goal:** cut the graph into clusters such that weights of cut edges is small compared to the total edge weight within each cluster.
 - find "blocks" of high affinity in the affinity matrix.
- Intuitively, consider a "mass-spring" system -- masses connected together with springs.
 - the graph nodes are masses, the edge weights are the spring stiffness.
 - if you hit the masses...
 - the masses that are tightly connected by stiff springs will move together in a low-frequency vibration mode.
- These low-frequency modes are found with the smallest non-zero eigenvectors of the graph Laplacian
- Graph Laplacian: $\mathbf{L} = \mathbf{D} - \mathbf{A}$
 - \mathbf{A} is the adjacency (affinity) matrix.
 - \mathbf{D} is the degree matrix.
 - diagonal matrix with entry $D_{ii} = \sum_j A_{i,j}$
- There are different ways to define the Laplacian, leading to different versions of Spectral Clustering
 - the one in sklearn is called "Normalized Cuts".



In [36]:

```
# spectral clustering
# rbf affinity
sc = cluster.SpectralClustering(n_clusters=4, affinity='rbf',
                                 gamma=1.0, assign_labels='discretize', n_jobs=-1)
Y = sc.fit_predict(X)
```

In [38]:

```
scfig
```

Out [38]:




Sensitivity to gamma

- gamma controls which structures are important
 - small gamma - far away points are still considered similar

```
In [52]: scfig2
```


```
Out[52]:
```



Another Example


```
In [45]: graphfig2
```

```
Out[45]:
```



```
In [47]: scfig
```

```
Out[47]:
```



Clustering Summary

- **Goal:** given set of input vectors $\{\mathbf{x}_i\}_{i=1}^n$, with $\mathbf{x}_i \in \mathbb{R}^d$, group similar x_i together into clusters.
 - estimate a cluster center, which represents the data points in that cluster.
 - predict the cluster for a new data point.

| Name | Cluster Shape | Principle | Advantages | Disadvantages |
|---------|---------------|-------------------------------------|------------------------------|-----------------------------------------------------------------------------|
| K-Means | circular | minimize distance to cluster center | - scalable (MiniBatchKMeans) | - sensitive to initialization; could get bad solutions due to local minima. |

| | | | | |
|------------------------|----------------------|-----------------------------|-----------------------------------------------------------|----------------------------------------------------------------------------------------------------|
| | | | | - need to choose K. |
| Gaussian Mixture Model | elliptical | maximum likelihood | - elliptical cluster shapes. | - sensitive to initialization; could get bad solutions due to local minima. - need to choose K. |
| Bayesian GMM | elliptical | maximum marginal likelihood | - automatically selects K via concentration parameter. | - can be slow. - sensitive to initialization; could get bad solutions due to local minima. |
| Mean-Shift | concentrated compact | move towards local mean | - automatically selects K via bandwidth parameter. | - can be slow. |
| Spectral clustering | irregular shapes | graph-based | - can handle clusters of any shape, as long as connected. | - need to choose K. - cannot assign novel points to a cluster. - can be slow (kernel matrix) |

Other Things


- *Feature normalization*
 - feature normalization is typically required clustering.
 - e.g., algorithms based on Euclidean distance (Kmeans, Mean-Shift, Spectral Clustering)

Example

- scaling down the x_1 feature makes its differences less important, compared to x_2 .

In [49]: `efig`

Out[49]:



In []: