

CS5489- Machine Learning

Lecture 9a - Convolutional Neural Networks (CNNs)

Prof. Antoni B. Chan

Dept. of Computer Science, City University of Hong Kong


Outline

- Convolutional neural network (CNN)
- Regularization


Signals

- So far we have assumed the input \mathbf{x} is a vector
 - or have turned 2D images into vectors.
- *What if the input has more structure?*
- For example:
 - **1-D signal** (time)

mono audio (1 feature)



stereo audio (2 features)




- **2-D signal** (space)

grayscale image (1 feature)



color image (3 features)




hyperspectral image (300 features)



- **3-D signal** (space+time, volume)

color video (3 features)



3D CT scan (1 feature)



Assumed Properties of Signals



- *Locality*

- at low-level, features from 1 region are independent (do not depend on) features from a far-away region.




- *Translation*

- the same features can appear anywhere in the signal.

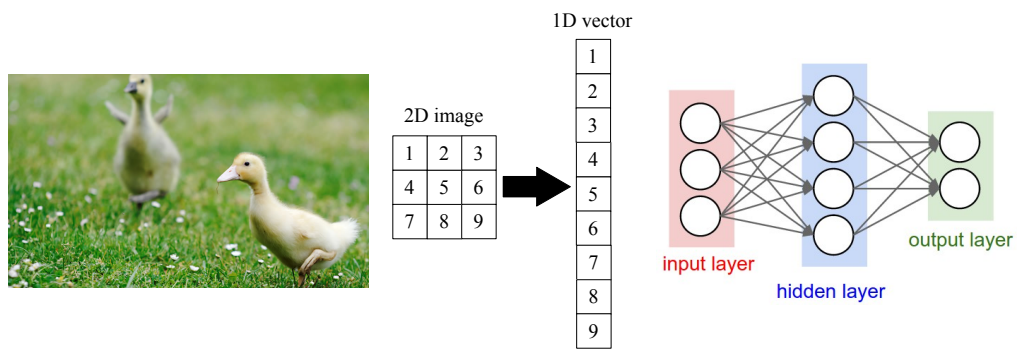


Using the standard MLP layer...

- Each feature z_i is computed from all the inputs, but we only want local features (locality).
- The pattern could appear anywhere, but weights are trained for each location z_i separately (translation).



- For image input, we transform the image into a vector, which is the input into the MLP.



- **Problem:** This ignores the spatial relationship between pixels in the image.

- Images contain local structures

- **How to model the local structure of the signal?**

- local features, feature translation


- **Answer:** Use a local feature extractor in the signal.

Convolution

- Consider 1D signal in *discrete* time: $x(t), t \in \mathbb{Z}$
- Define the filter $w(t)$
 - "flipped" filter: $\tilde{w}(t) = w(-t)$

In [9]: xfig


Out [9]:



- Convolution is a filtering operation
 - $s(t) = x * w = \sum_a x(a)w(t - a)$
 - "*" is the symbol for convolution
- It's related to cross-correlation with the "flipped" filter.
 - $s(t) = \sum_a x(a)\tilde{w}(a - t)$
 - for a given t :
 1. shift \tilde{w} by t
 2. multiply shifted \tilde{w} with x
 3. sum to get $s(t)$
- $s(t) = \sum_a x(a)\tilde{w}(a - t)$
 1. shift \tilde{w} by t
 2. multiply shifted \tilde{w} with x
 3. sum to get $s(t)$

In [11]: sfig

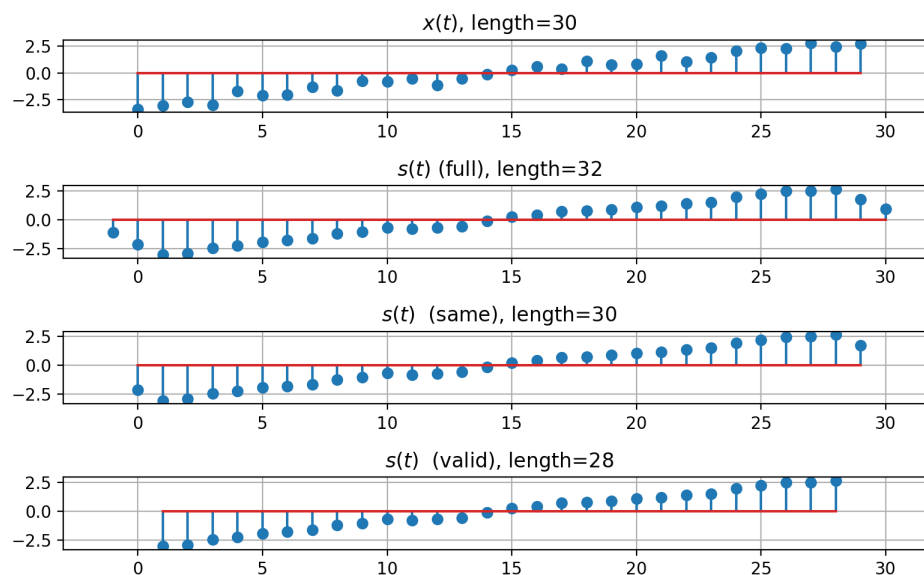
Out [11]:



- Boundary conditions
 - It is assumed that the rest of the signal is all 0.
- The convolution result near the ends of the signal uses these "artificial" zeros.
 - the filtered signal is longer than the original signal
 - length increases by $P - 1$, where P is the non-zero extent of the filter.
- Three ways to handle this:
 1. use everything with *non-zero* response ('full' mode)
 2. keep the response the *same* length as the signal ('same' mode)
 3. keep only responses where the *entire* filter is on the signal ('valid' mode)

In [13]: `cfig`

Out [13]:




2D Convolution

- Straightforward to extend to multiple dimensions
- 2D discrete convolution

$$s(n, m) = x * w = \sum_a \sum_b x(a, b) w(n - a, m - b)$$


In [15]: `sfig`

Out [15]:



Two Interpretations of Convolution

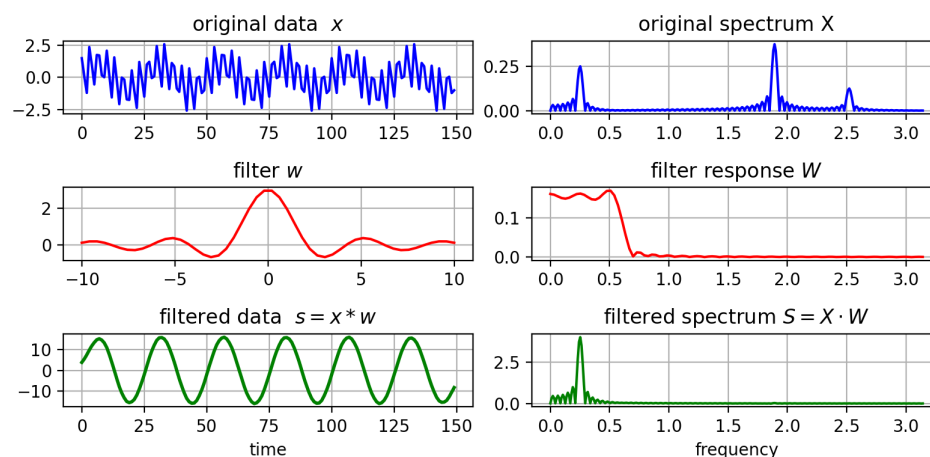
- Interpretation 1:
 - w is a filter on the frequency spectrum of signal x .
 - Example filters: *low-pass*, *high-pass*, *band-pass*, *moving-average*



- Analysis is in the frequency domain:
 - **Time domain** $x(t) \iff$ **Frequency domain** $X(\omega)$
- Discrete-time Fourier transform (DTFT)
 - represent data as sum of complex exponentials basis functions with different frequencies.
 - $e^{-i\omega t} = \cos(\omega t) - i \sin(\omega t)$
 - Imaginary number: $i^2 = -1$
 - $X(\omega) = \sum_t x(t)e^{-i\omega t} \iff x(t) = \frac{1}{2\pi} \int X(\omega)e^{i\omega t} d\omega$
- **Key results:**
 - convolution in the time domain is equivalent to multiplication in the frequency domain:
 - $s(t) = x(t) * w(t) \iff S(\omega) = X(\omega)W(\omega)$
 - we can design and analyze filters in the frequency domain, and then obtain their time-domain representation.

In [17]: `ffig`


Out [17]:



- Relationships between time and frequency spectrum
 - short-duration (high-frequency) signal \iff large frequency spectrum
 - long-duration (low-frequency) signal \iff small frequency spectrum
- For filters...
 - short (small) filter only captures short-range correlations (high frequencies).
 - long (large) filter can capture long-range correlations (low frequencies).


- Interpretation 2:
 - \tilde{w} is a pattern (template); try to find this pattern.
 - the maximum correlation occurs when pattern \tilde{w} matches the local x .
 - for a fixed energy $\|x\|^2 = 1$, $x = \frac{\tilde{w}}{\|\tilde{w}\|}$ has the maximum correlation with (response to) \tilde{w} .

In [20]: `pfig`



Convolution as a layer

- output layer is the convolution of input with filter (kernel) \mathbf{w} .
 - $\mathbf{z} = \mathbf{x} * \mathbf{w}$.
- filter \mathbf{w} acts locally on input \mathbf{x} .
 - \mathbf{w} also called a **kernel**.



- Equivalent to a linear transformation (layer) where \mathbf{A} has a particular form.
 - For example, if $\mathbf{w} = [w_1, w_2, w_3]$ and using "same" mode,

$$\mathbf{z} = \mathbf{x} * \mathbf{w}$$

$$= \mathbf{A}^T \mathbf{x} = \begin{bmatrix} w_2 & w_3 & 0 & 0 & 0 & 0 & \dots \\ w_1 & w_2 & w_3 & 0 & 0 & 0 & \dots \\ 0 & w_1 & w_2 & w_3 & 0 & 0 & \dots \\ 0 & 0 & w_1 & w_2 & w_3 & 0 & \dots \\ \dots & & & & & & \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_N \end{bmatrix}$$


- Note: \mathbf{A} has size $|\mathbf{x}| |\mathbf{z}|$, but only $|\mathbf{w}|$ parameters.

Translation equivariance

- Shifting x also shifts the response s .
- if $s = x * w$,
 - then $s(t - a) = x(t - a) * w(t)$
- We can find the pattern everywhere in x using the same filter.


In [22]: `cfig`

Out [22]:




Convolutional Neural Network (CNN)

- series of convolutional layers, sub-sampling layers, and MLP classifier.
 - convolutional and subsampling layers extract image features.
 - MLP uses extracted features for classification.




2D Convolution


- Use the spatial structure of the image
- 2D convolution filter
 - the weights \mathbf{W} form a 2D filter template
 - filter response: $h = f(\sum_{x,y} W_{x,y} P_{x,y})$
 - \mathbf{P} is an image patch with the same size as \mathbf{W} .
- Convolution feature map
 - pass a sliding window over the image, and apply filter to get a *feature map*.



- Convolution modes
 - "valid" mode - only compute feature where convolution filter has valid image values.
 - size of feature map is reduced.




- Convolution modes
 - **"same"** mode - zero-pad the border of the image
 - feature map is the same size as the input image.



- Usually "same" is better since it looks at structures around border. - position information is implicitly encoded in the CNN features based on the zero-padding border.


2D Convolutional layer

- **Input:** HxW image with C channels
 - For example, in the first layer, $C=3$ for RGB channels.
 - defines a 3D volume: $C \times H \times W$ (or $H \times W \times C$)
- **Features:** apply F convolution filters to get F feature maps.
 - Each feature map uses a 3D convolution filter ($C \times K \times K$) on the input
 - K is the spatial extent of the filter; total FCKK parameters
- **Activation:**
 - an activation function can be applied before output
- **Output:** a feature map with F channels
 - defines a 3D volume: $F \times H' \times W'$
 - H' and W' depend on various factors.




Combining Convolutional Layers

- Concatenate several convolutional layers.
 - From layer to layer
 - spatial resolution decreases
 - number of feature maps increases
 - Can extract high-level features in the final layers
- Feature map representation:




Spatial sub-sampling

- reduce the feature map size by subsampling feature maps between convolutional layers
 - *stride* for convolution filter - step size when moving the windows across the image.




- Spatial sub-sampling
 - *max-pooling layer* - use the maximum over a pooling window
 - gathers features together, summarizes features in local region.



- introduces **translation invariance** - it doesn't matter where the maximal feature is located locally, it is passed to the next layer. - increases robustness to small changes in configuration of features

Receptive field size

- Stacking convolutional layers increases the effective size of the pattern filter
 - called **receptive field** - what pixels in the input affect a particular node.
 - larger receptive fields can see larger patterns.
 - Example: 2 convolutional layers
 - $|\mathbf{w}_1| = 5, |\mathbf{w}_2| = 3$, receptive field size = $|\mathbf{w}_1| + |\mathbf{w}_2| - 1 = 7$




Advantages of Convolution Layers

- The convolutional filters extract the same features throughout the image.
 - Useful because the object can appear in different locations of the image (global translation equivariance).
- Pooling makes it robust to changes in feature configuration (local translation invariance).
- The number of parameters is small compared to Dense (Fully-connected) layer
 - Example: input is $C \times H \times W$, and output is $F \times H \times W$

- Number of MLP parameters: $(CHW+1) \times (FWH)$
- Number of CNN parameters: $E \times (CKK+1)$

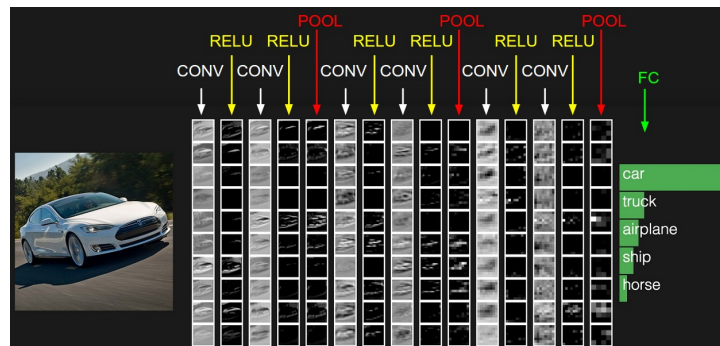
Fully-connected layers (MLP)

- after several convolutional layers, input the feature map into an MLP to get the final classification.
- also called "fully-connected" (FC) layers.



Example: Object classification CNN


- Each layer shows its feature maps for the example image.
 - early layers extract *low-level* (visual) features
 - e.g., corners, edges
 - middle layers extract *mid-level* (part) features
 - e.g., object parts
 - later layers extract *high-level* (semantic) features.
 - e.g., object



- The number of feature channels increases with each layer
 - combining low-level parts together to get more higher-level parts
 - e.g., {edges, corners} → {wheel, door, window}
 - trading off spatial resolution for semantic specificity
 - e.g., 512x512x3 RGB image → 8x8x512 semantic features
- the spatial resolution decreases with each layer
 - increase the window size (receptive field) on the object
 - high-level semantic correspond to large regions.

CNN on MNIST

```
In [23]: # Example images
plt.figure(figsize=(8,4))
show_imgs(training[0:50])
```



4D tensor format

- There are two common formats for the 4-D tensor:
 - "NCHW" - batch, channel, height, width - called 'channels_first'
 - "NHWC" - batch, height, width, channel - called 'channels_last'
- NHWC is required for CPU version of Tensorflow 2.

```
In [24]: # use keras backend (K) to force channels-last ordering (for CPU compatability)
K.set_image_data_format('channels_last')
```

Example on MNIST

- Pre-processing
 - scale to [0,1]
 - 4-D tensor: (sample, height, width, channel)
 - channel = 1 (grayscale)
 - create training/validation sets

```
In [25]: # scale to 0-1
trainI = (training.reshape((6000,28,28,1)) / 255.0)
testI = (testing.reshape((10000,28,28,1)) / 255.0)
print(trainI.shape)
print(testI.shape)

(6000, 28, 28, 1)
(10000, 28, 28, 1)
```

- Generate fixed training and validation sets

```
In [26]: # generate fixed validation set of 10% of the training set
vtrainI, validI, vtrainYb, validYb = \
    model_selection.train_test_split(trainI, trainYb,
    train_size=0.9, test_size=0.1, random_state=4487)

validsetI = (validI, validYb)
```

Shallow CNN Architecture

- 1 Convolution layer
 - 5x5x1 kernel, 10 features, *stride* = 2 (step-size between sliding windows)
 - No pooling here since the image input is small (28x28)
 - Input: 28x28x1 (grayscale image) -> Output: 14x14x10
- 1 fully-connected layer (MLP), 50 nodes
 - Input: 14x14x10=1960 -> Output: 50
- Classification output node

```
In [28]: # initialize random seed
K.clear_session(); random.seed(4487); tf.random.set_seed(4487)

# build the network
nn = Sequential()
nn.add(Conv2D(10, (5,5), strides=(2,2), # channel, kernel size, stride
             activation='relu', padding='same', # activation, convolution padding
             input_shape=(28,28,1)))
nn.add(Flatten()) # flatten the feature map into a vector to apply Dense layers
nn.add(Dense(units=50, activation='relu'))
nn.add(Dense(units=10, activation='softmax'))

# compile and fit the network
nn.compile(loss=keras.losses.categorical_crossentropy,
           optimizer=keras.optimizers.SGD(learning_rate=0.02, momentum=0.9, nesterov=True),
           metrics=['accuracy'])
history = nn.fit(vtrainI, vtrainYb, epochs=100, batch_size=50,
               callbacks=callbacks_list,
               validation_data=validsetI, verbose=False)
```

Metal device set to: Apple M1 Max

```
2023-01-22 21:21:06.335346: I tensorflow/core/common_runtime/pluggable_device/plug
gable_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, de
faulting to 0. Your kernel may not have been built with NUMA support.
2023-01-22 21:21:06.335478: I tensorflow/core/common_runtime/pluggable_device/plug
gable_device_factory.cc:271] Created TensorFlow device (/job:localhost/replica:0/t
ask:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (device: 0, name:
METAL, pci bus id: <undefined>)
2023-01-22 21:21:06.420007: W tensorflow/core/platform/profile_utils/cpu_utils.cc:
128] Failed to get CPU frequency: 0 Hz
2023-01-22 21:21:06.515089: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-22 21:21:07.207752: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

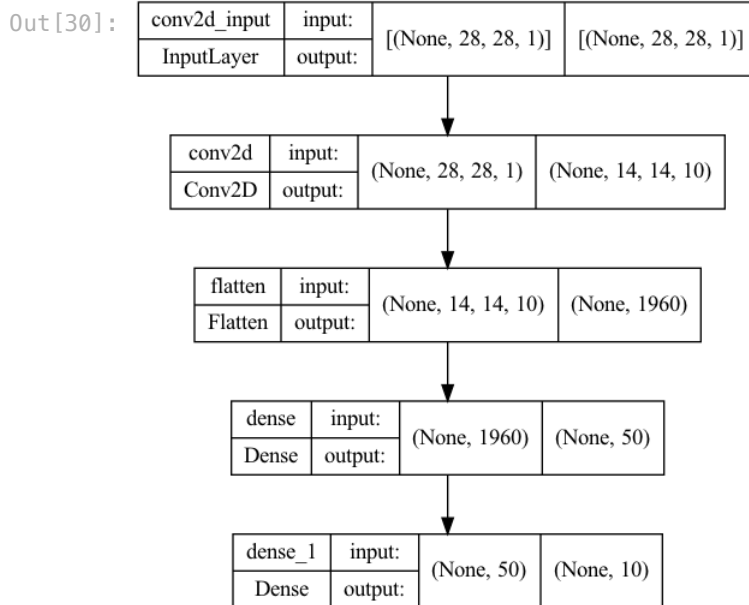
Epoch 16: early stopping

```
In [29]: nn.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 14, 14, 10)	260
flatten (Flatten)	(None, 1960)	0
dense (Dense)	(None, 50)	98050
dense_1 (Dense)	(None, 10)	510
Total params: 98,820		
Trainable params: 98,820		
Non-trainable params: 0		

```
In [30]: # visualize the network
tf.keras.utils.plot_model(nn, to_file='tmp_model.png', show_shapes=True)
```




- test results
 - for comparison, the best MLP from Lecture 8 was 0.9499 accuracy
 - 3 layer MLP: (500, 500, 10) with 648,010 parameters

```
In [31]: plot_history(history)

predY = argmax(nn.predict(testI, verbose=False), axis=-1)
acc = metrics.accuracy_score(testY, predY)
print("test accuracy:", acc)
```

2023-01-22 21:21:17.814170: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.


test accuracy: 0.9534



- Visualize the convolutional filters
 - filters are looking for local stroke features
 - corners, edges, lines

```
In [32]: W = nn.get_layer(index=0).get_weights()[0]
print(W.shape)
flist = [squeeze(W[:, :, :, c]) for c in range(10)]
show_imgs(flist)
```

(5, 5, 1, 10)



Deep CNN Architecture

- 3 Convolutional layers
 - 5x5x1 kernel, stride 2, 10 features (output 14x14x10)
 - 5x5x10 kernel, stride 2, 40 features (output 7x7x40)
 - 5x5x40 kernel, stride 1, 80 features (output 7x7x80)
 - set stride as 1 to avoid reducing the feature map too much)
- 1 fully-connected layer (7x7x80=3920 -> 50)
- Classification output node

```
In [33]: K.clear_session()
random.seed(4487); tf.random.set_seed(4487) # initialize random seed
# build the network
nn = Sequential()
nn.add(Conv2D(10, (5,5), strides=(2,2), activation='relu',
              padding='same', input_shape=(28,28,1)))
nn.add(Conv2D(40, (5,5), strides=(2,2), activation='relu', padding='same'))
nn.add(Conv2D(80, (5,5), strides=(1,1), activation='relu', padding='same'))
nn.add(Flatten())
nn.add(Dense(units=50, activation='relu'))
nn.add(Dense(units=10, activation='softmax'))

# compile and fit the network
nn.compile(loss=keras.losses.categorical_crossentropy,
           optimizer=keras.optimizers.SGD(learning_rate=0.02, momentum=0.9, nesterov=True),
           metrics=['accuracy'])
history = nn.fit(vtrainI, vtrainYb, epochs=100, batch_size=50,
                callbacks=callbacks_list,
                validation_data=validsetI, verbose=False)
```

```
2023-01-22 21:21:18.731616: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-22 21:21:19.505703: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

Epoch 10: early stopping

```
In [34]: nn.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 14, 14, 10)	260
conv2d_1 (Conv2D)	(None, 7, 7, 40)	10040
conv2d_2 (Conv2D)	(None, 7, 7, 80)	80080
flatten (Flatten)	(None, 3920)	0
dense (Dense)	(None, 50)	196050
dense_1 (Dense)	(None, 10)	510
=====		
Total params: 286,940		
Trainable params: 286,940		
Non-trainable params: 0		
=====		


```
In [35]: plot_history(history)

predY = argmax(nn.predict(testI, verbose=False), axis=-1)
```

```
acc = metrics.accuracy_score(testY, predY)
```

```
2023-01-22 21:21:26.269758: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

test accuracy: 0.9564



Summary

- Convolution operation
 - looks at the local structure of the signal (1D, 2D, 3D, etc).
 - two interpretations: filtering, pattern matching
- Convolutional neural network (CNN)
 - convolutional layer - use convolution instead of dense connections
 - learns to extract image features, and learns classifier.