

CS5489 - Machine Learning

Lecture 11b - Auto-Encoders and Deep Generative Models

Prof. Antoni B. Chan

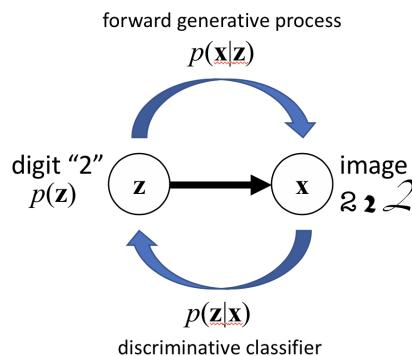
Dept. of Computer Science, City University of Hong Kong

Outline

- Unsupervised models
 - Autoencoders
 - Denoising Autoencoders
- Deep generative models
 - Variational Autoencoders
 - GANs
 - Diffusion Models

Deep Generative Models

- Why are they necessary?
 - typical mapping for discriminative classifiers:
 - many inputs to one output (e.g., images of "2" to class label "2")
 - typical mapping for synthesis:
 - one input to many outputs (e.g., label 2 to images of "2")



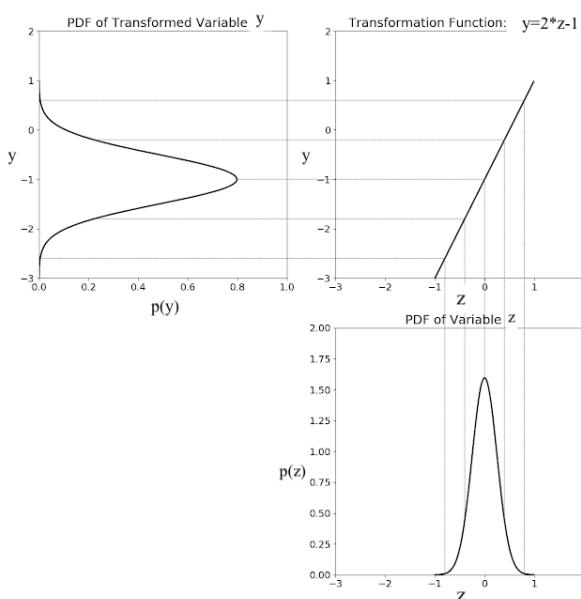
- typical NN can only predict one output
 - if multiple outputs are possible, it predicts the mean to minimize the loss
 - e.g. suppose \mathbf{x} is associated with K outputs $\{y_j\}$.
 - $\min \sum_{j=1}^K \|y_j - f(\mathbf{x})\|^2 \Rightarrow f(\mathbf{x}) = \frac{1}{K} \sum_j \mathbf{y}_j$
- How to extend NN to be generative models?

- **simple idea:** use NN to predict the parameters of a distribution
 - e.g., $y \sim \mathcal{N}(\mu(\mathbf{x}), \sigma(\mathbf{x})^2)$
 - sample from the Gaussian to obtain values of y for input \mathbf{x} .
- Problems:
 - How to do back propagation through the random variable?
 - How to take the derivative with respect to (μ, σ^2) ?

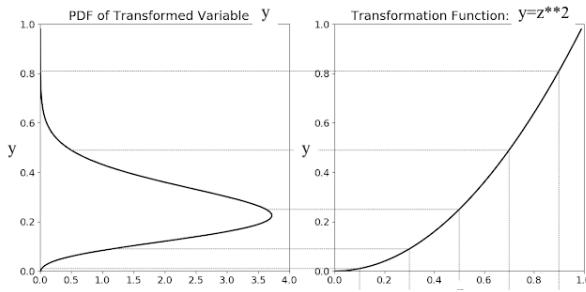
Reparameterization Trick

- **Solution:**
 - rewrite y as a function of another random variable with known parameters
 - e.g., $y = \mu(\mathbf{x}) + \sigma(\mathbf{x})z$, $z \sim \mathcal{N}(0, 1)$
 - decompose y into parameters and "randomness"
 - can now take derivatives w.r.t (μ, σ^2) since z is not a function of (μ, σ^2) .
 - z is an extra input (random samples)
- Example:

$$y = 2 * z - 1$$



- In general, we rewrite the distribution as
 - $y \sim p(y|x) \Rightarrow y = f(z, x), z \sim p(z)$
 - x are the inputs, z is the source of randomness.
 - f transforms samples of z to samples of y
 - modeled as a NN and learned.
- Example: $y = z^2$



What is the actual pdf?

- transform of a random variable: $y = f(z)$, $z \sim p_z(z)$
 - assume f is monotonically increasing, then $z = f^{-1}(y)$.
 - mapping density from z to y :

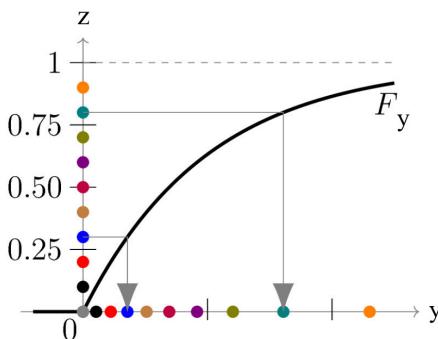
$$|p_y(y)dy| = |p_z(z)dz|$$

$$\Rightarrow p_y(y) = \left| \frac{dz}{dy} \right| p_z(z) = \left| \frac{df^{-1}(y)}{dy} \right| p_z(f^{-1}(y))$$

- hard to evaluate...
 - ... learning f is indirectly learning the density $p_y(y)$.

Closer look

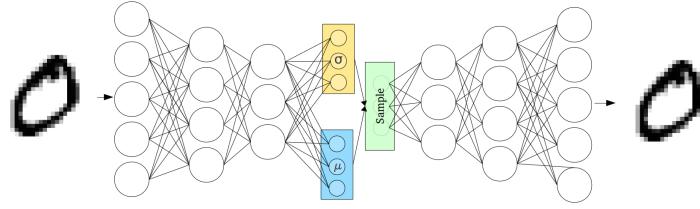
- Suppose $f^{-1}(y) = F_y(y) = \int_{-\infty}^y p_y(y)dy$ is the CDF of y
 - then $\frac{df^{-1}(y)}{dy} = \frac{dF_y(y)}{dy} = p_y(y)$
- Assume $z \sim U(0, 1)$ is a uniform distribution over $[0,1]$
 - then $p_z(f^{-1}(y)) = 1$.
- Finally, substituting we obtain
 - $\left| \frac{df^{-1}(y)}{dy} \right| p_z(f^{-1}(y)) = p_y(y)$
- Thus, when $y = f(z)$ and z is uniform(0,1),
 - learning f is equivalent to learning the inverse CDF of y .
- Example:
 - z samples are mapped to x



Variational AutoEncoder (VAE)

- The standard autoencoder can have difficulty encoding/decoding new images
 - the decoder never sees (encoded) latent vectors outside of the training set
- VAE fixes this by introducing noise in the latent vectors

- the noise lets the decoder network see slightly different latent vectors for each training image.
- improves the ability to interpolate between training samples

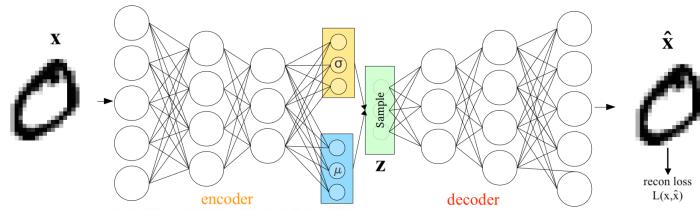


VAE model

- consider the probabilistic model
 - let \mathbf{z} be the latent variable, and \mathbf{x} the observation.
 - assume prior $p(\mathbf{z}) \sim \mathcal{N}(0, \mathbf{I})$
 - let $p(\mathbf{x}|\mathbf{z})$ be the likelihood function
 - i.e., $-\log p(\mathbf{x}|\mathbf{z})$ is the data loss function
- given a \mathbf{x}_i we want to infer \mathbf{z}_i
 - let $q(\mathbf{z}_i|\mathbf{x}_i) \sim \mathcal{N}(\mu(\mathbf{x}_i), \text{diag}(\sigma(\mathbf{x}_i)^2))$ be the posterior of \mathbf{z}_i given observation \mathbf{x}_i .
 - $\mu(\cdot), \sigma(\cdot)$ are the outputs of a NN.
 - samples from posterior are generated via reparameterization trick
 - $\mathbf{z}_{il} = \mu(\mathbf{x}_i) + \sigma(\mathbf{x}_i) \circ \epsilon_{il}$
 - where $\epsilon_{il} \sim \mathcal{N}(0, \mathbf{I})$

VAE learning

- Ideally we would want to learn the parameters by maximizing the marginal log-likelihood.
 - $\log p(\mathbf{x}) = \log \int p(\mathbf{x}|\mathbf{z})p(\mathbf{z})d\mathbf{z}$.
- Intractable, so we maximize a lower-bound using a **variational approximation**
 - ...more details about this approximation in the next lecture on Graphical Models.
- After some derivation, we obtain the VAE loss function
 - $L(\mathbf{x}_i) = -\frac{1}{L} \sum_{l=1}^L \underbrace{\log p(\mathbf{x}_i|\mathbf{z}_{il})}_{\text{reconstruction loss}} + \underbrace{KL(q(\mathbf{z}|\mathbf{x}_i), p(\mathbf{z}_i))}_{\text{regularizer}}$
 - where $\mathbf{z}_{il} = \underbrace{\mu(\mathbf{x}_i) + \sigma(\mathbf{x}_i) \circ \epsilon_{il}}_{\text{encoder}}, \epsilon_{il} \sim \mathcal{N}(0, \mathbf{I})$
 - decoder is the output of a NN fed into the reconstruction loss
 - (e.g., MSE, binary cross-entropy)
 - regularizer keeps the posterior close to the prior.



VAE Example

In [14]:

```
#some settings
K.clear_session()
random.seed(4487); tf.random.set_seed(4487)
original_dim = 784
input_shape = (original_dim, )
intermediate_dim = 512
batch_size = 128
latent_dim = 10
epochs = 50
```

- Build the encoder
- Map the input into the mean and log(sigma) of the Gaussian distribution
 - the mean is the encoded vector

In [15]:

```
# encoder mapping to distribution (mean and log_sigma)
x = Input(shape=(original_dim,), name='input')
h = Dense(intermediate_dim, activation='relu', name='enc_hidden')(x)

# the mean and log-sigma
z_mean = Dense(latent_dim, name='mu')(h)
z_log_sigma = Dense(latent_dim, name='log_sigma')(h)

# encoder, from inputs to latent space
encoder = Model(x, z_mean, name='encoder')
```

- Use the mean and log(sigma) to sample a latent variable z
 - the `Lambda` layer applies a function to an input, and outputs it.
 - use keras backend `K` for all math operations, so backprop can be used

In [16]:

```
# sampling function - draw Gaussian random noise
epsilon_std = 0.01 # fixed stddev
def sampling(args):
    z_mean, z_log_sigma = args
    epsilon = K.random_normal(shape=K.shape(z_mean), mean=0., stddev=epsilon_std)
    return z_mean + K.exp(z_log_sigma) * epsilon

# layer that samples according to mean and sigma
z = Lambda(sampling, name='sampler')([z_mean, z_log_sigma])
```

- Decode the latent variable z
- Construct the whole VAE

In [17]:

```
# create the layers since we need to use it later
decoder_h = Dense(intermediate_dim, activation='relu', name='dec_hidden')
decoder_mean = Dense(original_dim, activation='sigmoid', name='output')

# connect the latent variable and hidden states
h_decoded = decoder_h(z)
x_decoded_mean = decoder_mean(h_decoded)

# end-to-end variational autoencoder
vae = Model(x, x_decoded_mean)
```

- Construct the generator (decoder)
 - attach another input to the decoder layers, and connect them

In [18]:

```
# generator, from latent space to reconstructed inputs
decoder_input = Input(shape=(latent_dim,)) # make an input and attach it to the hidden stat
_h_decoded = decoder_h(decoder_input) # and other layers
_x_decoded_mean = decoder_mean(_h_decoded)

# the generator model
generator = Model(decoder_input, _x_decoded_mean)
```

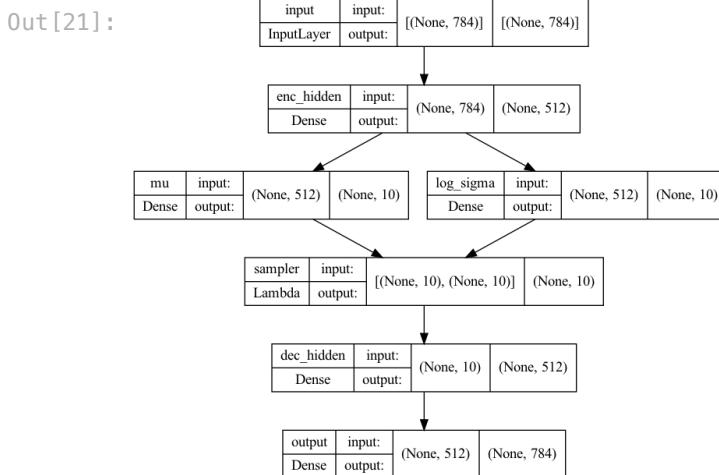
- define the VAE loss function

```
In [19]: # define the VAE loss
def vae_loss(x, x_decoded_mean):
    # cross-entropy loss (reconstruction loss)
    xent_loss = keras.losses.binary_crossentropy(x, x_decoded_mean)
    # KL divergence loss (regularizer)
    kl_loss = -0.5 * K.mean(1 + z_log_sigma - K.square(z_mean) - K.exp(z_log_sigma), axis=-1)
    return xent_loss + kl_loss

# compile the model for optimization
vae.compile(optimizer='rmsprop', loss=vae_loss)
```

- The final VAE model
 - note that there are two layers going into "sampler"

```
In [21]: tmp=tf.keras.utils.plot_model(vae, show_shapes=True); tmp.width=500; tmp
```



- The encoder and decoder

```
In [22]: encoder.summary()
```

Model: "encoder"

Layer (type)	Output Shape	Param #
<hr/>		
input (InputLayer)	[(None, 784)]	0
enc_hidden (Dense)	(None, 512)	401920
mu (Dense)	(None, 10)	5130
<hr/>		
Total params: 407,050		
Trainable params: 407,050		
Non-trainable params: 0		

```
In [23]: generator.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #
<hr/>		
input_1 (InputLayer)	[(None, 10)]	0
dec_hidden (Dense)	(None, 512)	5632
output (Dense)	(None, 784)	402192
<hr/>		

```
Total params: 407,824
Trainable params: 407,824
Non-trainable params: 0
```

- Fit the model

```
In [24]: history = vae.fit(vtrainXraw, vtrainXraw,
    shuffle=True,
    epochs=epochs,
    batch_size=batch_size,
    validation_data=(validXraw, validXraw),
    callbacks=[TensorBoard(log_dir='./logs/vae')],
    verbose=False
)

Metal device set to: Apple M1 Max

2023-01-25 20:22:47.299282: I tensorflow/core/common_runtime/pluggable_device/plug
gible_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, de
faulting to 0. Your kernel may not have been built with NUMA support.
2023-01-25 20:22:47.299421: I tensorflow/core/common_runtime/pluggable_device/plug
gible_device_factory.cc:271] Created TensorFlow device (/job:localhost/replica:0/t
ask:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (device: 0, name:
METAL, pci bus id: <undefined>)
2023-01-25 20:22:47.307580: W tensorflow/core/platform/profile_utils/cpu_utils.cc:
128] Failed to get CPU frequency: 0 Hz
2023-01-25 20:22:47.307688: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:22:47.321445: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:22:47.375246: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:22:47.630258: I tensorflow/core/profiler/lib/profiler_session.cc:11
0] Profiler session initializing.
2023-01-25 20:22:47.630272: I tensorflow/core/profiler/lib/profiler_session.cc:12
5] Profiler session started.
2023-01-25 20:22:47.641492: I tensorflow/core/profiler/lib/profiler_session.cc:67]
Profiler session collecting data.
2023-01-25 20:22:47.641904: I tensorflow/core/profiler/lib/profiler_session.cc:14
3] Profiler session tear down.
2023-01-25 20:22:47.642590: I tensorflow/core/profiler/rpc/client/save_profile.cc:
136] Creating directory: ./logs/vae/plugins/profile/2023_01_25_20_22_47

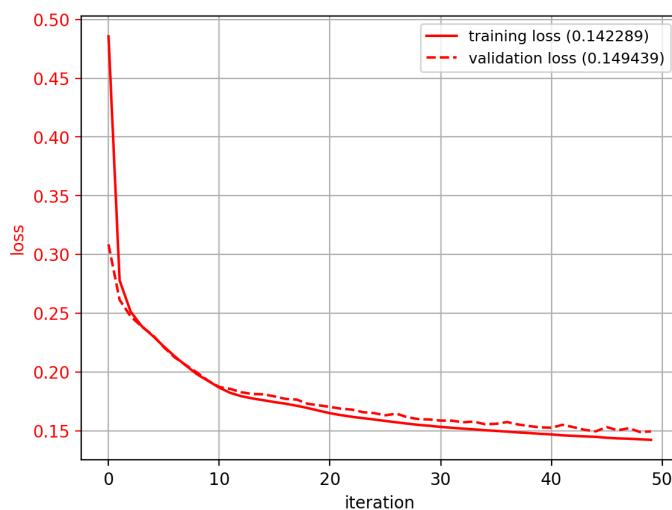
2023-01-25 20:22:47.642864: I tensorflow/core/profiler/rpc/client/save_profile.cc:
142] Dumped gzipped tool data for trace.json.gz to ./logs/vae/plugins/profile/2023
_01_25_20_22_47/Kili.trace.json.gz
2023-01-25 20:22:47.643809: I tensorflow/core/profiler/rpc/client/save_profile.cc:
136] Creating directory: ./logs/vae/plugins/profile/2023_01_25_20_22_47

2023-01-25 20:22:47.643965: I tensorflow/core/profiler/rpc/client/save_profile.cc:
142] Dumped gzipped tool data for memory_profile.json.gz to ./logs/vae/plugins/pro
file/2023_01_25_20_22_47/Kili.memory_profile.json.gz
2023-01-25 20:22:47.644336: I tensorflow/core/profiler/rpc/client/capture_profile.
cc:251] Creating directory: ./logs/vae/plugins/profile/2023_01_25_20_22_47
Dumped tool data for xplane.pb to ./logs/vae/plugins/profile/2023_01_25_20_22_47/K
ili.xplane.pb
Dumped tool data for overview_page.pb to ./logs/vae/plugins/profile/2023_01_25_20_
22_47/Kili.overview_page.pb
Dumped tool data for input_pipeline.pb to ./logs/vae/plugins/profile/2023_01_25_20_
22_47/Kili.input_pipeline.pb
Dumped tool data for tensorflow_stats.pb to ./logs/vae/plugins/profile/2023_01_25_
20_22_47/Kili.tensorflow_stats.pb
Dumped tool data for kernel_stats.pb to ./logs/vae/plugins/profile/2023_01_25_20_2
2_47/Kili.kernel_stats.pb

/Users/abc/miniforge3/envs/py38tfm28/lib/python3.8/site-packages/keras/engine/trai
ning_v1.py:2057: UserWarning: `Model.state_updates` will be removed in a future ve
rsion. This property should not be used in TensorFlow 2.0, as `updates` are applie
d automatically.
    updates = self.state_updates
```

```
2023-01-25 20:22:48.054219: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

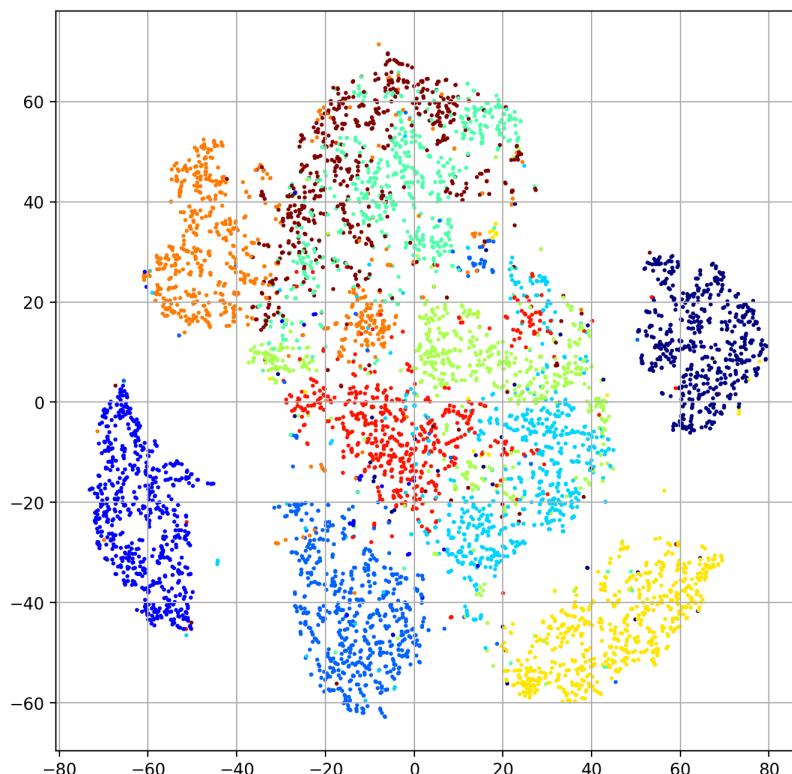
In [25]: `plot_history(history)`



- View t-SNE embedding fo the latent space.

In [27]: `z = encoder.predict(trainXraw)
plot_embedding(z, trainY)`

```
/Users/abc/miniforge3/envs/py38tfm28/lib/python3.8/site-packages/keras/engine/training_v1.py:2079: UserWarning: `Model.state_updates` will be removed in a future version. This property should not be used in TensorFlow 2.0, as `updates` are applied automatically.  
    updates=self.state_updates,  
2023-01-25 20:23:05.313357: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```



- View some reconstruction results

In [29]: `rfig`

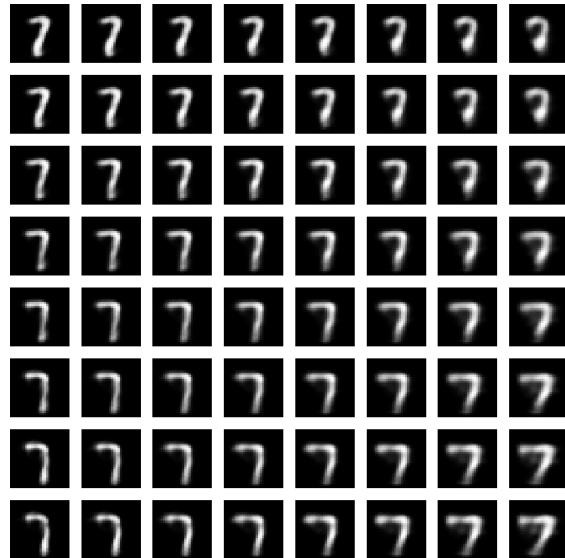
Out[29]:



- visualize the latent space by interpolating between

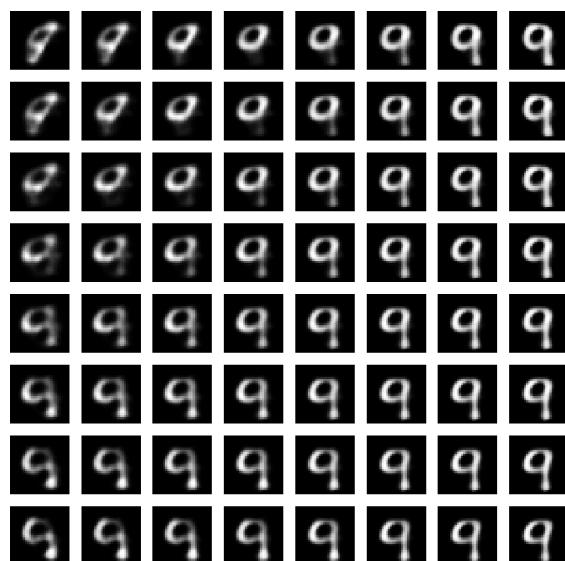
In [32]: rfig

Out[32]:



In [34]: rfig

Out[34]:



Convolutional VAE

- The previous VAE is using fully-connected layers
- Since the inputs are images, then replace the Dense layers with Conv2D and Pooling
 - The encoder has 3 outputs: the latent mean, log-sigma, and the sampled z
 - Latent dimension is 10

In [35]:

```
# the Conv2D encoder
K.clear_session()
random.seed(4487); tf.random.set_seed(4487)
```

```

latent_dim = 10
input_img2 = Input(shape=(28, 28, 1), name='input')
x = Conv2D(16, (3, 3), activation='relu', padding='same', name='enc_conv1')(input_img2)
x = MaxPooling2D((2, 2), padding='same', name='enc_pool1')(x)
x = Conv2D(8, (3, 3), activation='relu', padding='same', name='enc_conv2')(x)
x = MaxPooling2D((2, 2), padding='same', name='enc_pool2')(x)
x = Conv2D(8, (3, 3), activation='relu', padding='same', name='enc_conv3')(x)
x = MaxPooling2D((2, 2), padding='same', name='enc_pool3')(x)
x = Flatten(name='enc_hidden')(x)

# the mean and log-sigma
z_mean = Dense(latent_dim, name='mu')(x)
z_log_sigma = Dense(latent_dim, name='log_sigma')(x)

```

In [36]:

```

# sampling step
epsilon_std = 0.01
def sampling(args):
    z_mean, z_log_sigma = args
    epsilon = K.random_normal(shape=K.shape(z_mean), stddev=epsilon_std)
    return z_mean + K.exp(z_log_sigma) * epsilon

z = Lambda(sampling, name='sampler')([z_mean, z_log_sigma])

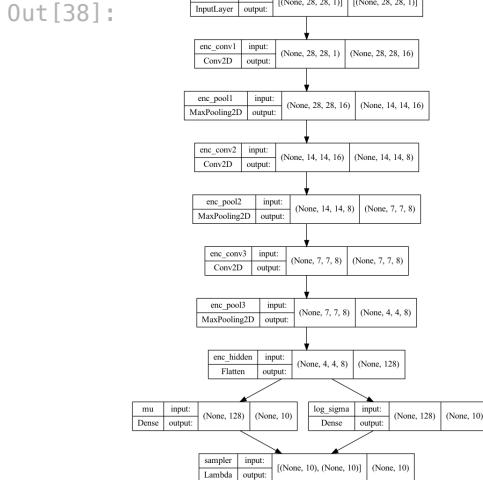
# build the encoder (outputs mean, log-sigma, and z)
encoder = Model(input_img2, [z_mean, z_log_sigma, z], name='encoder')

```

- Encoder summary

In [38]:

```
tmp=tf.keras.utils.plot_model(encoder, show_shapes=True); tmp.width=300; tmp
```



- Same for the decoder

In [39]:

```

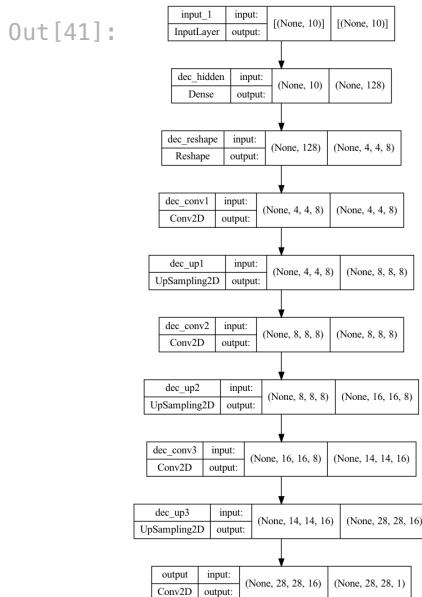
# the Conv2D decoder
encoded_input2 = Input(shape=(latent_dim,))
x = Dense(128, activation='relu', name='dec_hidden')(encoded_input2)
x = Reshape((4,4,8), name='dec_reshape')(x)
x = Conv2D(8, (3, 3), activation='relu', padding='same', name='dec_conv1')(x)
x = UpSampling2D((2, 2), name='dec_up1')(x)
x = Conv2D(8, (3, 3), activation='relu', padding='same', name='dec_conv2')(x)
x = UpSampling2D((2, 2), name='dec_up2')(x)
x = Conv2D(16, (3, 3), activation='relu', name='dec_conv3')(x)
x = UpSampling2D((2, 2), name='dec_up3')(x)
x = Conv2D(1, (3, 3), activation='sigmoid', padding='same', name='output')(x)

decoder = Model(encoded_input2, x, name='decoder')

```

In [41]:

```
tmp=tf.keras.utils.plot_model(decoder, show_shapes=True); tmp.width=250; tmp
```



- make VAE:
 - connect the encoder and decoder
 - select the sampled z of the encoder output as the input into the decoder

In [42]:

```
vae = Model(input_img2, decoder(encoder(input_img2)[2]))
```

In [43]:

```
vae.summary()
```

```

Model: "model"

Layer (type)          Output Shape         Param #
=====
input (InputLayer)     [(None, 28, 28, 1)]   0
encoder (Functional)  [(None, 10),  
                      (None, 10),  
                      (None, 10)]   4484
decoder (Functional)  (None, 28, 28, 1)       3889
=====
Total params: 8,373
Trainable params: 8,373
Non-trainable params: 0
=====
```

- The VAE loss as before

In [44]:

```

def vae_loss(x, x_decoded_mean):
    xent_loss = keras.losses.binary_crossentropy(x, x_decoded_mean)
    kl_loss = - 0.5 * K.mean(K.mean(1 + z_log_sigma - K.square(z_mean)
                                    - K.exp(z_log_sigma), axis=-1), axis=-1)
    return xent_loss + kl_loss
vae.compile(optimizer='adam', loss=vae_loss)

```

- Train the model

In [45]:

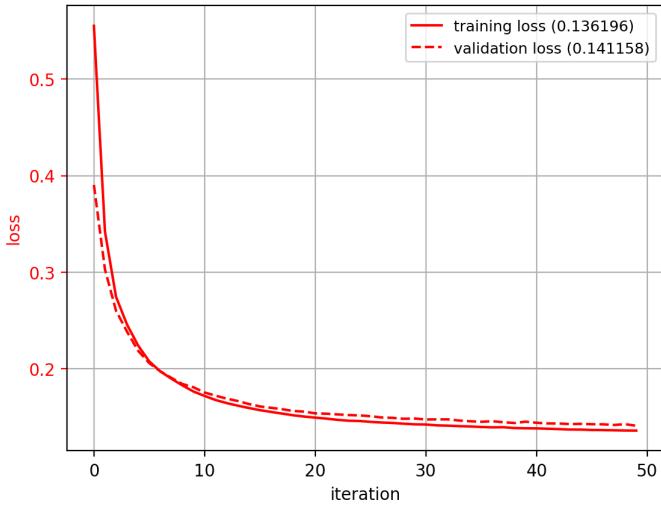
```

# early stopping criteria
earlystop = keras.callbacks.EarlyStopping(monitor='val_loss',
                                           min_delta=0.0001, patience=10,
                                           verbose=1, mode='auto')

history = vae.fit(vtrainI, vtrainI,
                   shuffle=True,
                   epochs=epochs,
                   batch_size=batch_size,
```

```
validation_data=(validI, validI),
callbacks=[earlystop, TensorBoard(log_dir='./logs/vae')],  
verbose=False  
  
2023-01-25 20:23:21.735836: I tensorflow/core/common_runtime/pluggable_device/plug  
gible_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, de  
faulting to 0. Your kernel may not have been built with NUMA support.  
2023-01-25 20:23:21.735855: I tensorflow/core/common_runtime/pluggable_device/plug  
gible_device_factory.cc:271] Created TensorFlow device (/job:localhost/replica:0/t  
ask:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (device: 0, name:  
METAL, pci bus id: <undefined>)  
2023-01-25 20:23:21.750214: I tensorflow/core/grappler/optimizers/custom_graph_opt  
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.  
2023-01-25 20:23:21.775544: I tensorflow/core/grappler/optimizers/custom_graph_opt  
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.  
2023-01-25 20:23:21.849793: I tensorflow/core/grappler/optimizers/custom_graph_opt  
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.  
2023-01-25 20:23:22.262396: I tensorflow/core/profiler/lib/profiler_session.cc:11  
0] Profiler session initializing.  
2023-01-25 20:23:22.262412: I tensorflow/core/profiler/lib/profiler_session.cc:12  
5] Profiler session started.  
2023-01-25 20:23:22.287540: I tensorflow/core/profiler/lib/profiler_session.cc:67]  
Profiler session collecting data.  
2023-01-25 20:23:22.288099: I tensorflow/core/profiler/lib/profiler_session.cc:14  
3] Profiler session tear down.  
2023-01-25 20:23:22.288603: I tensorflow/core/profiler/rpc/client/save_profile.cc:  
136] Creating directory: ./logs/vae/plugins/profile/2023_01_25_20_23_22  
  
2023-01-25 20:23:22.288956: I tensorflow/core/profiler/rpc/client/save_profile.cc:  
142] Dumped gzipped tool data for trace.json.gz to ./logs/vae/plugins/profile/2023  
_01_25_20_23_22/Kili.trace.json.gz  
2023-01-25 20:23:22.289450: I tensorflow/core/profiler/rpc/client/save_profile.cc:  
136] Creating directory: ./logs/vae/plugins/profile/2023_01_25_20_23_22  
  
2023-01-25 20:23:22.289796: I tensorflow/core/profiler/rpc/client/save_profile.cc:  
142] Dumped gzipped tool data for memory_profile.json.gz to ./logs/vae/plugins/pro  
file/2023_01_25_20_23_22/Kili.memory_profile.json.gz  
2023-01-25 20:23:22.290133: I tensorflow/core/profiler/rpc/client/capture_profile.  
cc:251] Creating directory: ./logs/vae/plugins/profile/2023_01_25_20_23_22  
Dumped tool data for xplane.pb to ./logs/vae/plugins/profile/2023_01_25_20_23_22/K  
ili.xplane.pb  
Dumped tool data for overview_page.pb to ./logs/vae/plugins/profile/2023_01_25_20_  
23_22/Kili.overview_page.pb  
Dumped tool data for input_pipeline.pb to ./logs/vae/plugins/profile/2023_01_25_20  
_23_22/Kili.input_pipeline.pb  
Dumped tool data for tensorflow_stats.pb to ./logs/vae/plugins/profile/2023_01_25_  
20_23_22/Kili.tensorflow_stats.pb  
Dumped tool data for kernel_stats.pb to ./logs/vae/plugins/profile/2023_01_25_20_2  
3_22/Kili.kernel_stats.pb  
  
/Users/abc/miniforge3/envs/py38tfm28/lib/python3.8/site-packages/keras/engine/trai  
ning_v1.py:2057: UserWarning: `Model.state_updates` will be removed in a future ve  
rsion. This property should not be used in TensorFlow 2.0, as `updates` are applie  
d automatically.  
    updates = self.state_updates  
2023-01-25 20:23:23.467634: I tensorflow/core/grappler/optimizers/custom_graph_opt  
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```

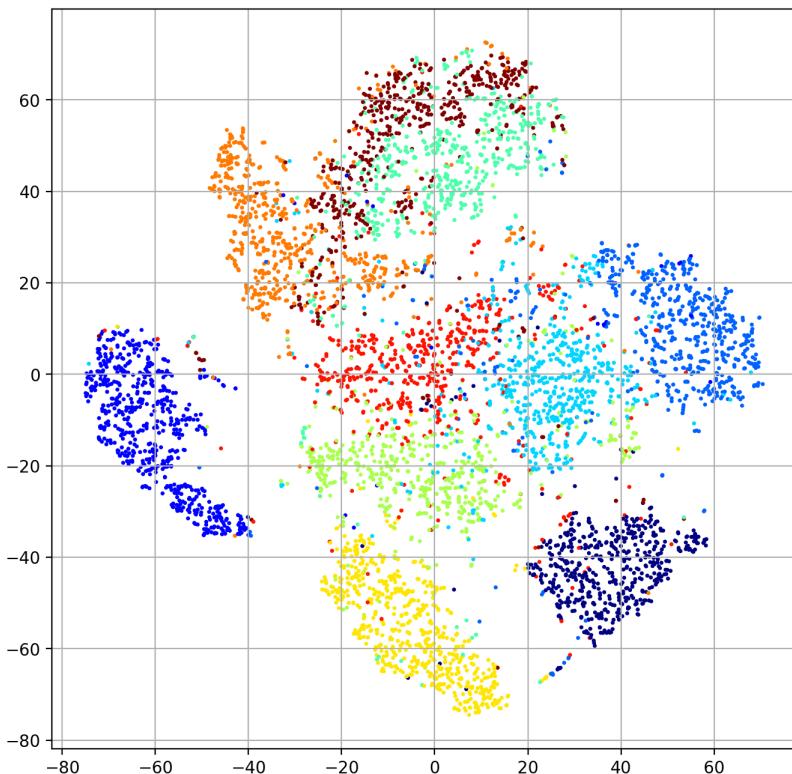
In [46]: `plot_history(history)`



- Visualize the 10-dim latent space using t-SNE

```
In [47]: # select z_mean from the encoder
z = encoder.predict(trainI)[0]
plot_embedding(z.reshape((z.shape[0], -1)), trainY)

/Users/abc/miniforge3/envs/py38tfm28/lib/python3.8/site-packages/keras/engine/training_v1.py:2079: UserWarning: `Model.state_updates` will be removed in a future version. This property should not be used in TensorFlow 2.0, as `updates` are applied automatically.
    updates=self.state_updates,
2023-01-25 20:24:14.464274: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```



- Visualize the reconstruction

```
In [49]: rfig
```

Out[49]:

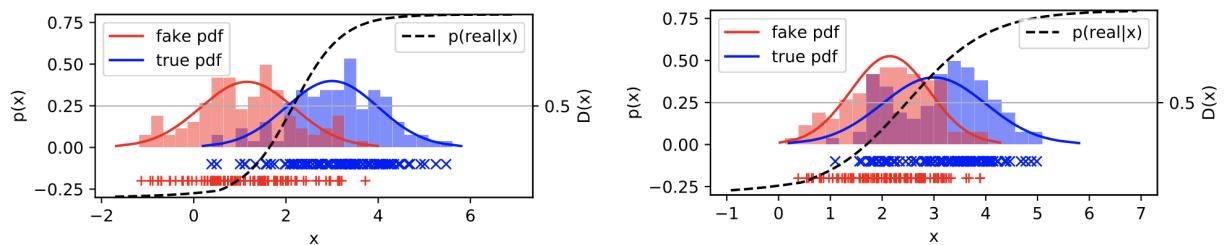
- Visualize the latent space between a 7, 1, 9, and 4

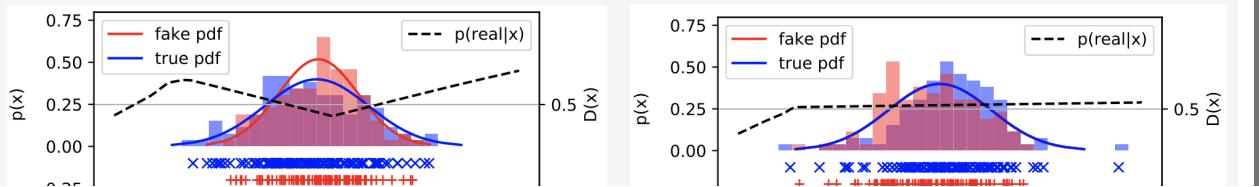
In [51]:

rfig

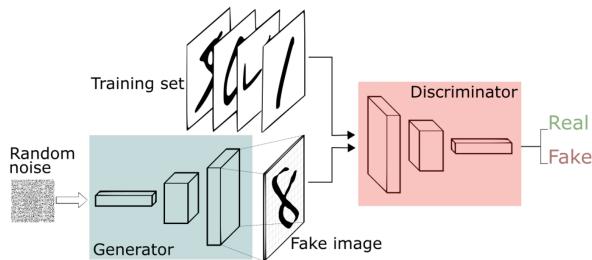
Generative Adversarial Networks (GAN)

- **Goal:** generate samples from a probability density $p_{data}(x)$ which is represented by a training set $\{x_i\}$.
- **Setup**
 - *true data samples: $x \sim p_d(x)$*
 - we have samples $\mathcal{X} = \{x_i\}$, but not p_d .
 - *generated samples: $x \sim p_m(x)$*
 - use the reparam trick to generate samples
 - $x = g(z)$, $z \sim N(0, I)$
 - $g(\cdot)$ is a neural network
 - Usually maximize the likelihood of $p_m(x)$ using samples \mathcal{X} .
- **Problem:**
 - cannot explicitly compute $p_m(x)$ because of the reparam trick.
- **Solution:** look at the samples directly
- Build a classifier to predict whether a sample is *real* (from training samples) or *fake* (generated)





- Two models:
 - Generator (G) - make fake samples to fool the discriminator
 - Discriminator (D) - classify x as real or fake



- **Advantages:**
 - don't need to approximate the likelihood function or inference procedure
 - don't need the partition function (normalization constant)
 - easy to generate samples (good for synthesis)
- **Disadvantages:**
 - cannot compute the likelihood value (probability density) of a sample
 - difficult to train
 - ...

GAN Training

- Define two class labels:
 - $y = 1$ means "real" sample
 - $y = 0$ means "fake" sample
- Define discriminator (classifier): $D(x) = p(y = 1|x)$
 - probability of real sample
- Cross-entropy loss to train the discriminator
 - let $\{x_i, y_i\}$ be a dataset of true/fake samples.

$$\begin{aligned} & \min_D - \sum_i y_i \log D(x_i) + (1 - y_i) \log(1 - D(x_i)) \\ &= \max_D \underbrace{\sum_{y_i=1} \log D(x_i)}_{\text{real examples}} + \underbrace{\sum_{y_i=0} \log(1 - D(x_i))}_{\text{generated examples}} \end{aligned}$$

- Replace the sum of samples with expectations

$$\begin{aligned} & \max_D \sum_{y_i=1} \log D(x_i) + \sum_{y_i=0} \log(1 - D(x_i)) \\ &= \max_D \mathbb{E}_{x \sim p_d} [\log D(x)] + \mathbb{E}_{x \sim p_m} [\log(1 - D(x))] \\ &= \max_D \mathbb{E}_{x \sim p_d} [\log D(x)] + \mathbb{E}_{z \sim \mathcal{N}(0, I)} [\log(1 - D(G(z)))] \end{aligned}$$

- The G 's objective is to maximize the cross entropy loss, i.e., make D predict wrongly.
- Final training problem:

$$\min_G \max_D \mathbb{E}_{x \sim p_d} [\log D(x_i)] + \mathbb{E}_{z \sim \mathcal{N}(0, I)} [\log(1 - D(G(z)))]$$

- **Notes:**

- it is a *minimax* game (zero-sum game)
 - the values for D and G move in opposite directions
 - the equilibrium point is a *saddle* point
 - i.e., a local max for D, and local min for G
- G can also be trained to make the D predict "real"
 - $\min_G -\mathbb{E}_z [\log D(G(z))]$
 - i.e., pass the generated samples with "real" class label to D

Analysis of GAN

- Assuming that G and D have enough capacity
- Considering the defined training problem:
 - the optimal discriminator is
 - $D^*(x) = \frac{p_d(x)}{p_d(x) + p_m(x)} = p(\text{real}|x)$
 - given the optimal discriminator, the optimal generator is
 - $p_m^*(x) = p_d(x)$
- thus, ideally, we can obtain the true distribution.

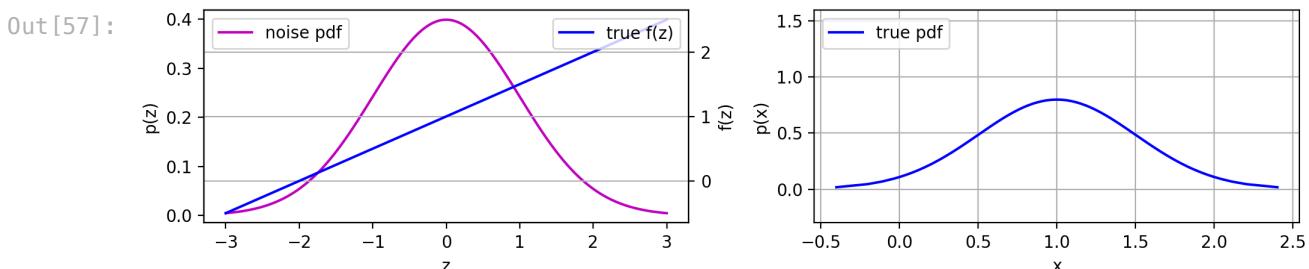
1-D GAN example

- the true distribution
 - transform `xform` of $\text{Gaussian}(0,1)$

```
In [55]: # transformed gaussian
xform = lambda x: (0.5*x)+1
sample_range = get_bounds(xform)
```

```
In [56]: # samples from the true distribution
def sample(N, mu=0.0, sigma=1):
    samples = xform(random.normal(mu, sigma, N))
    return reshape(samples, (len(samples), 1))
```

```
In [57]: plot_epoch(onlygt=True)
```



- samples from the noise source

```
In [58]: # samples from the generator noise source (Gaussian(0,1))
def g_sample(N, rng=5):
    a = random.normal(0, 1, N)
    return reshape(a, (len(a), 1))
```

- function to build a generator
 - the input is a 1-D noise source
 - output is a 1-D sample

```
In [59]: # build the generator
def gen(hidden_size, n_layers=2, activation="tanh"):
    g_input = Input(shape=[1], name="in1")
    x = g_input
    # hidden layers
    for n in range(n_layers):
        x = Dense(hidden_size, activation=activation,
                  name="g{}".format(n))(x)
    # output layer
    ho = Dense(1, name="go")(x)
    g = Model(g_input, ho, name='generator')
    return g
```

- function to build a discriminator
 - the input is a 1-D sample
 - the output is class label (0=fake, 1=real)
 - need to `compile` since D is trained by itself

```
In [60]: # build the discriminator
def disc(hidden_size, n_layers=2, activation='tanh'):
    inputd = Input(shape=[1], name="in2")
    x = inputd
    for n in range(n_layers):
        x = Dense(hidden_size * 2, activation=activation,
                  name="d{}".format(n))(x)
    ho = Dense(1, activation="sigmoid", name="do")(x)
    d = Model(inputd, ho, name='discriminator')
    d.compile(loss="binary_crossentropy", optimizer="sgd",
               metrics=['accuracy'])
    return d
```

- we will alternate training D and G
- this function sets the trainable flags in a `Model` to prevent/allow updating the weights

```
In [61]: # set trainable flag for a model
def make_trainable(net, val):
    net.trainable = val # this only works before compiling
    for l in net.layers:
        l.trainable = val # this works before/after compiling
```

- build the G and D
- connect G and D to make GAN
 - compile GAN for training G

```
In [62]: K.clear_session()
random.seed(5489); tf.random.set_seed(5488)

# the generator
noise = Input(shape=[1])
generator = gen(10, activation='softplus')
g = generator(noise)

# the discriminator
discriminator = disc(10, activation='tanh')
d = discriminator(g)
make_trainable(discriminator, False)

# the GAN w/ noise input and discrim output
GAN = Model(noise, d)
GAN.compile(loss="binary_crossentropy", optimizer="sgd", metrics=['accuracy'])
GAN.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
<hr/>		

```

input_1 (InputLayer)      [(None, 1)]          0
generator (Functional)    (None, 1)            141
discriminator (Functional) (None, 1)           481
=====
Total params: 622
Trainable params: 141
Non-trainable params: 481

```

In [63]: `generator.summary()`

```

Model: "generator"
-----  

Layer (type)        Output Shape       Param #
-----  

in1 (InputLayer)    [(None, 1)]        0  

g0 (Dense)          (None, 10)         20  

g1 (Dense)          (None, 10)         110  

go (Dense)          (None, 1)          11
-----  

Total params: 141
Trainable params: 141
Non-trainable params: 0

```

In [64]: `discriminator.summary()`

```

Model: "discriminator"
-----  

Layer (type)        Output Shape       Param #
-----  

in2 (InputLayer)    [(None, 1)]        0  

d0 (Dense)          (None, 20)         40  

d1 (Dense)          (None, 20)         420  

do (Dense)          (None, 1)          21
-----  

Total params: 962
Trainable params: 481
Non-trainable params: 481

```

GAN Training

- Loop over training the discriminator and the generator
 - train the discriminator well (20 updates), vs 1 update for the generator

In [66]:

```

def train_for_n(nb_epoch=50, plt_frg=25, BATCH_SIZE=100, lr=0.01, g_inner=1, d_inner=20):
    epochs = {'g_loss':[], 'g_acc':[], 'd_loss':[], 'd_acc':[], 'figs':[]}

    for e in range(nb_epoch):
        # Make generative sample
        sample_batch = sample(BATCH_SIZE)
        noise_gen = g_sample(BATCH_SIZE)
        generated_sample = generator.predict(noise_gen)

        # Train discriminator on generated samples
        X = concatenate((sample_batch, generated_sample))
        # class 1=real, class 0 = fake
        y = concatenate((ones((BATCH_SIZE,)), zeros((BATCH_SIZE,))))

```

```

# train the discriminator a few steps on this batch
make_trainable(discriminator, True)           # make it trainable
K.set_value(discriminator.optimizer.lr, lr) # set the learning rate
for i in range(d_inner):                      # train...
    d_loss = discriminator.train_on_batch(X,y)

epochs['d_loss'].append(d_loss[0])
epochs['d_acc'].append(d_loss[1])

# Train the generator
y2 = ones((BATCH_SIZE,)) # target for generator is real class (1)
noise_gen = g_sample(BATCH_SIZE)

# keep discriminator fixed

# run a few updates on this batch
make_trainable(discriminator, False) # freeze discriminator
K.set_value(GAN.optimizer.lr, lr) # set learning rate
for i in range(g_inner):          # train...
    g_loss = GAN.train_on_batch(noise_gen, y2)

epochs['g_loss'].append(g_loss[0])
epochs['g_acc'].append(g_loss[1])

# some plots
if (e % plt_frq==0) or (e+1==nb_epoch):
    fig=plot_epoch(noise_gen, sample_batch, "epoch {}".format(e))
    epochs['figs'].append(fig)

return epochs

```

- train the GAN for 3000 epochs

In [67]: `epochs = train_for_n(nb_epoch=3000, lr=0.01, plt_frq=500)`

```

2023-01-25 20:24:30.842743: I tensorflow/core/common_runtime/pluggable_device/plug
gable_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, de
faulting to 0. Your kernel may not have been built with NUMA support.
2023-01-25 20:24:30.842764: I tensorflow/core/common_runtime/pluggable_device/plug
gable_device_factory.cc:271] Created TensorFlow device (/job:localhost/replica:0/t
ask:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (device: 0, name:
METAL, pci bus id: <undefined>)
2023-01-25 20:24:30.847946: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:30.854803: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:30.887429: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:30.921538: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:30.925908: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:30.929600: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:30.954262: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:30.958606: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:30.961102: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:31.060969: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:31.168323: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:31.174022: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:31.177585: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:31.206336: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:31.212040: I tensorflow/core/grappler/optimizers/custom_graph_opt

```

```

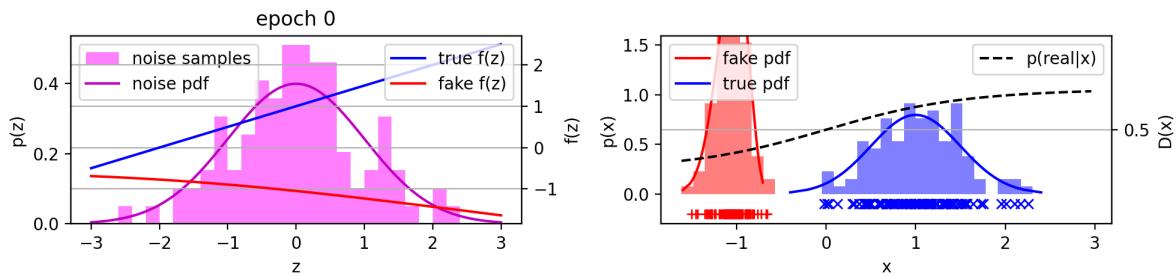
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:31.214883: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:31.335185: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:24:31.470135: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

```

- Visualization:
 - (left) noise source and transformations
 - (right) true/fake pdfs; discriminator function $p(y = 1|x)$ (prob of true sample)

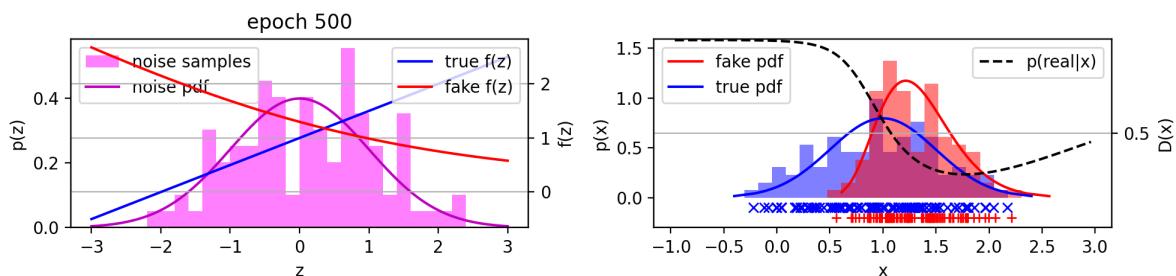
In [68]: `epochs['figs'][0]`

Out[68]:



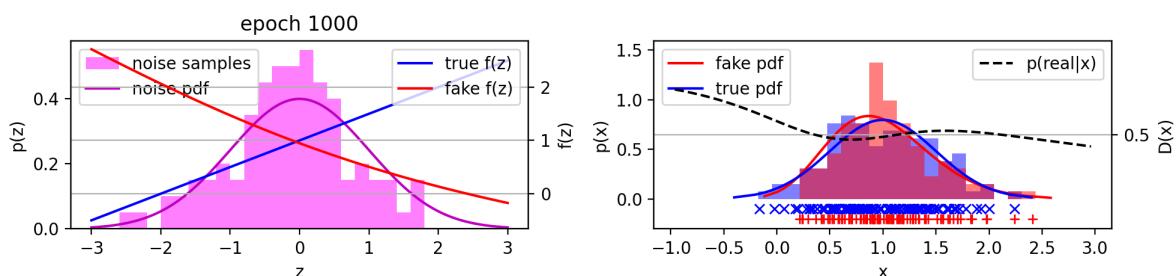
In [69]: `epochs['figs'][1]`

Out[69]:



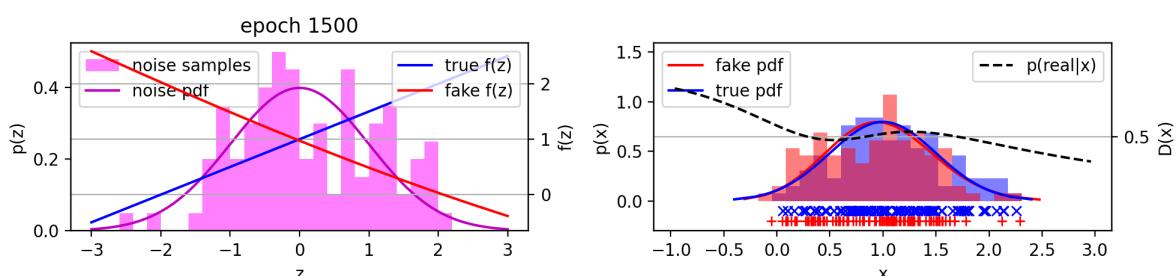
In [70]: `epochs['figs'][2]`

Out[70]:



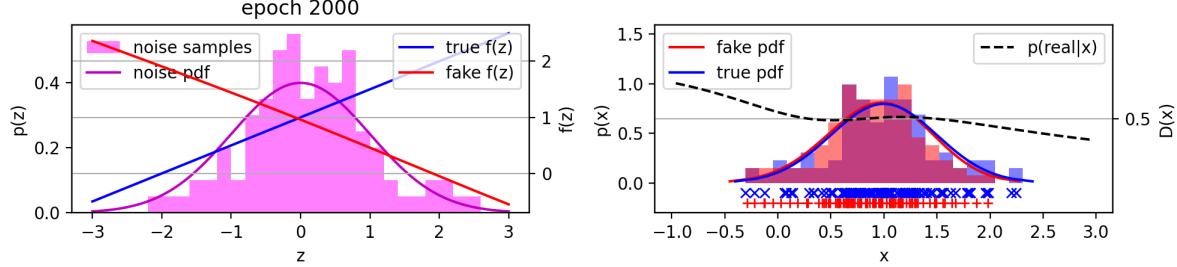
In [71]: `epochs['figs'][3]`

Out[71]:



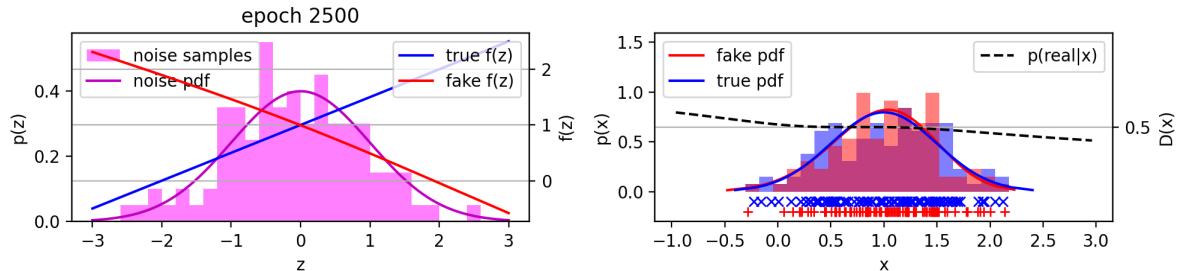
In [72]: `epochs['figs'][4]`

```
Out[72]:
```



```
In [73]: epochs['figs'][5]
```

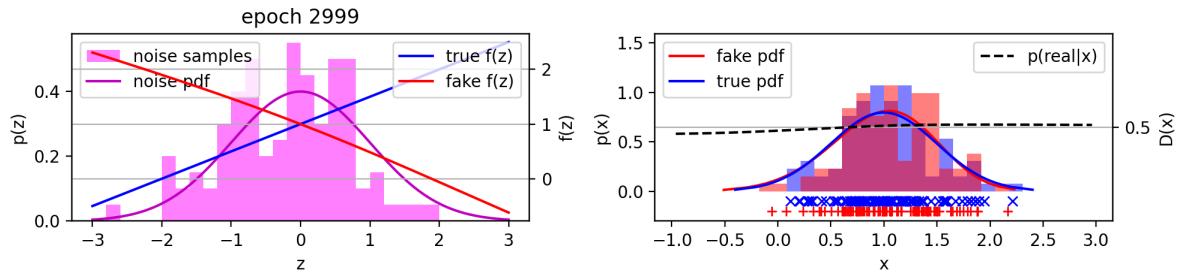
```
Out[73]:
```



- Converges to the true transformation.
 - discriminator can't tell the difference

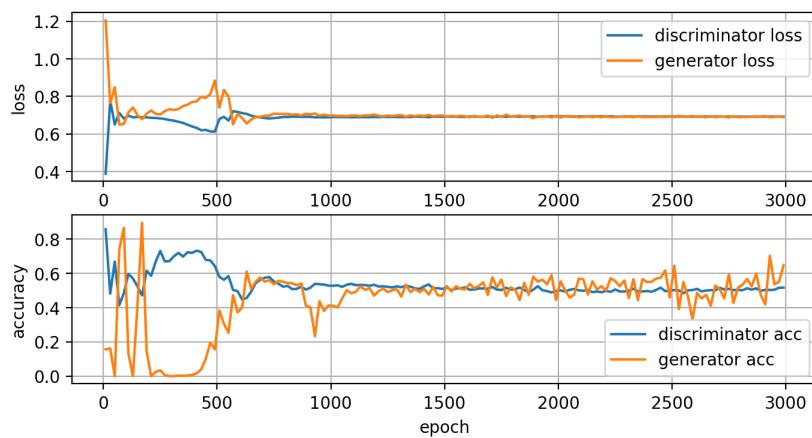
```
In [74]: epochs['figs'][6]
```

```
Out[74]:
```



- loss/accuracy during training.
 - the values will oscillate, since the two networks are playing against each other.
 - at convergence, both G & D obtain 50% accuracy (random chance).

```
In [75]: plot_epochs(epochs)
```

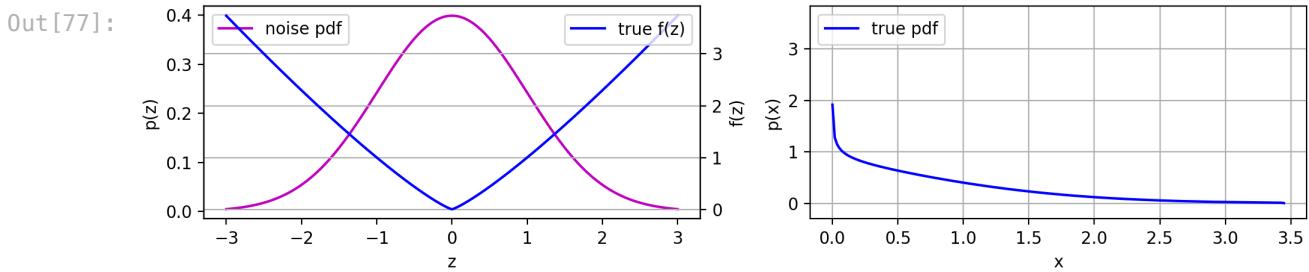


Another example

- distribution with a long-tail

```
In [76]: xform = lambda x: abs(x)**1.2
sample_range = get_bounds(xform)
```

```
In [77]: plot_epoch(onlygt=True)
```



- build the GAN

```
In [78]: K.clear_session()
random.seed(5489); tf.random.set_seed(5489)
noise = Input(shape=[1])
generator = gen(10, activation='relu')
g = generator(noise)
discriminator = disc(10, activation='relu')
d = discriminator(g)
make_trainable(discriminator, False)
GAN = Model(noise, d)
GAN.compile(loss="binary_crossentropy", optimizer="sgd", metrics=['accuracy'])
```

- train for 3000 epochs

```
In [79]: epochs = train_for_n(nb_epoch=3000, lr=0.01, plt_frg=500)
```

2023-01-25 20:30:24.777467: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel may not have been built with NUMA support.
2023-01-25 20:30:24.777488: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id: <undefined>)
2023-01-25 20:30:24.782135: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:24.789004: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:24.823720: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:24.846783: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:24.850821: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:24.853938: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:24.877972: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:24.882310: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:24.884685: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:24.964359: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:25.074930: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:25.080190: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:25.083465: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:25.110549: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:25.116362: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

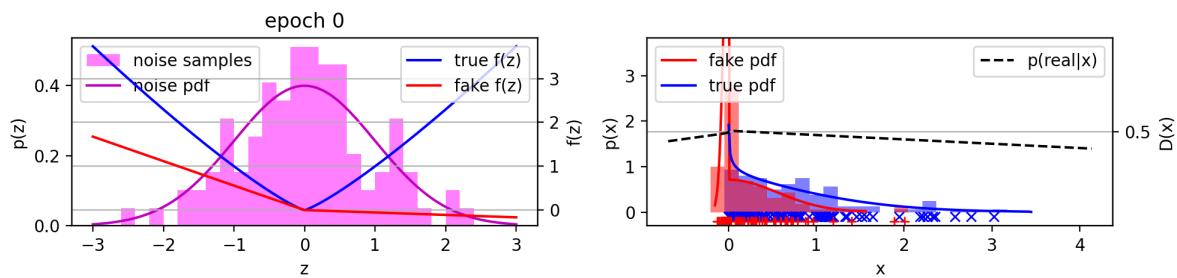
```

2023-01-25 20:30:25.118978: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:25.216480: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:30:25.424121: I tensorflow/core/grappler/optimizers/custom_graph_opt
imizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

```

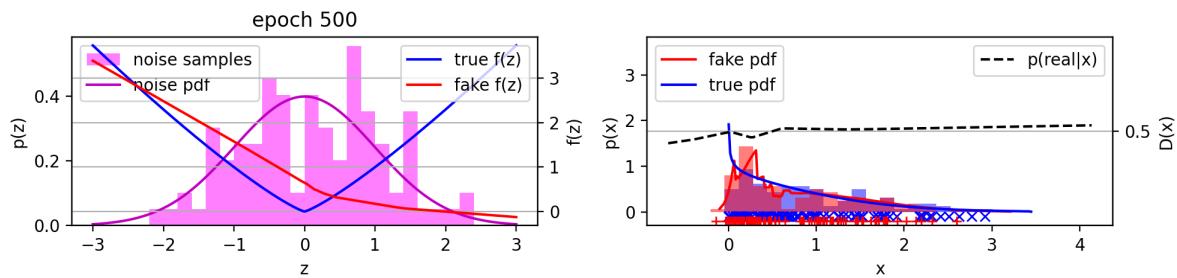
In [80]: `epochs['figs'][0]`

Out[80]:



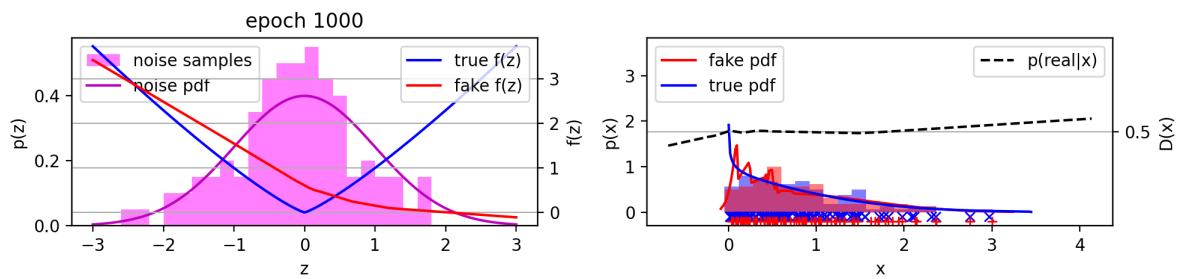
In [81]: `epochs['figs'][1]`

Out[81]:



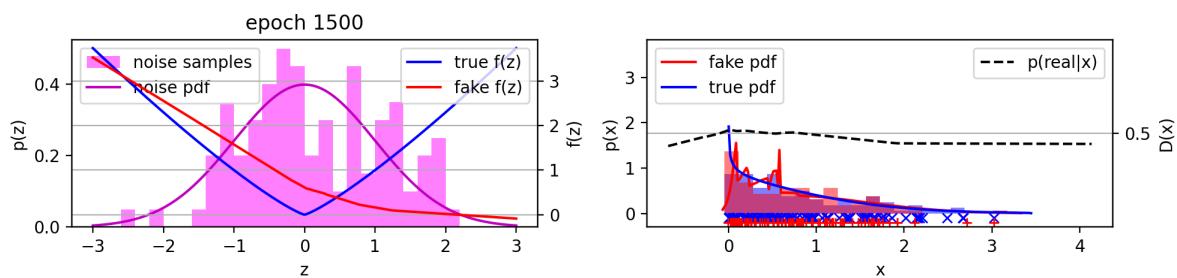
In [82]: `epochs['figs'][2]`

Out[82]:



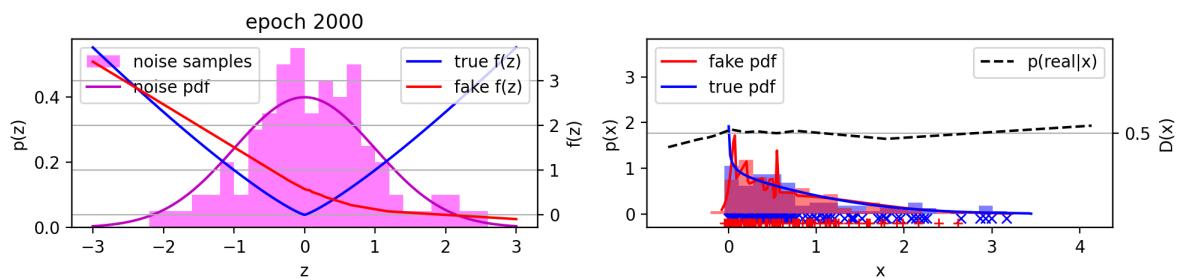
In [83]: `epochs['figs'][3]`

Out[83]:



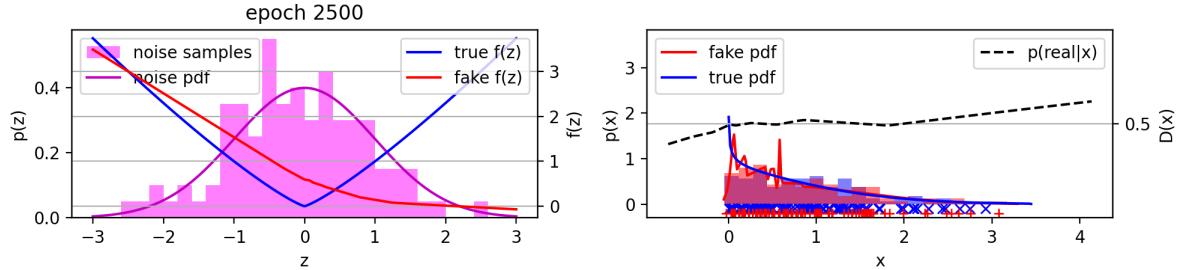
In [84]: `epochs['figs'][4]`

Out[84]:



In [85]: `epochs['figs'][5]`

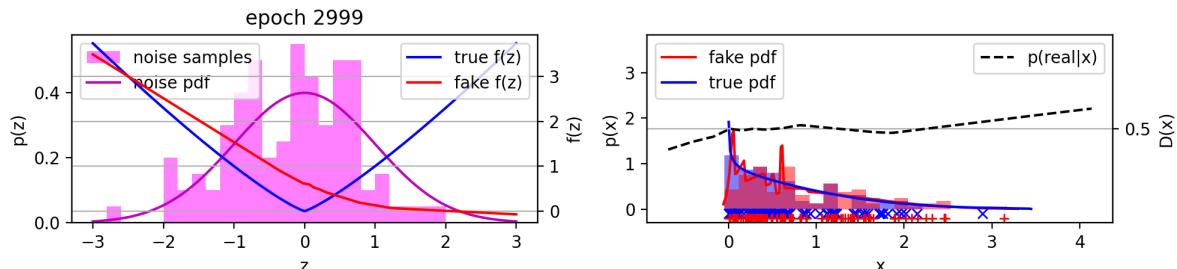
```
Out[85]:
```



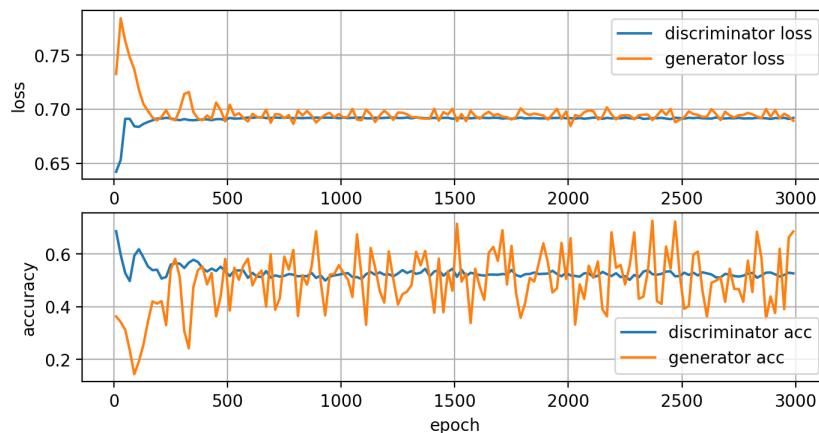
- converges to an equivalent transformation to the real one.

```
In [86]: epochs['figs'][6]
```

```
Out[86]:
```



```
In [87]: plot_epochs(epochs)
```



Mode collapse problem

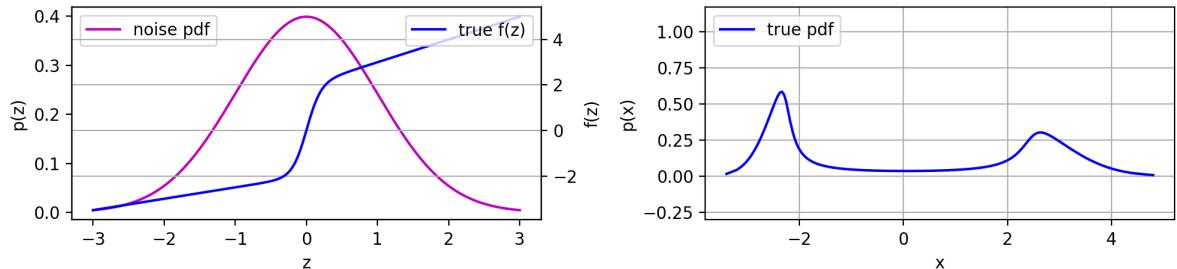
- GANs also work to learn multi-modal distributions

```
In [88]: # two modes
```

```
sigmoid = lambda x,a: 1/(1+exp(-a*x))
xform = lambda x: (0.5*x-2)*sigmoid(-x,10) + (x+2)*sigmoid(x,10)
sample_range = get_bounds(xform)
```

```
In [89]: plot_epoch(onlygt=True)
```

```
Out[89]:
```

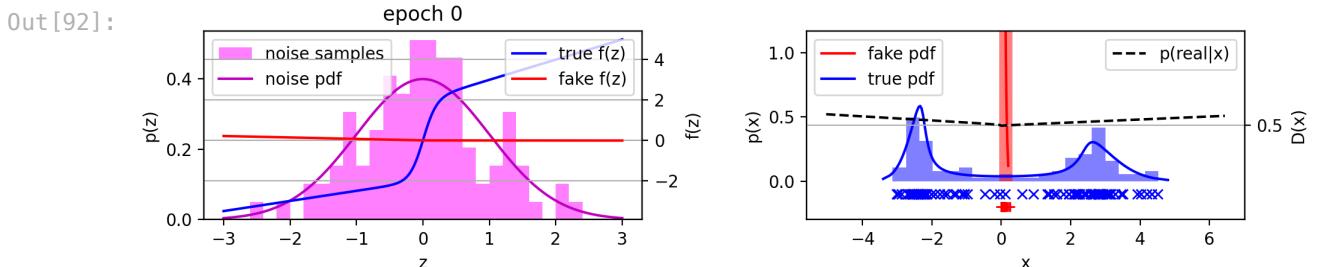


- train 3000 epochs

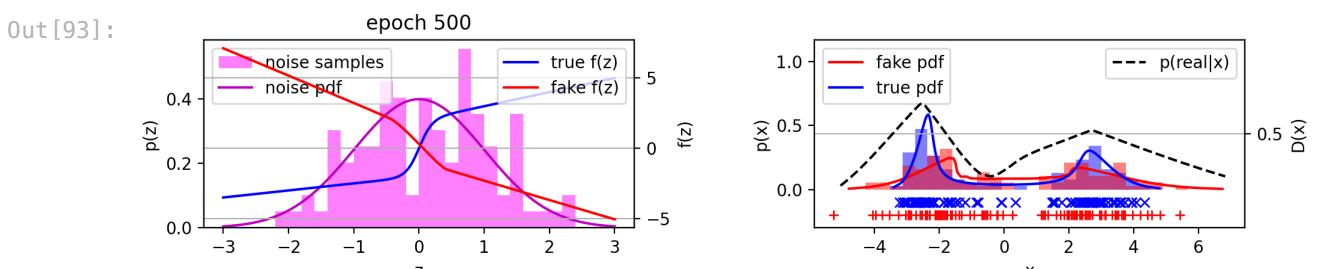
```
In [91]: epochs = train_for_n(nb_epoch=3000, lr=0.01, plt_frg=500)
```

2023-01-25 20:36:01.060915: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel may not have been built with NUMA support.
2023-01-25 20:36:01.060939: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id: <undefined>)
2023-01-25 20:36:01.065167: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.071950: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.105224: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.132430: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.136558: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.139555: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.165359: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.170034: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.172477: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.266105: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.388049: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.394342: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.397976: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.427604: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.434035: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.437259: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.548121: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
2023-01-25 20:36:01.653903: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.

```
In [92]: epochs['figs'][0]
```

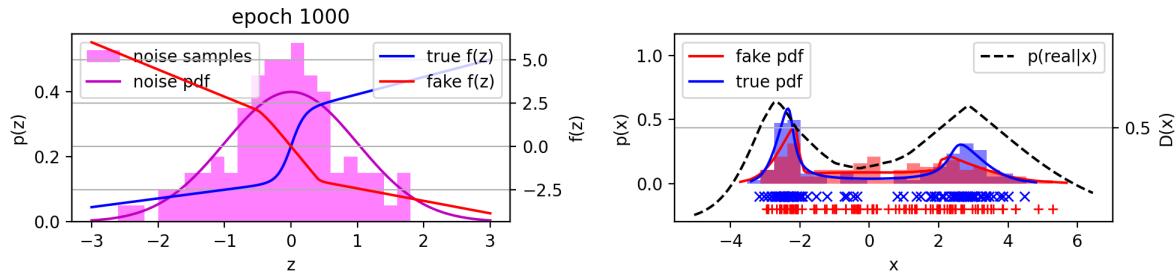


```
In [93]: epochs['figs'][1]
```



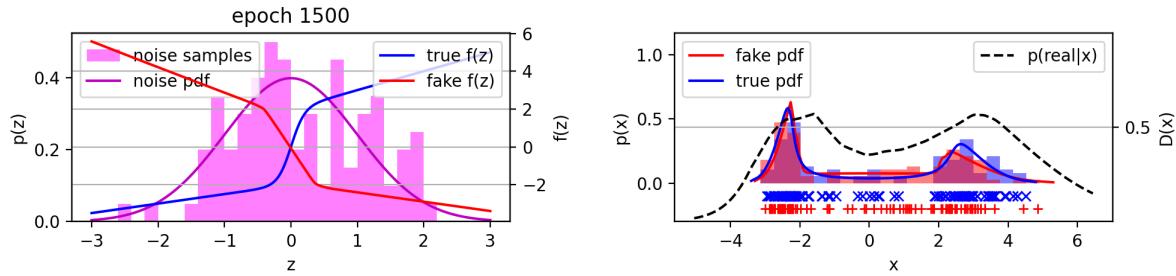
```
In [94]: epochs['figs'][2]
```

```
Out[94]:
```



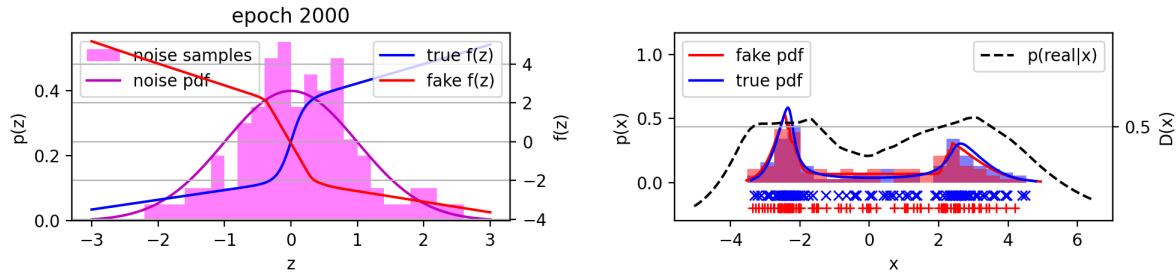
```
In [95]: epochs['figs'][3]
```

```
Out[95]:
```



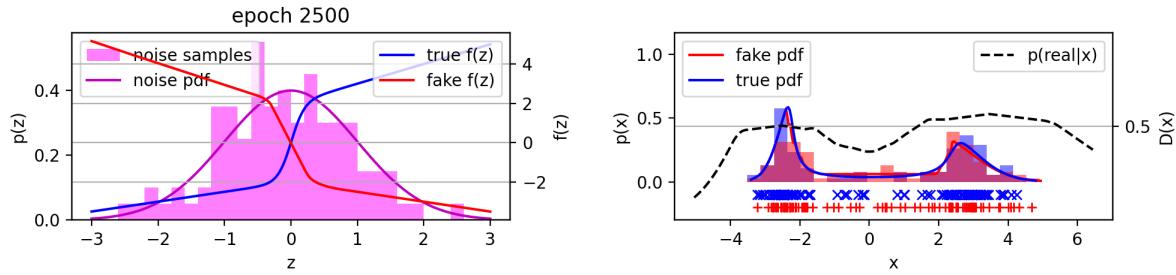
```
In [96]: epochs['figs'][4]
```

```
Out[96]:
```



```
In [97]: epochs['figs'][5]
```

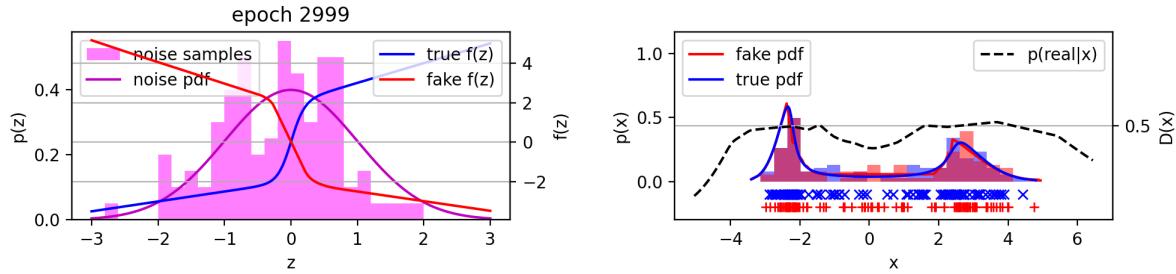
```
Out[97]:
```



- represents both modes of the target density
- converges to true transformation

```
In [98]: epochs['figs'][6]
```

```
Out[98]:
```



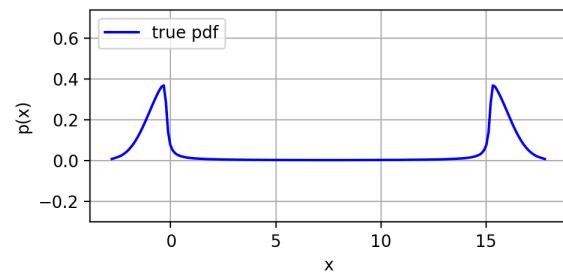
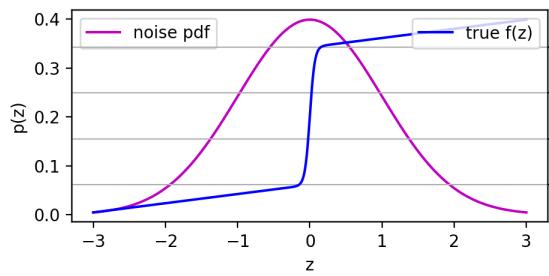
- try another example
 - modes are farther apart

```
In [99]: # two modes
```

```
sigmoid = lambda x,a: 1/(1+exp(-a*x))
xform = lambda x: (x)*sigmoid(-x,30) + (x+15)*sigmoid(x,30)
sample_range = get_bounds(xform)
```

```
In [100]: plot_epoch(onlygt=True)
```

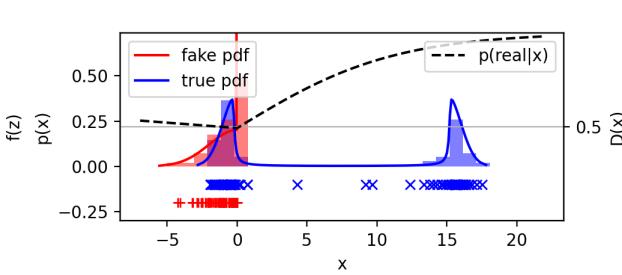
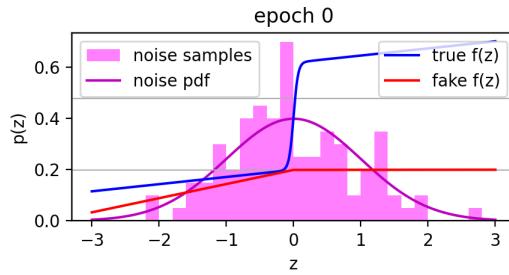
```
Out[100]:
```



- G collapses to sample from only one of the modes
 - G can always fool D with this one mode

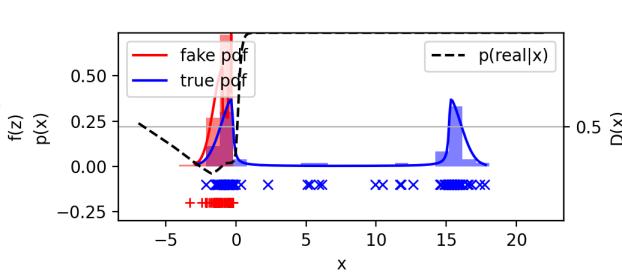
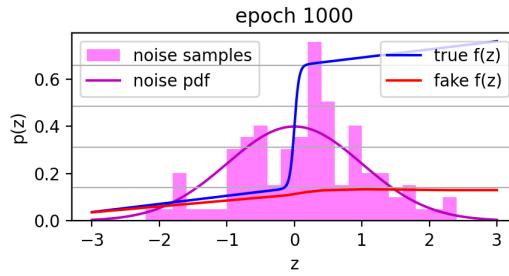
```
In [103]: epochs['figs'][0]
```

```
Out[103]:
```



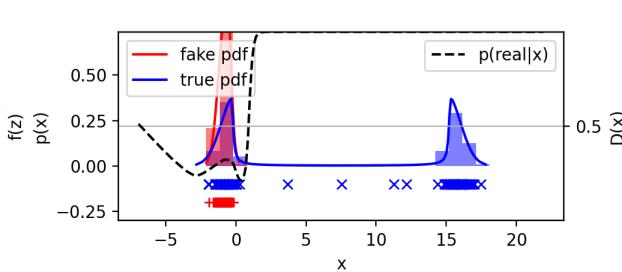
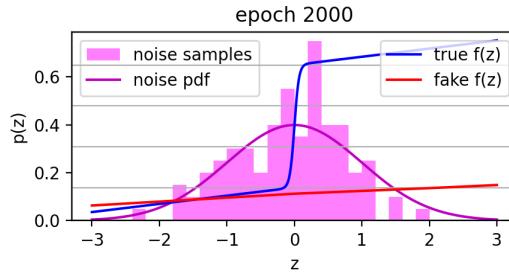
```
In [104]: epochs['figs'][1]
```

```
Out[104]:
```



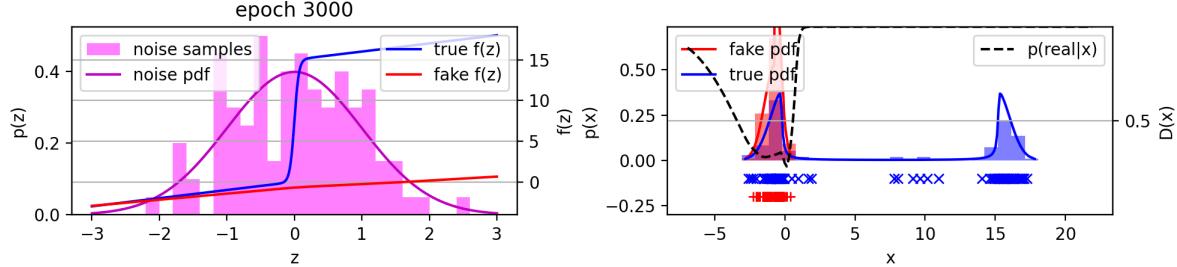
```
In [105]: epochs['figs'][2]
```

```
Out[105]:
```



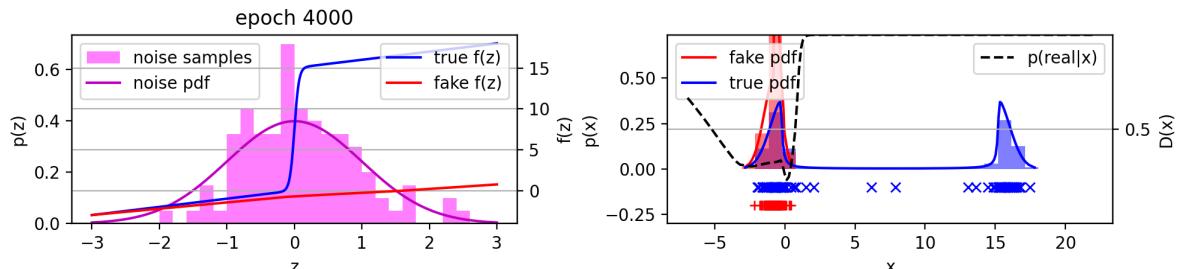
```
In [106]: epochs['figs'][3]
```

```
Out[106]:
```



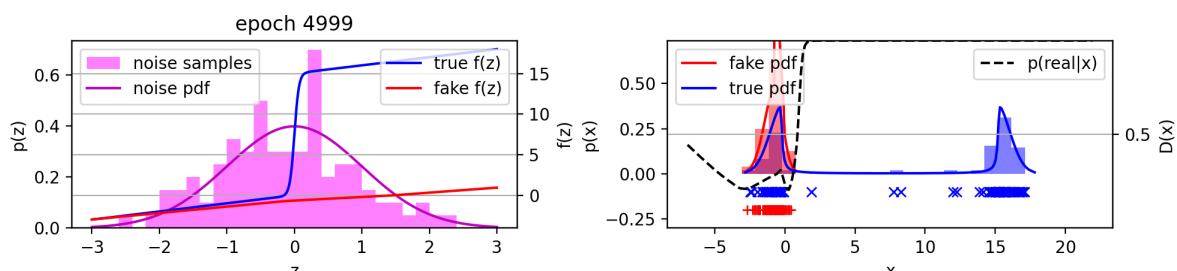
```
In [107]: epochs['figs'][4]
```

```
Out[107]:
```



```
In [108]: epochs['figs'][5]
```

```
Out[108]:
```



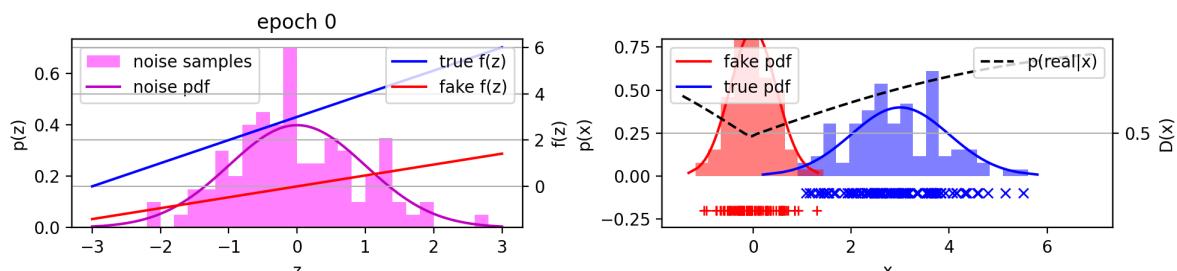
- **Mode collapse problem**

- G samples from one mode.
- when D catches this, G switches to another mode
- and so on

example from beginning

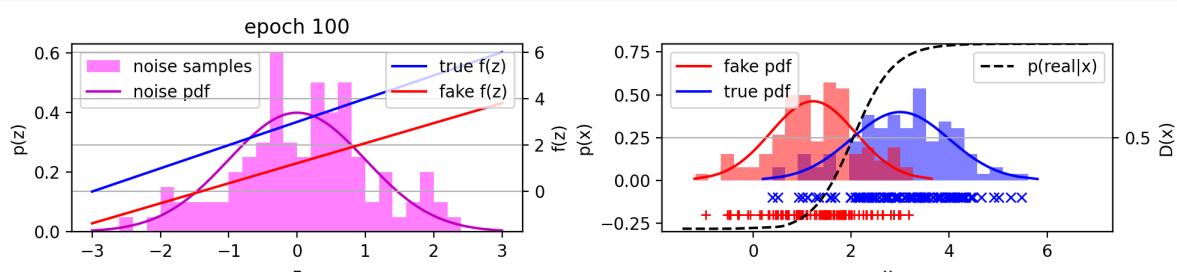
```
In [111]: epochs['figs'][0]
```

```
Out[111]:
```



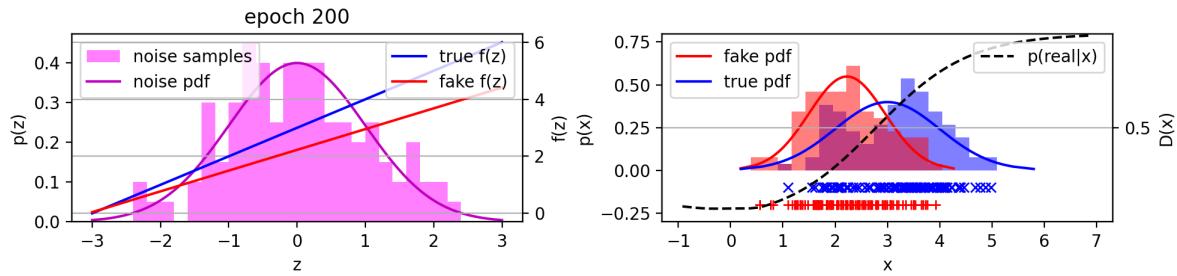
```
In [112]: epochs['figs'][1]
```

```
Out[112]:
```



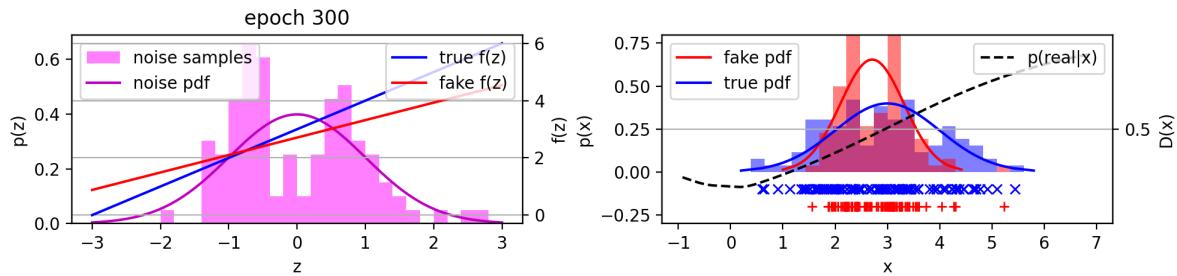
```
In [113]: epochs['figs'][2]
```

Out[113]:



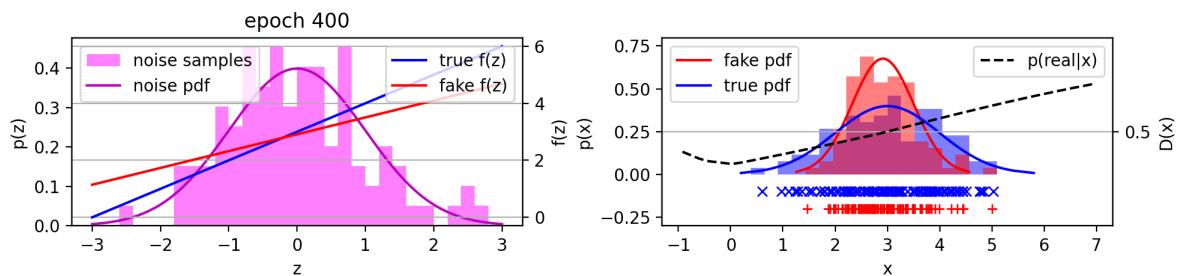
```
In [114]: epochs['figs'][3]
```

Out[114]:



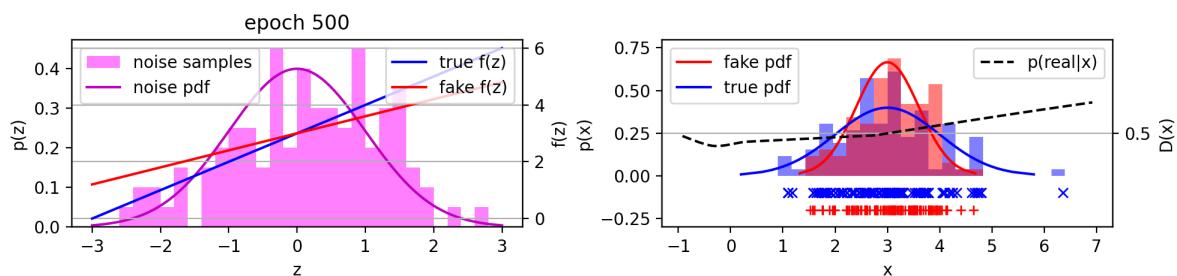
```
In [115]: epochs['figs'][4]
```

Out[115]:



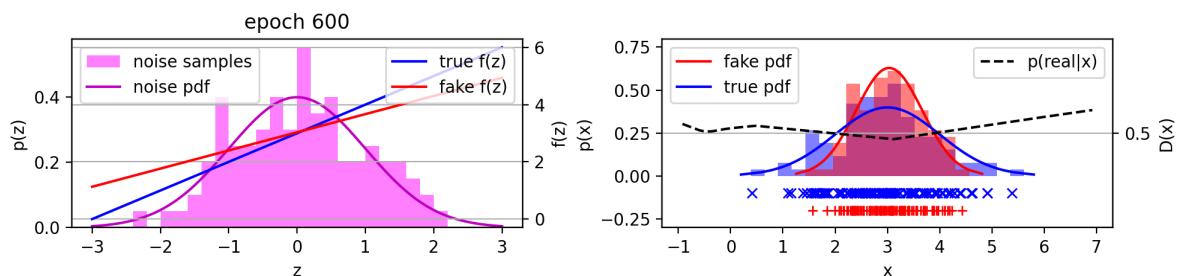
```
In [116]: epochs['figs'][5]
```

Out[116]:



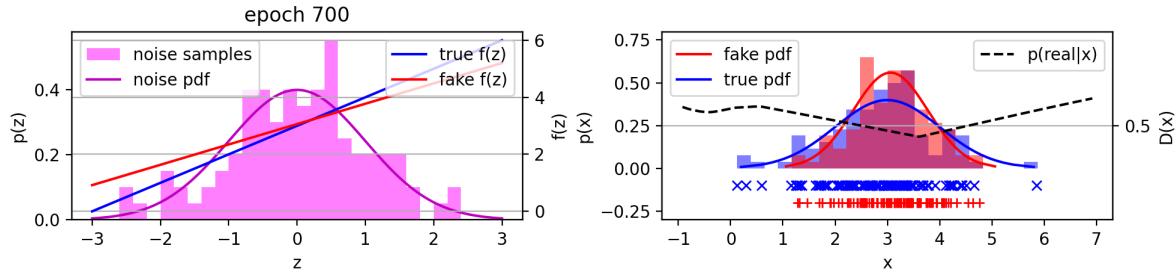
```
In [117]: epochs['figs'][6]
```

Out[117]:



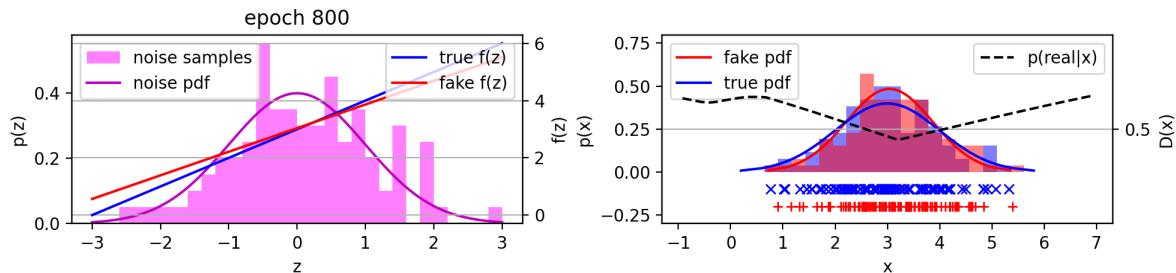
```
In [118]: epochs['figs'][7]
```

```
Out[118]:
```



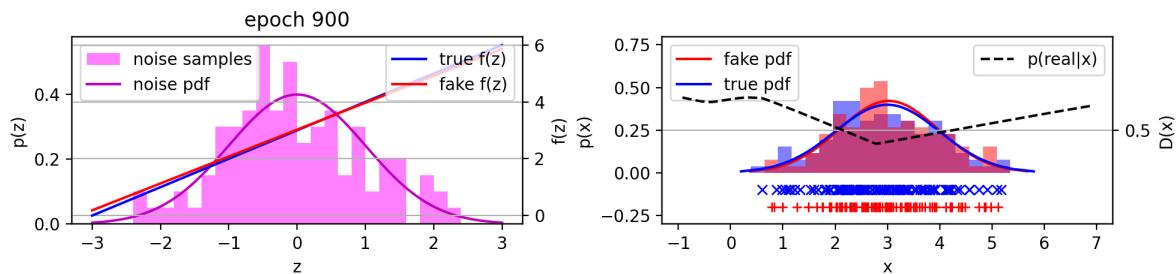
```
In [119]: epochs['figs'][8]
```

```
Out[119]:
```



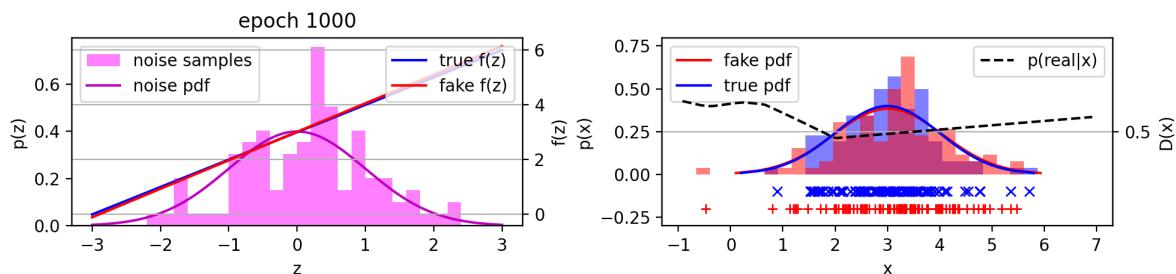
```
In [120]: epochs['figs'][9]
```

```
Out[120]:
```



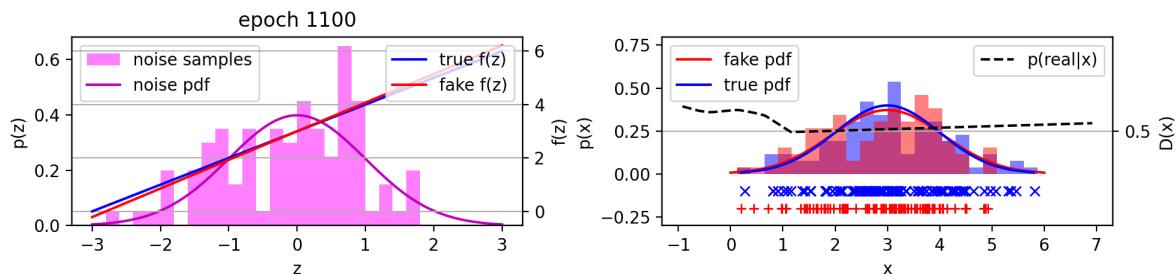
```
In [121]: epochs['figs'][10]
```

```
Out[121]:
```



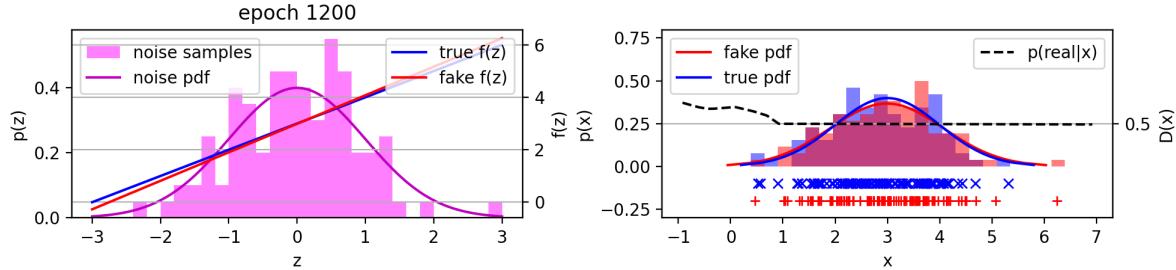
```
In [122]: epochs['figs'][11]
```

```
Out[122]:
```



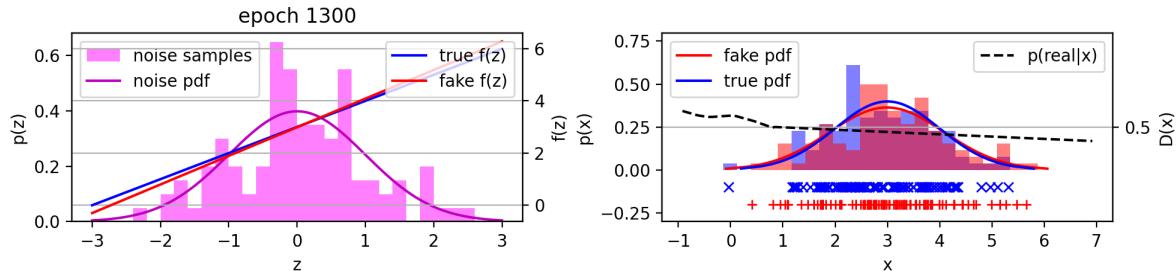
```
In [123]: epochs['figs'][12]
```

```
Out[123]:
```



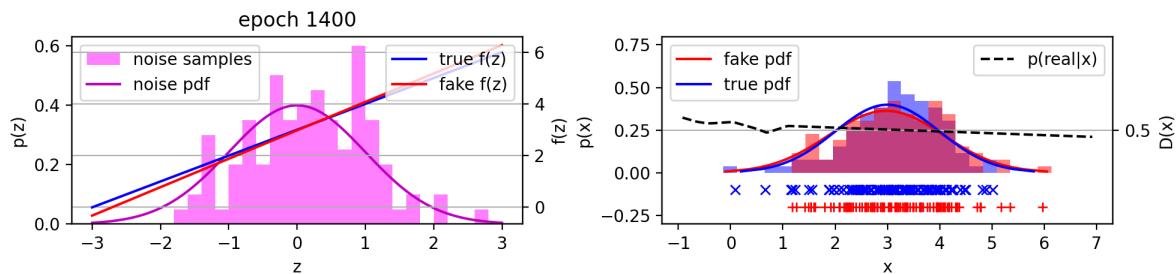
```
In [124]: epochs['figs'][13]
```

```
Out[124]:
```



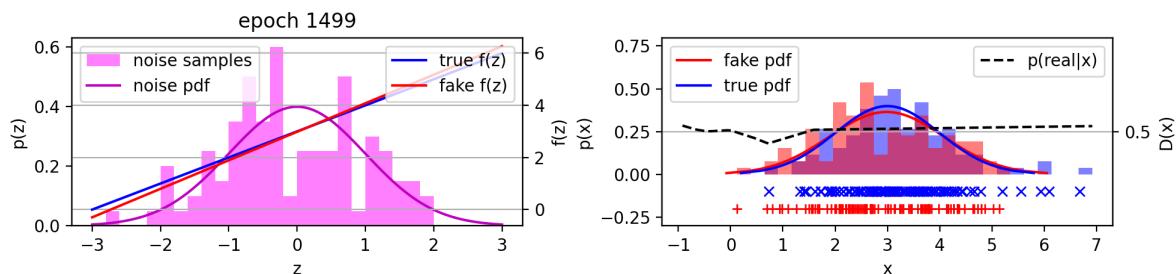
```
In [125]: epochs['figs'][14]
```

```
Out[125]:
```



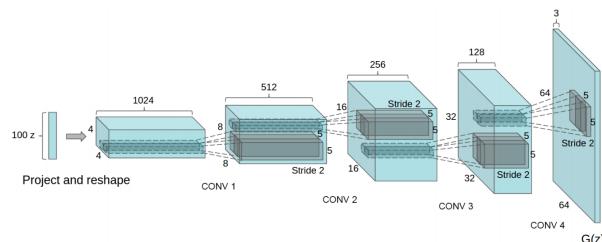
```
In [126]: epochs['figs'][15]
```

```
Out[126]:
```



Example on MNIST

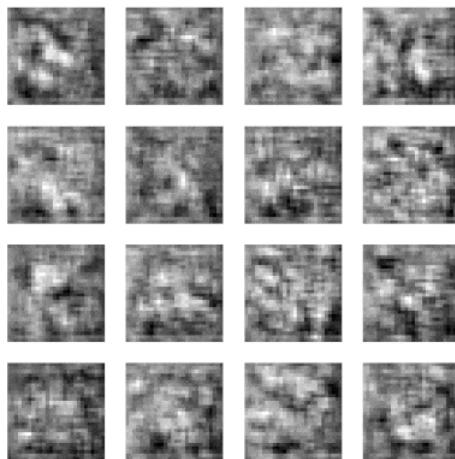
- use a deep convolutional network as the generator (Source: [post](#), [code](#))



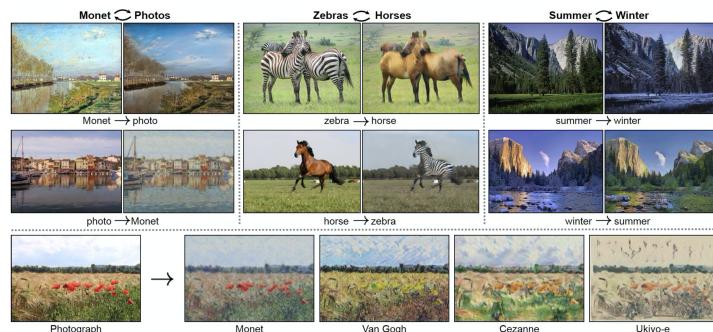
- generated images

z **y** **5** **!** **3** **9** **9** **7** **4** **9**

- generated examples during training



Example: Style transfer (CycleGAN)



Example: Face Synthesis

- Which faces are fake?

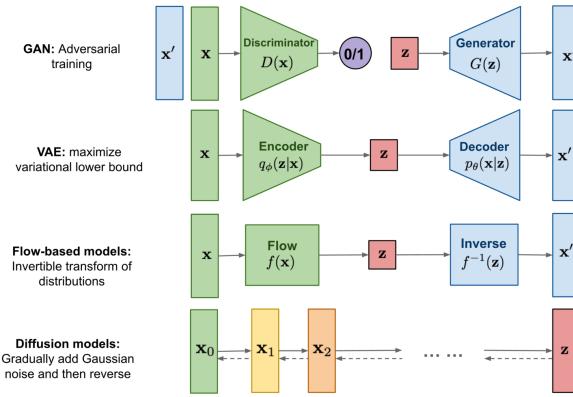


- Try it out: <https://thispersondoesnotexist.com>

Diffusion Models

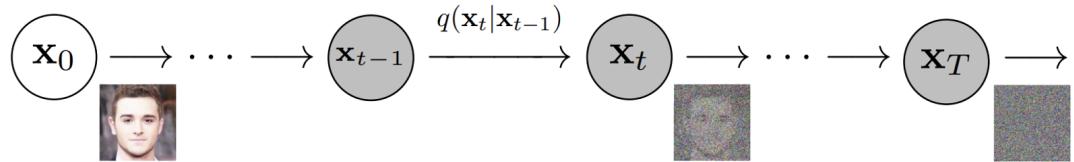
- Define a forward diffusion process that gradually turns an image into noise.
- Learn the reverse process that can gradually remove the noise and recover the image.

- key difference: supervision on each step.



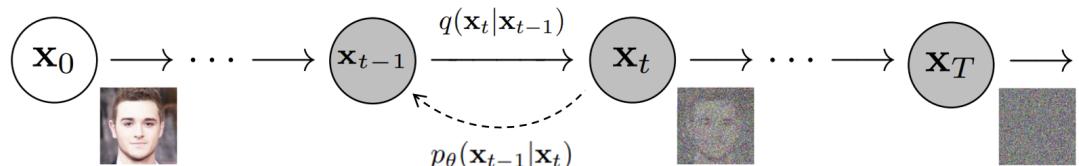
Forward diffusion process

- In each step, fade the image and add Gaussian noise.
 - \mathbf{x}_0 = image
 - $\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \epsilon_t$, $\epsilon_t \sim \mathcal{N}(0, \mathbf{I})$
 - Noise variance schedule: $\{\beta_1, \dots, \beta_T\}$
- Probability distributions
 - for one step: $q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t | \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$
 - for all steps: $q(\mathbf{x}_{1:T} | \mathbf{x}_0) = \prod_{t=1}^T q(\mathbf{x}_t | \mathbf{x}_{t-1})$



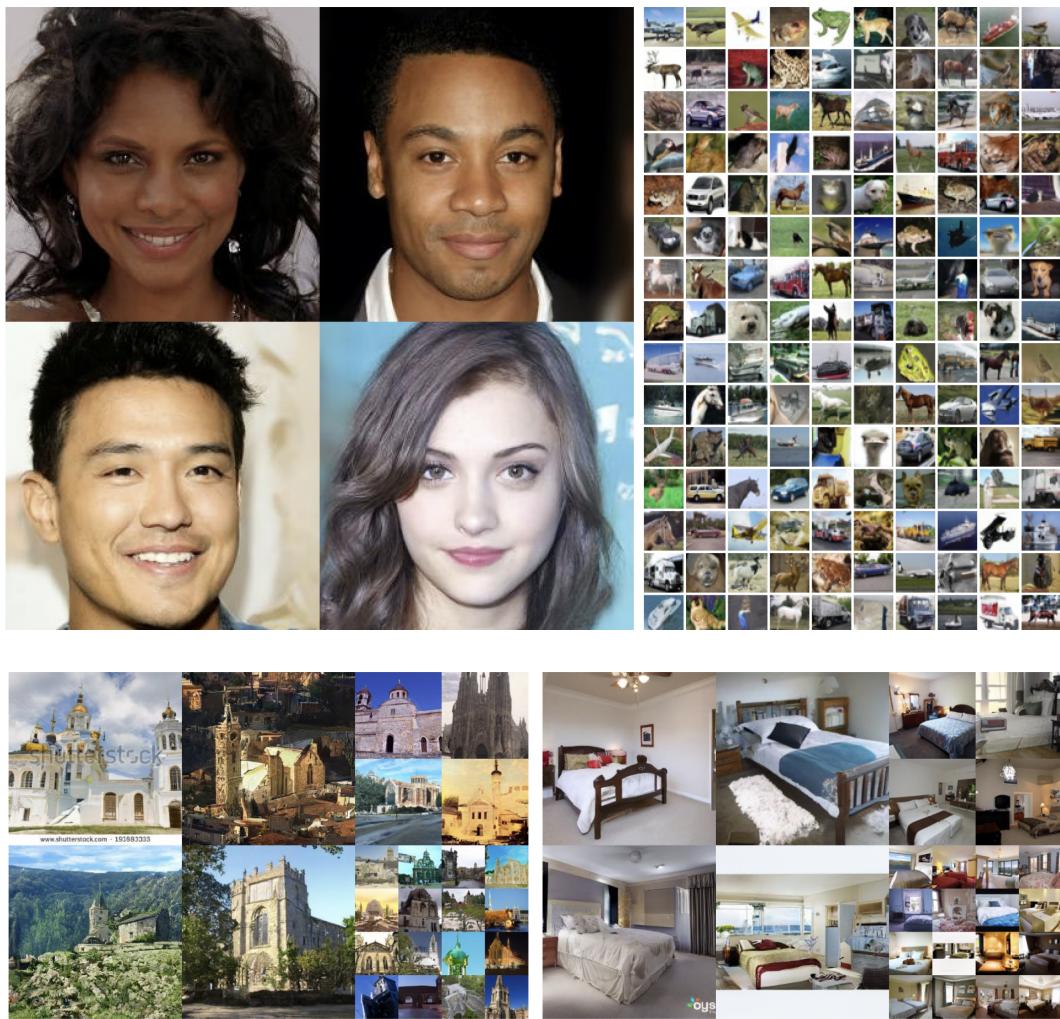
Reverse diffusion process

- We want recover the image from the noise.
 - the true reverse process $q(\mathbf{x}_{t-1} | \mathbf{x}_t)$ is not tractable.
- Approximate it with a parametric model (NN) $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$
 - assume Gaussian: $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1} | \mu_\theta(\mathbf{x}_t, t), \Sigma_\theta(\mathbf{x}_t, t))$
 - for all steps: $p(\mathbf{x}_{0:T}) = p_\theta(\mathbf{x}_T) \prod_{t=1}^T p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$



Training

- minimize an upper-bound on the negative data log-likelihood.
 - $\mathbb{E}[-\log p_\theta(\mathbf{x}_0)] \leq \mathbb{E}_q \left[-\log \frac{p_\theta(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \right] = \sum_{t=0}^T L_t$
 - $L_0 = -\log p_\theta(\mathbf{x}_0 | \mathbf{x}_1)$
 - $L_{t-1} = KL(q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0), p(\mathbf{x}_{t-1} | \mathbf{x}_t))$
 - $L_T = KL(q(\mathbf{x}_T | \mathbf{x}_0, p(\mathbf{x}_T))$
- Usually use a U-net to model $\mu_\theta, \Sigma_\theta$.



Summary

- **Unsupervised Learning**

- Autoencoder - unsupervised dimensionality reduction and clustering.
- Convolutional autoencoder - AE for images.
- Masked Autoencoder (MAE) - self-supervised representation learning
- Variational autoencoder (VAE) - improve interpolation ability, generative model
- Generative Adversarial Networks (GAN) - learn to generate samples from a distribution
- Diffusion model - transform noise into image, step by step.